

Beyond One-Model-Fits-All: A Survey of Domain Specialization for Large Language Models

CHEN LING, Emory University, USA and NEC Labs America, USA

XUJIANG ZHAO^{*†}, NEC Labs America, USA

JIAYING LU^{*}, Emory University, USA

CHENGYUAN DENG^{*}, NEC Labs America, USA and Rutgers University, USA

CAN ZHENG^{*}, NEC Labs America, USA and University of Pittsburgh, USA

JUNXIANG WANG^{*}, NEC Labs America, USA

TANMOY CHOWDHURY^{*} and YUN LI, George Mason University, USA

HEJIE CUI, Emory University, USA

XUCHAO ZHANG, Microsoft, USA

TIANJIAO ZHAO, AMIT PANALKAR, Blackrock Aladdin Financial Engineer Group, USA

WEI CHENG, HAOYU WANG, YANCHI LIU, ZHENGZHANG CHEN, NEC Labs America, USA

HAIFENG CHEN, CHRIS WHITE, NEC Labs America, USA

QUANQUAN GU, University of California, Los Angeles, USA

CARL YANG and LIANG ZHAO[†], Emory University, USA

Large language models (LLMs) have significantly advanced the field of natural language processing (NLP), providing a highly useful, task-agnostic foundation for a wide range of applications. The great promise of LLMs as general task solvers motivated people to extend their functionality largely beyond just a “chatbot”, and use it as an assistant or even replacement for domain experts and tools in specific domains such as healthcare, finance, and education. However, directly applying LLMs to solve sophisticated problems in specific domains meets many hurdles, caused by the heterogeneity of domain data, the sophistication of domain knowledge, the uniqueness of domain objectives, and the diversity of the constraints (e.g., various social norms, cultural conformity, religious beliefs, and ethical standards in the domain applications). To fill such a gap, explosively-increase research, and practices have been conducted

^{*}Equal Contribution.

[†]Corresponding author.

Authors’ addresses: Chen Ling, chen.ling@emory.edu, Emory University, Atlanta, GA, USA and NEC Labs America, Princeton, NJ, USA; Xujiang Zhao, xuzhao@nec-labs.com, NEC Labs America, Princeton, NJ, USA; Jiaying Lu, jiaying.lu@emory.edu, Emory University, Atlanta, GA, USA; Chengyuan Deng, cd751@rutgers.edu, NEC Labs America, Princeton, NJ, USA and Rutgers University, New Brunswick, NJ, USA; Can Zheng, caz51@pitt.edu, NEC Labs America, Princeton, NJ, USA and University of Pittsburgh, Pittsburgh, PA, USA; Junxiang Wang, junwang@nec-labs.com, NEC Labs America, Princeton, NJ, USA; Tanmoy Chowdhury, tchowdh6@gmu.edu; Yun Li, yli38@gmu.edu, George Mason University, Fairfax, VA, USA; Hejie Cui, hejie.cui@emory.edu, Emory University, Atlanta, GA, USA; Xuchao Zhang, xuchaozhang@microsoft.com, Microsoft, Redmond, WA, USA; Tianjiao Zhao, Amit Panalkar, Blackrock Aladdin Financial Engineer Group, New York, NY, USA; Wei Cheng, Haoyu Wang, Yanchi Liu, Zhengzhang Chen, NEC Labs America, Princeton, NJ, USA; Haifeng Chen, Chris White, NEC Labs America, Princeton, NJ, USA; Quanquan Gu, qgu@ucla.edu, University of California, Los Angeles, Los Angeles, CA, USA; Carl Yang, j.carlyang@emory.edu; Liang Zhao, liang.zhao@emory.edu, Emory University, Atlanta, USA, GA.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2018 Association for Computing Machinery.

Manuscript submitted to ACM

Manuscript submitted to ACM

in very recent years on the domain specialization of LLMs, which, however, calls for a comprehensive and systematic review to better summarize and guide this promising domain. In this survey paper, first, we propose a systematic taxonomy that categorizes the LLM domain-specialization techniques based on the accessibility to LLMs and summarizes the framework for all the subcategories as well as their relations and differences to each other. We also present a comprehensive taxonomy of critical application domains that can benefit from specialized LLMs, discussing their practical significance and open challenges. Furthermore, we offer insights into the current research status and future trends in this area.

Additional Key Words and Phrases: Large Language Models, Natural Language Processing, Domain Specialization

ACM Reference Format:

Chen Ling, Xujiang Zhao, Jiaying Lu, Chengyuan Deng, Can Zheng, Junxiang Wang, Tanmoy Chowdhury, Yun Li, Hejie Cui, Xuchao Zhang, Tianjiao Zhao, Amit Panalkar, Wei Cheng, Haoyu Wang, Yanchi Liu, Zhengzhang Chen, Haifeng Chen, Chris White, Quanquan Gu, Carl Yang, and Liang Zhao. 2018. Beyond One-Model-Fits-All: A Survey of Domain Specialization for Large Language Models. 1, 1 (June 2018), 42 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 INTRODUCTION

The evolution of natural language processing (NLP) and artificial intelligence (AI) models has witnessed a remarkable trajectory, beginning with the rule-based systems of the 1950s and 1960s, transitioning to statistical models in the 1990s, followed by the emergence of neural networks in the 2010s. Owing to the success of self-attention and Transformer-based neural network architecture [240], Pre-trained Language Models (PLMs) emerged and swiftly gained popularity in the late 2010s due to their ability to learn universal language representations from large-scale data in an unsupervised manner, which can be beneficial for many downstream NLP tasks such as commonsense reasoning [270], multiple-choice question answering [206], and story generation [30], while avoiding training new models from scratch. In the last few years, with the fast growth of large corpus and hardware capacities, researchers have found scaling up model and training data can continuously improve the model capacity, following the scaling law [99], eventually resulting in Large Language Models (LLMs) [259], such as GPT-3 [28] (175B parameters), PaLM [39] (540B parameters), and LLaMA [235] (65B parameters). LLMs, significantly outperforming smaller models in understanding and generating human-like text, have emerged as a promising AI research trend. Their potential to revolutionize natural and social sciences through efficient literature analysis, novel hypothesis generation, and complex data interpretation could accelerate research, enhance the discovery process, and facilitate interdisciplinary collaboration.

The great promise of LLMs as general task solvers motivated people to extend their functionality largely beyond just a “chatbot” [173], but use it as an assistant or even replacement of human or existing *de facto* tools in specific domains such as healthcare, finance, and education. However, directly applying LLMs for domain-specific problem solving will meet many hurdles. First, there are significant difference in conversation and language styles in different fields, roles, and tasks ranging from medical prescriptions, to legal sentences, to online chatting, etc. The acquisition of such capabilities and experience even require human beings many years of training, a lot of which are hands-on and proprietary. Moreover, different fields, institutions, and teams have their own “business models” about which response will maximize their own utility function for their tasks, which is not directly replaceable by a single general-purpose LLMs solver with no customization. More importantly, the requirement of domain knowledge for professional-level usage also need to be very in-depth, in-real-time, and accurate, none of which can be easily achieved by pre-trained LLMs. Many domain knowledge resources are proprietary assets and core competitiveness of the organizations that can never be leaked to general-purpose LLMs. Last but not the least, languages are constrained by social norms, cultural conformity, religious beliefs, legal requirements, and ethical practice, all of which are changing parameters in different

locations, countries, populations, races, communities, etc., which make general-purpose LLMs impossible to be a one-fits-all solver without any customization. All the above hurdles lead to the necessity of “specializing LLMs into the domains” or “domain-specialization of LLMs”, which is to customize the general-purpose LLMs to the domain contextual data, augmented by domain knowledge, optimized by the domain objective, and regulated by domain constraints. To achieve this goal, this topic is currently experiencing extremely rapid growth.

Domain Specialization of LLMs is a critical yet challenging problem that requires the invention and integration of effective techniques to address the serious challenges caused by its unique characteristics, including: **1) Knowledge Gaps and Domain-specific Expertise.** The power of LLMs is attributed mainly to their massive amount of training corpus. Yet, it also indicates LLMs tend to have a knowledge cut-off (i.e., LLMs lack access to the latest information, events, or discoveries). In many specialized domains, new discoveries, regulations, and best practices continuously emerge, making it difficult for LLMs to stay up-to-date. For instance, more than 30 thousand mainstream news articles are published every day [247]. For social media analysis and fact-checking, LLMs may not handle them since the knowledge extracted from the training corpus is offline. This indicates that regular re-training or continuous learning mechanisms are required to maintain LLMs’ relevance and accuracy in these dynamic fields. However, ensuring the model freshness can be resource-intensive, as it necessitates continuous high-quality and up-to-date data collection, processing, and computationally extensive model re-training. **2) Domain Knowledge Elicitation from LLMs.** LLMs, by default, possess general knowledge across a wide range of topics and may have seen and obtained specific knowledge for most domains. However, more popular or widely-discussed topics may be over-represented, while some domain-specific topics might be under-represented, which makes it difficult to be effectively extracted for domain-specific tasks. In addition, domain-specific tasks often involve complex concepts, specialized terminology, and intricate relationships between different entities. Without proper guidance, LLMs may generate plausible-sounding but inconsistent answers to similar queries (i.e., LLM’s hallucination) or slightly rephrased questions [15]. This issue arises because LLMs are designed to predict the most likely word sequences based on the input rather than providing a definitive answer based on a structured knowledge base. Researchers have found users can guide the model to produce more relevant, accurate, and task-specific responses, enhancing the overall utility and effectiveness of AI systems across numerous domains by providing LLMs with a few task-specific demonstrations [259]. Nevertheless, providing LLMs with adequate demonstrations is not trivial since user instructions can often be vague, incomplete, or ambiguous, making it difficult to discern the intended meaning or desired outcome. Let alone LLMs tend to have a finite context window, typically determined by the maximum token length they can process (e.g., ChatGPT can only handle 4097 tokens). **3) Model Complexity and Extensive Computational Resources for Fine-tuning.** To better adapt to specific domain applications, fine-tuning is historically a commonly used practice to specialize language models. However, different from traditional language models, fine-tuning an LLM needs vast amounts of high-quality, domain-specific data for effective fine-tuning. Acquiring, cleaning, and pre-processing such data can be time-consuming and resource-intensive. Moreover, the sheer complexity of LLMs makes it challenging to identify the most appropriate fine-tuning strategy, as the choice of hyperparameters, learning rate, and training duration can significantly impact the model’s performance. Chen et al. [34] have also discussed fine-tuning LLMs may lead to severe *catastrophic forgetting* since the LLM with a complex architecture is more likely to forget previously learned knowledge and overfits to target domains during fine-tuning. In addition to the data requirement and complex model architecture, LLMs typically consist of billions of parameters, e.g., both Generative Pre-trained Transformer 3 (GPT-3) [28] and Pathways Language Model (PaLM) [39] contains more than 100 billion parameters, which require substantial computational power to train. Fine-tuning or

re-training these models necessitates access to high-performance GPUs or specialized hardware, such as TPUs, which can be expensive and difficult to obtain, especially for individual researchers or smaller organizations.

Over the past few years, significant research has been conducted on domain specialization techniques for LLMs. Many methods focus on generic technical contributions, adaptable to specific domains with minor modifications and access to domain-specific information. However, cross-referencing these techniques across different application domains remains a challenge, as does the absence of a systematic standardization and summary of methods for evaluating various domain specialization techniques. This lack of clarity creates obstacles for non-AI professionals and obfuscates existing bottlenecks, pitfalls, open problems, and potential future research directions. To surmount these obstacles and harness artificial intelligence for more effectively accomplishing tasks across various domains, this survey paper offers a comprehensive and systematic review of the current state-of-the-art LLM domain specialization. The major contributions of this paper include:

- **A systematic categorization and taxonomy of LLMs domain specialization techniques:** We comprehensively classify existing methods based on different levels (i.e., black-box, grey-box, and white-box) of accessibility to the LLM and organize their corresponding techniques into a taxonomy. We discuss details, relationships, pros, and cons among different subcategories. The proposed taxonomy is designed to assist domain experts in identifying the most suitable techniques for their target problem settings.
- **A comprehensive categorization and summarization of major application domains:** We debut the taxonomy of representative application domains that domain-specialized LLMs can enhance. The practical significance and open challenges for each application domain or subdomain are elucidated, allowing for easy mapping to the proposed technique taxonomy. Researchers and various domain experts could cross-reference additional application domains for evaluating their newly proposed methods while expanding their advanced techniques to encompass new application domains.
- **An insightful discussion of the current status of research in this area and future trends.** An overall picture and trends of LLM domain specialization have been outlined and discussed. The paper concludes by presenting fresh insights into the bottlenecks, open problems, as well as a discussion of possible future directions.

1.1 Related Surveys

This section briefly outlines previous surveys relevant to the domain specialization of LLMs in three categories: (1) fundamental overview of PLMs and LLMs; (2). techniques on domain adaptation and generalization of PLMs; and (3). Specializing language models for specific domains.

Fundamental overview of PLMs and LLMs. While comprehensive reviews [157, 193] of PLMs and their use in diverse NLP tasks exist, they don't necessarily apply to LLMs due to differences between the two. Given the recent growth in popularity and effectiveness of LLMs, several review papers have emerged, addressing various LLM aspects. Some focus on fundamental LLM components [270, 292], others on the history and potential applications of generative AI [30, 283], and a few [156] on enhancing LLMs with reasoning capabilities. However, a comprehensive review and technical taxonomy of LLM domain specialization are yet to be provided.

Domain adaptation and generalization of PLMs. Surveys [51, 70] examine how to effectively and efficiently adapt PLMs to specific domains, such as adding a layer to the model or updating the model parameters. However, most of these

techniques don't apply to LLMs because of the inaccessibility of their architecture and parameter space. Also, updating knowledge in LLMs is challenging due to computational costs and the need for efficient optimization strategies.

Specializing language models for specific domains. Recent review papers have emphasized the benefits and necessity of customizing LLMs for specific domains. Risks linked with applying generic LLMs to areas like medical education have been noted in [209], including lack of originality and inaccuracies. Practical considerations for legal domain-specific language models have also been suggested in [217]. In the finance sector, initial steps towards a finance-specialized LLM have shown improved performance on financial tasks without compromising general benchmarks [262]. These advances highlight the need for a comprehensive review and technical taxonomy of domain specialization techniques to assist different sectors in effectively employing LLMs for their unique tasks.

2 TAXONOMY OF DOMAIN SPECIALIZATION

Large language models are typically referred to as large-scale pre-trained language models (PLMs) based on the Transformer architecture [157, 193]. Empirical evidence suggests that scaling pre-trained language models, such as increasing the model size or data size, frequently results in enhanced model capacity for downstream tasks. In this section, we begin by reviewing the fundamental concepts of PLMs and proceed to present a comprehensive taxonomy of existing techniques aimed at specializing large language models for specific domains.

2.1 Background

Specifically, PLM is a type of neural network pre-trained on a large corpus of text data to learn linguistic patterns, structures, and semantics. The input and output of PLMs can be described as follows. In LLMs, the *input* is a text sequence that serves as context for understanding and processing. To clarify the task, a *prompt*, an additional sentence or query, is often included. These prompts, designed based on the NLP task, provide a premise or task explanation. For instance, in text summarization, a prompt like "Summarize the key points in the following passage:" could precede the input passage. The *output* is the text sequence or prediction generated in response to the input. Depending on the task, this could be an answer to a question or a sentiment label, and may require post-processing like token decoding or label extraction for final presentation. As LLMs are typically scaled-up versions of PLMs, they follow the similar architecture design of PLMs, which come in three main flavors: *encoder-only*, *encoder-decoder*, and *decoder-only* architectures. This brief introduction will provide an overview of these PLMs architectures, and discuss their differences and commonalities.

- *Decoder-only Language Models*, like GPT [195], are autoregressive language models that generate the next word in a sequence based on previous words. They map a sequence of tokens to a vector representation and generate contextually relevant content autoregressively, calculating the probability of the next token based on the context. This autoregressive modeling approach is particularly suitable for text generation tasks.
- *Encoder-Decoder Language Models* consist of an encoder that processes input text into vector representations and a decoder that generates output text from these representations. They employ cross-entropy loss as the objective function, comparing the actual and predicted target sequences. These PLMs are often used for sequence-to-sequence tasks like machine translation and summarization, with T5 [196] being a notable example.
- *Decoder-only Language Models*, like GPT [195], are autoregressive language models that generate the next word in a sequence based on previous words. They map a sequence of tokens to a vector representation and generate contextually relevant content autoregressively, calculating the probability of the next token based on the context. This autoregressive modeling approach is particularly suitable for text generation tasks.

2.2 Taxonomy of Domain Specialization Techniques

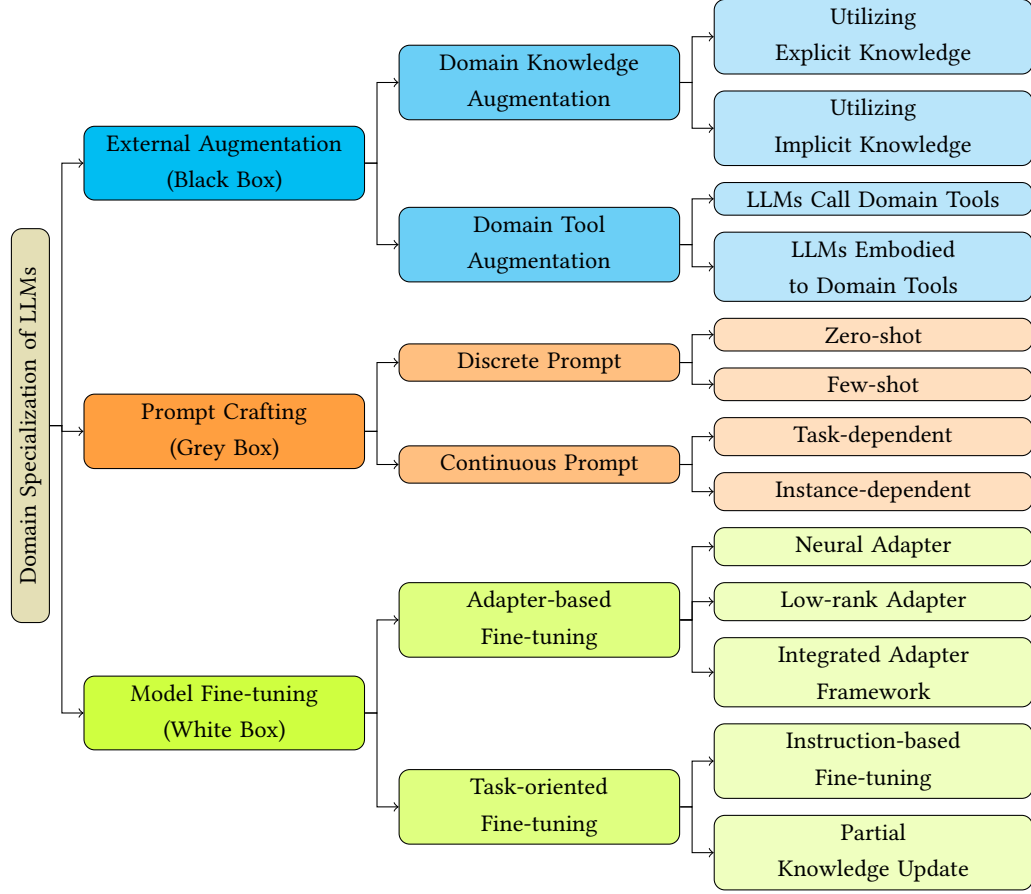


Fig. 1. The taxonomy for current techniques on LLM domain specialization.

Approaches that specialize LLMs into domains are categorized into three approaches according to the level of accessibility to LLMs, namely, no access (black box), partial access (grey box), and full access (white box). Black box typically indicates we only have access to the model API (e.g., ChatGPT and GPT4) without knowing any information but the generated output; grey box denotes we have limited information (e.g., the probability of generated tokens in GPT-3 API), such information can guide us to design and fine-tune a suitable prompt to better elicit domain knowledge; and white box indicates we have full access to the LLM (e.g., LLaMA and its variants), including the parameter setting, training data, and full model architecture.

We provide an overview of each approach in Figure 2. To be more specific, 1) *External augmentation (black box)* does not necessarily require access to the LLM’s inner parameter space, making it the most accessible for users with limited resources (e.g., computational resources, domain-specific data). As shown in Figure 2 (b), by using external resources or tools, domain-specific knowledge is incorporated into the input prompt, generated output, or both, effectively adapting the LLM’s performance without modifying its internal structure. 2) *prompt crafting (grey box)* involves designing various

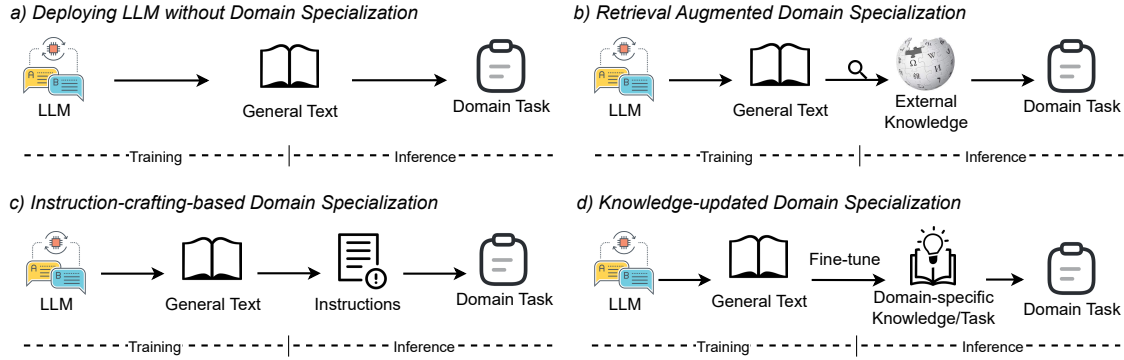


Fig. 2. This exposition discusses different approaches for tailoring LLMs to domain-specific tasks: (a) using an LLM trained on general corpora without modifications, (b) enhancing the LLM’s performance through retrieving relevant external knowledge, (c) utilizing domain-specific and task-relevant instructions to improve LLM’s capabilities, and (d) updating the LLM’s internal knowledge with domain-specific text and tasks.

types of prompts by accessing the gradient or loss values of LLMs, allowing for finer control over the model’s behavior. 3) *model fine-tuning* (white box) demands the most access and resources, as it involves updating the LLM’s parameters to incorporate domain-specific knowledge directly into the model. (Figure 2 (d)).

Relations between approaches in different categories.

- **Different levels of specialization:** Each approach operates at a different level of specialization (i.e., black box, grey box, and white box). Augmenting with external knowledge provides a focused injection of domain-specific information while prompt engineering works at the input level, shaping the model’s inference process. Fine-tuning modifies the LLM’s internal parameters, leading to more profound changes in the model’s behavior.
- **Trade-offs:** The approaches exhibit different trade-offs regarding computational cost, ease of implementation, and generalization. Augmenting with external information and crafting task-specific instructions are often less computationally expensive than knowledge updates of LLMs but may not yield the same level of performance improvement. Fine-tuning and neural adapters can provide more substantial performance gains but can be more challenging to implement and may suffer from reduced generalization capabilities if overfitting occurs.
- **Complementary nature:** The three approaches can be used independently or in combination to achieve better performance on domain-specific tasks. For instance, external knowledge can be integrated with a fine-tuned LLM to leverage both specialized knowledge and optimized parameters. Similarly, carefully designed prompts can be used alongside neural adapters to guide the model’s output while taking advantage of the newly learned domain-specific knowledge.

Researchers can utilize these methods independently or in combination to achieve optimal performance on specific tasks while considering the unique requirements and constraints of each approach. In the following sections, we review representative works in each approach in detail.

3 EXTERNAL AUGMENTATION FOR DOMAIN SPECIALIZATION

Retrieval augmentation aims to enhance LLMs by retrieving relevant information from external sources, without fine-tuning model parameters. There are two primary categories: (1) **Domain Knowledge Augmentation**, where LLMs are provided with domain-specific context from an external knowledge source, and (2) **Domain Tool Augmentation**,

which integrates LLMs with external systems or tools, often via APIs. Domain Knowledge Augmentation supplements the model’s responses with external information, while Domain Tool Augmentation expands the model’s capabilities for tasks it couldn’t perform otherwise. Domain Knowledge improves depth and accuracy within a specific field, while Domain Tools allow the model to perform tasks beyond its inherent abilities. This section discusses both approaches, their limitations, and advantages.

3.1 Domain Knowledge Augmentation

Domain knowledge, in the broadest sense, is a comprehensive understanding of a specific field or subject area. It includes concepts, principles, facts, and patterns that are unique to a particular domain. The knowledge can be represented in various forms, including a set of documents, a domain-specific knowledge graph, or a neural network that contains parametric domain knowledge. Domain knowledge augmentation in LLM specification refers to the process of enriching an LLM’s performance in specific domains by incorporating additional information from domain knowledge. Two categories of external knowledge typically can facilitate LLMs in their domain specialization: **explicit knowledge** refers to knowledge that is clearly defined, easily expressed, and structured in a manner that can be directly understood and utilized; and **implicit knowledge** refers to knowledge that is not directly stated or easily expressed, but is embedded within the data or the system, often in a latent, non-obvious form.

3.1.1 Utilizing Explicit Knowledge with LLM. A conventional method for customizing language models to domain-specific tasks is to retrieve domain-specific information from external context. When presented with an explicit knowledge source containing domain-specific information, it is crucial for LLMs to prioritize the context if the data source holds task-relevant details that contradict the model’s memorized knowledge. This strategy ensures that model predictions are anchored in the context, allowing for the refinement or correction of specific model predictions without the need for frequent retraining.

Current techniques often employ a neural retriever to acquire task-relevant information from either a large corpus (e.g., Wikipedia) or a knowledge base (e.g., Wikidata) [27, 46, 73, 89, 117, 118, 129, 133, 224]. Specifically, given a task-specific query, early works [27, 117, 129, 224] designed neural retrievers to vectorize the query and all information in the external knowledge source to search for relevant information based on various similarity metrics (e.g., cosine similarity) in the latent space. The searched information can then be concatenated with the query for downstream tasks. With the prevalence of LLMs, researchers have been using LLMs to replace the neural network based retriever [46, 89, 197], and Izacard et al. [89] demonstrated that coupling a rather lightweight LLM (around 11 billion parameter size) with an external knowledge base can achieve similar performance when using a 540B LLM (i.e., PaLM). Furthermore, in order to enhance the transparency and explainability of the retrieval, He et al. [73] proposed to leverage LLMs to decompose the information retrieval process with detailed reasoning steps, and Lu et al. [133] explored to utilize LLMs to verify whether information obtained by a pre-trained neural-network-based retriever is relevant or not.

3.1.2 Utilizing Implicit Knowledge with LLM. Implicit domain knowledge in machine learning refers to latent, non-obvious information embedded within data or the system, often represented as vectorized knowledge or embeddings learned during pre-training. Such embeddings capture intricate data patterns, symbolizing domain knowledge in an abstract form. Previous research [58, 67, 155, 244] suggests the use of attention mechanisms to enable PLMs to retrieve task-related information from this implicit knowledge. These studies transform task-specific queries into latent embeddings, calculating attention scores between the query vector and each knowledge entry. A softmax function is used to generate a weight or probability distribution across all knowledge entries concerning the input query. The

retrieved memory vector is then obtained via a weighted sum of the memory entries, using attention weights. This method enhances traditional neural networks with implicit knowledge, permitting the model to access relevant, current information during inference.

While LLMs can store a substantial amount of information in their parameters to generate high-quality responses, augmentation with implicit knowledge isn't always required. Unlike explicit knowledge, implicit knowledge requires extra processing, such as transforming domain-specific data into latent vectors, making it less practical. Despite the limited work in augmenting LLMs with implicit knowledge, researchers are exploring its potential, including its use in storing instructional knowledge about a domain. This approach involves creating an instruction cycle that retrieves the next input prompt from implicit knowledge, parses the LLM's output to recover variable assignments, and stores these back into the memory for retrieving the next instruction. Augmenting LLMs with this instruction cycle allows them to process large inputs and potentially solve complex domain-specific problems [215].

3.1.3 Open Challenges. By incorporating external knowledge, LLMs function like librarians, finding relevant information without needing to memorize all resources. This enhances performance in specialized tasks without extensive retraining, enabling more adaptable and efficient AI systems capable of lifelong learning and knowledge updating. However, augmenting LLMs with external knowledge for domain-specific tasks presents several open challenges.

- (1) *Seamless integration*: Seamless integration of external knowledge into LLMs is crucial, whether the knowledge is explicit or implicit. Existing methods typically concatenate retrieved knowledge to the LLM's input or intermediate layers. However, it's important for the LLM to have the option of accepting or rejecting retrieved information, given that such information may be incomplete or conflicting.
- (2) *Scalability and adaptability*: Designing systems capable of scaling to manage large amounts of domain-specific data and adapting to new or changing information is challenging. With rapidly expanding knowledge bases, computing pairwise knowledge similarity will become increasingly computationally unfeasible.

3.2 Domain Tool Augmentation

Domain tools refer to specialized software, libraries, or frameworks that are developed specifically for a particular domain or field (e.g., NCBI Web APIs for genomics question answering [98], automated formal theorem prover for mathematical proofs [95], sandbox environment for social behavior simulation [176], etc.). These tools are designed to handle domain-specific tasks, data, or knowledge effectively, which often incorporate algorithms, techniques, or data structures that are tailored to the unique requirements of that domain. However, the utilization of these domain tools often demands strict adherence to input formats or extensive training, making them less accessible to general users. On the other hand, LLMs are artificial general intelligence models that exhibit intelligence and cognitive capabilities across a wide range of tasks and domains. Despite their versatility, current LLMs are constrained in tasks that require domain specialization. These limitations [156, 213] include: (1) unstable result formats depending on the random seeds, generation hyperparameters, and input contents [214]; (2) inability to access up-to-date information [165] since LLMs are solely capable of acquiring information from their training data; (3) a tendency to make up facts observed by researchers [93]; (4) lack of precision in certain tasks such as arithmetic [64].

Researchers propose a collaborative integration approach to overcome the limitations of solely using either domain tools or LLMs for complex domain-specific tasks. This approach combines the strengths of both, utilizing domain-specific knowledge, algorithms, and functionalities from the tools, while offering a user-friendly interface through LLMs. This

collaboration optimizes the use of domain-specific resources and eases user engagement by allowing LLMs to guide effective tool usage.




Fig. 3. A toy example for LLMs call domain tools.

LLMs Call Domain Tools. One straightforward way for domain tool augmentation is to allow LLMs to call domain tools. Essentially, this type of approach follows a multi-stage pipeline, given an LLM $f_{\Theta}(\cdot)$ and a domain tool $\mathcal{T}(\cdot)$: (1) elicit an executable command c for the domain tool from the LLM by curated or constructed prompts p , denoted as $c = f_{\Theta}(p)$. (2) execute the command c in the domain tool and get the outputs, denoted as $r = \mathcal{T}(c)$. (3) post-process the domain tool outputs by pre-defined rules or the LLM, denoted by $y = \text{post-process}(r)$.

This pipeline provides a general diagram and can be easily expanded to multi-LLMs multi-tools collaboration scenarios. The key technical challenge is to ensure the instruction-following and validity of generated commands c from LLMs, so that domain tools can accurately solve desired tasks. Most existing works propose to utilize zero-shot or few-shot prompting for executable commands generation (please refer to Sec. 4 for more details). Figure 3 shows a toy example, where the task is to solve an arithmetic question “There are 72 heads and 200 feet inside a cage. How many rabbits and chickens are there?”. To elicit LLMs to generate an executable Python program, we can formulate the prompt as “Please write a Python script to solve the arithmetic question. Question: {question_text}”. Then, a snippet of scripts is returned by LLM as the executable command c for the Python interpreter. Finally, the Python interpreter responds with the program outputs “44, 28”, and further post-processed into desired results “44 Chickens and 28 Rabbits”.

Depending on the types of domain tools, LLMs can generate corresponding commands that adhere to the syntax and format requirements to call them. Many domain tools provide APIs for easy and precise access. Early exploration in this direction is to elicit search engine queries (e.g., WebGPT [165], WizInternet [109], GopherCite [153]) written in natural language or database queries (e.g., Binder-SQL [37], DIN-SQL [188], BIRD [121]) in the programming language. Later, researchers study how to elicit LLMs to write executable codes that can be executed in program interpreters such as Python [35, 64, 231], Wolfram [261], and so on. Other than the widely-used search queries, database queries, and Python programs, there exist many domain-specialized APIs that have unique syntax. For instance, chatGPT plugin system [172] introduces how to utilize tools for travel booking, restaurant reservation, e-commerce shopping, and workflow automation. These API calling scripts are typically generated by zero-shot or few-shot prompting techniques, as stated in the toy example.

Complex tasks often require multiple tools for completion, prompting researchers to generalize LLMs as “task planners” or “API selectors” that utilize various domain tools. These approaches not only generate executable commands but also focus on task decomposition and tool coordination. For instance, the Draft, Sketch, and Prove (DSP) framework is used for automated theorem proofs, employing an LLM or oracle for drafting informal proofs and another for generating formal sketches, with off-the-shelf provers finalizing open conjectures within each sketch [95]. TaskMatrix.AI uses LLMs to create high-level solution outlines for specific tasks and matches subtasks to appropriate off-the-shelf models

or systems [126]. HuggingGPT leverages LLMs as controllers to manage existing domain models for complicated tasks [219]. Lastly, Qin et al. [192] propose a tool-augmented LLMs framework that dynamically adjusts execution plans and effectively completes each subtask using suitable tools.

LLMs Embodied to Domain Tools. LLMs can also be called by domain tools to serve as smart agents in interactive environments, namely *LLMs embodied to domain tools*. LLMs, when embodied in interactive robots, can serve as the decision-making module for domain-specific applications. For example, PROG PROMPT [225] investigates LLMs’ ability to assist robots in completing tasks, when the robot’s perception module observes surrounding objects and the LLM is prompted with available action specifications. Results indicate the LLM can generate situated actions for simulated household and real-world tabletop tasks. Furthermore, Murali et al. [164] employ LLMs as the primary component for identifying different speakers in multiparty conversations involving a social robot. The robotics community is progressively exploring these areas, studying LLM utility in human-robot interfaces, planning, grounding, and more [48, 125, 257]. Furthermore, researchers start to investigate how multiple LLMs can *interact with the environment* or *communicate and collaborate* together for real-world task-solving. Mind’s eye [130] studies how LLMs can benefit from the interaction with simulated physics engines to inject grounded rationale for physics alignment tasks. CAMEL [119] proposes a communicative agent framework to assign different roles to LLM agents so that multiple AI agents can collaboratively communicate by chatting with each other in an instruction-following fashion to solve the specified task. Park et al. [176] utilize twenty-five LLMs as generative agents in a game-based sandbox environment to create believable simulations of human behavior for interactive applications.

3.2.1 Open Challenges. By leveraging the power of LLMs, domain tools can assist in a variety of tasks across multiple fields, including robotics, virtual agents, and problem-solving in real-world scenarios. This allows for more intuitive and seamless human-machine collaboration, leading to increased efficiency and adaptability in tackling complex problems.

Augmenting LLMs with domain tools poses several open challenges:

- (1) *Automated integration:* At present, augmenting LLMs with domain-specific tools requires a significant amount of effort to ensure proper integration. A promising future direction involves utilizing LLMs as a unified interface through standardized protocols to connect various applications and services, thereby enabling seamless communication and interaction between them.
- (2) *Getting rid of domain tools:* Another direction for the future development of LLMs is to focus on creating a powerful artificial general intelligence (AGI) model that is not dependent on external tools or domain-specific knowledge. An AGI model would have the potential to revolutionize the way we use language models, enabling more complex and sophisticated tasks to be performed with greater ease and efficiency.

4 PROMPT CRAFTING FOR DOMAIN SPECIALIZATION

While LLMs trained on large-scale corpora are powerful, further pre-training on prompts can enhance their ability to adhere to user intentions and generate accurate and less toxic responses [174, 253]. Prompts, or task-specific input texts designed to elicit specific model responses, help guide the LLM’s content generation process and set expectations for the desired output. Approaches generally fall into two categories: (1) **Discrete Prompt** involves creating task-specific natural language instructions to prompt LLMs, eliciting domain-specific knowledge from their parameter space, and (2) **Continuous Prompt** uses learnable vectors to prompt LLMs, eliminating the need for manually designed text instructions. This section delves into both approaches and the merits and limitations of domain specialization.

4.1 Discrete Prompt

Recent works [28, 173] allow LLMs to quickly adapt to unseen domains by discrete prompting, and GPT-3 [28] is the first work that introduces how to perform an unseen task using an LLM with zero-shot/few-shot discrete prompts without updating the LLM’s inner parameter. We give a formal definition of the discrete prompt framework below.

Problem Setup. Given an LLM $f_{\Theta}(\cdot)$ where Θ denotes pre-trained model parameters, the task is to elicit desired output \mathbf{y} from LLM with a discrete prompt p and a test query, denoted as $\mathbf{y} = f_{\Theta}([p; c])$, when freezing Θ . It is worth noting that both \mathbf{y} , p and c are sequences of tokens (i.e., natural language sentences). The rationale behind using discrete prompts is that they can serve as instructions to elicit the generalized reasoning abilities of LLMs. By following such instructions, LLMs can perform domain-specific tasks that they have not been specifically trained for. This approach allows LLMs to demonstrate their ability to apply previously learned knowledge to new and diverse situations, thus enhancing their overall effectiveness and utility.

Depending on the prompting types, discrete prompts can be divided into two categories: (1) zero-shot [17, 108, 211], where the prompt p consist of only the task description; and (2) few-shot [140, 158, 260], where the prompt p consists of the task description and few illustrative examples. The key difference between zero-shot and few-shot prompts is whether or not illustrative examples are provided.

4.1.1 Zero-shot Discrete Prompts. The zero-shot setting represents the cold-start scenario, where not a single supportive labeled example is available. Figure 4 presents a toy example of how zero-shot discrete prompts work. The task description that comprises the prompt p can be curated by human users or automatically generated by templates, where the intent of the task and the expected outcomes are described in natural language. However, as stated in [214], post-process is sometimes required to extract the rigorous prediction results from the unbounded raw outputs. Researchers demonstrate that instruction alignment pre-training enables decent zero-shot performance on various unseen tasks [28, 211, 258], where different tasks can be represented in a unified sequence generation format. PADA [17] is one of the pioneering works that explore how to elicit the domain adaptation ability of LLMs for domains unseen during the training phase. PADA first generates the target domain name followed by a set of domain-related features related to the test query, and then uses them together as the prompt to predict task labels. Follow-up works explore how to utilize zero-shot discrete prompt for domain adaptation in sentiment analysis [94], image classification [65], semantic segmentation [56], and rumor detection [128]. Later, Kojima et al. [108] extend the few-shot-Chain-of-Thoughts(Few-shot-CoT) [260] into zero-shot-CoT to elicit multi-step reasoning ability of LLMs. The core idea of Zero-shot-CoT is a two-stage prompting, where the 1st stage simply adds the same prompt “*Let’s think step by step*” before each answer to derive the reasoning process sentences, and the 2nd stage takes the generated reasoning sentences to generate the final answer. Zero-shot-CoT has achieved significantly stronger performance than the standard zero-shot prompting method on arithmetic, symbolic reasoning, and other logical reasoning tasks.

4.1.2 Few-shot Discrete Prompts. The few-shot setting reflects the characteristics of sparse training samples of many domain-specific applications (i.e., only a few annotated examples are available). Figure 5 presents a toy example of how few-shot discrete prompts work. Different from zero-shot prompts, a few examples that further convey the task intention and provide illustrations of the desired output formats are included in the prompt p . Researchers have observed that few-shot prompts yield more stable output formats and more decent performance on downstream tasks [28, 173]. Chain-of-Thought (CoT) [260] improves the domain specialization ability of LLMs by introducing a series of intermediate reasoning steps for complex reasoning tasks, but it also brings extra cost in manually designing

```
# Task description
Please determine if the two sentences entail, contradict, or are neutral to each other.
# Test query
Premise: She emerged vigorous with Apgar of 7 and 8.
Hypothesis: She had low APGAR scores.
Answer:
# LLM response
LLM: Contradiction
```

Fig. 4. An example (adapted from [113]) of zero-shot discrete prompts, where task description, and/or a test query are provided to LLMs. No illustrative examples are provided in zero-shot prompts.

```
# Task description
Please determine if the two sentences entail, contradict, or are neutral to each other. Below are several
examples.
# Example 1
Premise: ALT, AST, and lactate were elevated as noted above.
Hypothesis: The patient has abnormal lfts.
Answer: Entailment
# Example 2
Premise: Chest x-ray showed mild congestive heart failure.
Hypothesis: The patient complains of cough.
Answer: Neutral

# Test query
Premise: She emerged vigorous with Apgar of 7 and 8.
Hypothesis: She had low APGAR scores.
Answer:
# LLM response
LLM: Contradiction
```

Fig. 5. An example (adapted from [113]) of few-shot discrete prompts, where task description, illustrative examples, and/or a test query are provided to LLMs. Compared to zero-shot prompts, a few supportive labeled samples are provided in few-shot prompts.

CoTs for each test example. As a follow-up, Auto-CoT [288] eliminates manual designs by appending the “*Let’s think step by step*” prompt to the given task context and letting LLMs to generate reasoning chains by themselves. Other than the natural language format instructions, CoCoGen [140] studies the programming language format instructions to tackle structured reasoning tasks.

More advanced techniques are then proposed to further improve discrete instruction of LLMs for domain specialization and customization. For instance, *ensemble-based instruction* [134, 251, 252] utilizes multiple different instructions to derive multiple model outputs and then aggregates these outputs to achieve better task performance. Another line of research proposes *recursive instruction* [9, 53, 54, 106, 269, 293] to breaks down a complex unseen task into a series of subtasks that are relatively easier to solve and then employs LLMs with specific instructions for each subtask.

4.1.3 Open Challenges. Utilizing discrete prompts helps LLMs leverage their inherent knowledge to adapt to new and diverse situations. This approach not only demonstrates the flexibility and adaptability of LLMs but also enhances their

overall effectiveness and utility across a wide range of domains and tasks. However, crafting discrete prompts of LLMs for domain specialization poses several open challenges:

- (1) *Effectiveness*: Often the discrete instructions are curated by domain experts or follow some types of templates. It is arguable whether the instructions used are the most effective ones. Therefore, there is a need for evaluation of these instructions. This can be achieved through collaboration between domain experts and data scientists, who can analyze the performance of the LLMs and adjust the instructions accordingly. An automatic evaluation would be even better.
- (2) *Scalability and adaptability*: Automated ways to generate and select/combine discrete instructions without excessive human intervention is another promising direction to improve discrete instructions of LLMs.

4.2 Continuous Prompt

Similar to discrete prompts, the continuous prompt is a sequence of tokens proposed to attach with the input sentence and guide LLMs with extra knowledge but can be learned from the downstream dataset by continuous *prompt tuning*. In this case, the continuous prompt serves as a *soft* parameterized prompt instead of the hard-coded instruction as discrete language phrases. Prompt tuning is to optimize the prompt that adapts an LLM to customized tasks or domains with the preservation of LLM’s general language understanding ability. Here, one can solely update prompt-related parameters, whose quantity is around only 0.01% of the total number of the LLM’s parameters, while freezing the LLM itself during the fine-tuning phase.

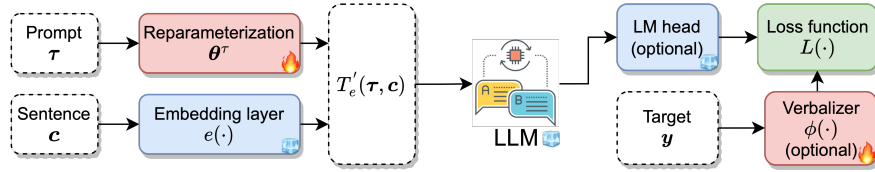


Fig. 6. An illustration of soft prompt tuning. Fire icon represents tunable modules and ice icon represents that parameters of those modules are frozen during tuning. Verbalizer are only used for classification task where a mapping from class label to label words is required, which can be one-one mapping, trainable tokens [72], or enhanced with extra knowledge [83].

A general framework of continuous prompt tuning (Figure 6) can be concisely described in the following stages: (1) Given an input sentence c and its corresponding target y , a template function $T(\cdot)$ organizes them along with a prompt τ of length m into a new sentence $T(\tau, c) = \{\tau_{0:i}, c, \tau_{i+1:m}\}$. (2) Subsequently, the sequence $T(\tau, c)$ is mapped into an embedding space using model’s input layer $e(\cdot)$, resulting in the sequence of token embeddings: $T_e(\tau, c) = \{e(\tau_1), \dots, e(\tau_i), e(\omega_1), \dots, e(\omega_n), e(\tau_{i+1}), \dots, e(\tau_m)\}$, where τ_i is the i -th token in the prompt and $T_e(\cdot)$ denotes the sequence in the embedding space. To perform prompt tuning, τ is considered as pseudo tokens without explicit semantic meanings, and thus $e(\tau_i)$ is replaced with a trainable tensor $h(\tau_i)$ reparameterized by θ^τ . This modifies the template to: $T'_e(\tau, c) = \{h(\tau_1), \dots, h(\tau_i), e(x_1), \dots, e(x_n), h(\tau_{i+1}), \dots, h(\tau_m)\}$. (3) Finally, we can feed the embedding sequence to an LLM, and optimize the continuous prompts θ^τ using the downstream loss function \mathcal{L} as follows:

$$\theta^{\tau*} = \arg \max_{\theta^\tau} \mathcal{L}(f_\Theta(T'_e(\tau, c)), y),$$

where $f_\Theta(\cdot)$ is the LLM function parametrized by Θ . For a cloze-style input reformulated from general tasks, for example, the sentiment analysis task for the sentence “I like the movie!” can be rephrased as a cloze-completion problem: “I like the movie! It was [MASK]”. The predicted words at the masked position are then employed for subsequent classification.

In this case, a unique token [MASK] is integrated during the generation of the template in step (1), and a verbalizer ϕ is required to map class labels to words in the language model’s vocabulary, e.g., *positive* \rightarrow ‘great’, resulting in:

$$\theta^{\tau\star} = \arg \max_{\theta^{\tau}} \sum_c \log P \left([\text{MASK}] = \phi(\mathbf{y}) | T'_e(\tau, c) \right)$$

The information condensed by the prompt falls into two categories: (1) **task-dependent prompt tuning**, and (2) **instance-dependent prompt tuning**. Each category encompasses general and specific enhancements for domain and task adaptation. Although some studies are based on PLMs, the advantages apply to LLMs, given the correlation between prompt tuning enhancements and model size [115] and successful implementations on large-scale PLMs. Moreover, it provides a parameter-efficient, fully controllable tuning method to adapt PLMs for more customized purposes.

4.2.1 Task-dependent Prompt Tuning. Task-dependent prompt tuning optimizes a shared prompt for all instances within a specific task, enabling it to encapsulate information from extensive datasets comprising thousands or millions of examples. However, training a naïve prompt is hard to converge and suboptimal for different scenarios, leaving room for improvement for specific tasks and domains.

Prompt Content Enhancement. We refer prompt content as the embedding values of continuous prompt, enhancements are developed in terms of task-specific initialization and prior knowledge transfer. Pilot works have validated that in contrast to many optimizers that begin with a random distribution applied in general ML tasks, the optimization process of soft prompt is significantly influenced by its initial value. For language models, word embeddings are pre-trained to be quite distinct. Consequently, a standard optimizer such as stochastic gradient descent (SGD) can only update the parameters in a limited vicinity, leading to the possibility of falling into a local minimum [5]. Therefore, a more effective initialization approach would involve using embeddings of concrete task-specific words.

One of the pioneering works, WARP [72] initializes the prompt by the embedding of special token “[MASK]”. KnowPrompt [36] designed learnable prompts as virtual type words and virtual answer words, which are initialized by the aggregated representation of concrete label words and disassembling words based on their frequency in the dataset. In addition, random initialization has been proven to be the least efficient, especially for small model, while Prompt-tuning [115] presented no significant gap between initialization strategies when the model size grows to 11B, indicating that LLMs is robust for prompt’s initialization values in general tasks.

Further studies have revealed that retrained prompts on source domains can enhance performance in unseen target domains, illustrating the ability of prompt transfer [242]. SPoT [242] initialize the prompt with a single generic source prompt learnt from multiple sources tasks, and then fine-tune it on target task in a classic way [115]. PPT [68] also pre-trains a prompt using self-supervised learning on extensive unlabeled corpora, which then serves as the initial prompt for the target task. Su et al. [228] demonstrated the transferability of continuous prompts in both cross-task as well as cross-model settings, and find that a well-initialized prompt can significantly accelerate training convergence. Furthermore, take the advantage of transferability, LFPT5 [191] employed soft prompt for lifelong learning. It continuously trains the prompt that simultaneously learns to solve the current task and generate training samples of previous tasks to overcome the *catastrophic forgetting*. Progressive prompts [201] introduce the prompt tuning into continuous learning. The prompt for current task is defined as the concatenation of prompts that are optimized on previous tasks and a tunable current prompt.

Prompt Construction Enhancement. We refer to the prompt construction about the positioning and length of the prompt, and combinations of additional templates or discrete prompts. Continuous prompts can be simply prepended,

appended, and inserted to the original input sentences without extra language phrases. The pioneering study, WARP [72] adopted all three intersections with a “[MASK]” token for classification tasks. In Prefix-tuning [122], tunable prompts are prepended to the sentence embedding as well as activations of all attention blocks, capitalizing on the left-to-right nature of the autoregressive model: the prepended prompt can efficiently affect the subsequent words through attention mechanisms. Similarly, Prompt-tuning [115] only prepends prompts at the input layer, achieving prompt-tuning results comparable to that of fine-tuned models.

Template is widely used to leverage the adaptation performance [214], for example, reformulating an NLP task (e.g., sentence classification) into the masked words prediction task that is employed during LM pre-training. Based on the predefined task-specific templates, soft prompts can be inserted and thus offers flexibility for conditional tuning. KnowPrompt [36] designed a template appending to the input sentence with a “[MASK]” between subject and object for relation extraction and incorporates trainable prompts of “virtual type words” surrounding these two entities. Output embeddings of ‘virtual type words’ are trained to align logically with the target relation at the masked position, conditioning the optimization with the information of entity type. KiPT [120] developed a knowledge extractor for event detection tasks, which identifies trigger words in sentences based on their semantic similarity to event concepts above a threshold. Identified trigger words as well as the corresponding event labels will then be prepended to a randomly-initialized soft prompt with the input sentence. KiPT also reformulates the sequence tagging tasks: trigger-identification and trigger classification, into the generative task by outputting structured event records.

4.2.2 Instance-dependent prompt tuning. A shared task-dependent prompt is static against the change in input sentence, which ignores semantic difference as well as specific knowledge of individual instance, and thus be suboptimal in fine-grained objectives. Instance-dependent prompt tuning however conditionally generates prompt for individual instances, incorporating both contextual information and task instructions.

Prompt Content Enhancement. Enhancement of prompt content for instance-dependent tuning focus on learning a joint and adaptive representation of tasks as well as instance context. IDPG [263] proposed an additional two-layer perceptron as a prompt generator, down and up project the sentence embedding to the adaptive soft prompt. ATTEMPT [11] first train multiple prompts on large-scale source tasks, and calculate an aggregated prompt base on a sentence-wise attention network, which will then be mixed with a newly initialized target task prompt as the final instance-dependent prompt. Jin et al., [97] assume that prompt tokens differently contribute the instance, and thus designed a look-up module to score the association of prompt tokens to instance tokens, which is then used to calculate the aggregated prompt embeddings. Bhardwaj et al., [19] generate context-aware prompts by a transformer-based sentence encoder, but further quantize the contextual prompt into a more compact representation to avoid optimization collapse. Levine et al., [116] learn the joint representation of prompt and input by a frozen T5 encoder following cross- and self-attention layers. Liu et al., [131] propose an instance-aware prompt that is applied to the intermediate layers of LM. The proposed prompt generator is a simple feed-forward layer with bottleneck architecture which take the embedding of [CLS] token or pooling of embeddings of sentence tokens.

Prompt Construction Enhancement. Similar to the purpose of construction enhancement discussed before, that for instance-dependent prompt tuning introduce instance-dependent knowledge as concrete words or learn adaptive prompt in terms of positioning and length. OntoPrompt [276] enrich the template with instance-related knowledge from external ontology as an additional text, and tune continuous prompts surrounding the “[MASK]” to help prediction. Recently, to give a comprehensive discussion of the effect of content and structure of prompts, dynamic prompting

[273] proposed a unified framework to learn an instance-dependent prompt by dynamically defining prompt position, length, and values for each instance. It also proves the effectiveness of post-fix prompt, given most prior works prepend the prompt to the input sentence.

4.2.3 Open Challenges. Continuous prompt tuning presents a streamlined method to utilize the broad language understanding capacity of LLMs for specific tasks across different domains. It efficiently tackles issues inherent in discrete prompt methods, such as (1) significant reliance on the prompt for LLM performance, where minor wording or template changes can greatly affect the result, (2) computational complexity in identifying the optimal natural language-based prompt from a large search space, and (3) the time-consuming and labor-intensive process of manually designing instructions, particularly in expertise-required domains. However, continuous prompt tuning has its limitations.

- (1) *Interpretability* is often criticized as a weakness of soft prompt tuning. By discretizing the optimal continuous prompts into nearby token vectors in LM's vocabulary, studies such as WARP [72] have found these prompts to be non-interpretable and lacking meaningful content. In KnowPrompt [36] and Prompt-tuning [115], prompt tokens are discovered in close proximity to domain-related terms. For example, Prompts trained on the BoolQ dataset revealed that *science*, *technology*, and *engineering* were the nearest neighbors of the optimal prompt, as approximately 20% of the questions pertain to the "Nature/Science" category [115]. However, the interpretability of continuous prompt as a coherent sequence is still unclear. In addition, continuous prompt is not confined to directing LLMs using compact textual information. OPTIMA [69] achieves domain adapting with prompt by tuning it help regularizes the decision boundary to be smooth around regions where source and target data distributions are similar with an adversarial learning framework.

5 MODEL FINE-TUNING FOR DOMAIN SPECIALIZATION

LLMs, despite being trained on extensive general text data, might not encode adequate knowledge for specific tasks or domains. In such scenarios, fine-tuning the model on a smaller, domain-specific dataset can enhance its performance within that particular area. This fine-tuning can be divided into two main approaches: Adapter-based Fine-tuning and Task-oriented Fine-tuning. (1) **Adapter-based Fine-tuning:** This approach, as illus-




Fig. 7. The visualization of two approaches to fine-tune LLMs based on domain-specific knowledge, where the blue rectangle denotes the set of parameters in LLM. (a) the adapter-based LLM fine-tuning aims to fine-tune LLMs on specific domains with a small number of extra parameters (i.e., adapter); and (b) the task-oriented model fine-tuning aims to fine-tune LLMs based on specific tasks.

trated in Figure 7 (a), employs neural adapters or modular components to enhance the LLM's performance on domain-specific tasks without major modifications to the LLM's inner parameters. These adapters, typically integrated into the existing LLM architecture, allow for task-specific learning while keeping the original model largely intact. (2) **Task-oriented Fine-tuning:** As represented in Figure 7 (b), this method focuses on modifying the LLM's inner parameters to improve alignment with specific tasks. However, entirely updating all parameters of an LLM may be impractical due to hardware limitations and potential performance degradation. Therefore, the challenge for researchers lies in identifying which parameters require alteration within the expansive parameter space, or in efficiently updating a subset of these parameters. These two approaches allow LLMs to be tailored to specific tasks or domains, offering flexibility and efficiency in handling specialized applications.

5.1 Adapter-based Fine-tuning

Adapter-based fine-tuning aims to add a small number of extra parameters to the LLM for achieving better performance in specific tasks. Typically the additional parameters are encoded in simple modules to *guide* the language model’s adaptation to target domains or tasks. The golden spots for the added modules include: (1) simple with a small number of parameters; (2) extensible to the original language models; (3) flexible with sequential training on each specific domain. Most of the proposed strategies with the above favorable properties are built on adapters [79, 202] under the umbrella of parameter-efficient fine-tuning.

5.1.1 Adapters. Adapters are trainable modules inserted between layers of a pre-trained model [79]. The key property of adapters highlights that the parameters of the original language model keep frozen, thus provide sustainable parameter sharing even with varying domains and tasks. Suppose $f_{\Theta}(\cdot)$ denotes the function of LLM parametrized with the set of parameters Θ and $g_{\Delta\Theta}(\cdot)$ denotes the function of adapters with parameter $\Delta\Theta$, then $f_{\Theta} \circ g_{\Delta\Theta}$ represents the fine-tuned language model by adapters. Let X be general input data with task performance metric ϕ , and D be the domain training data with domain-specific task performance ϕ_D (for both ϕ and ϕ_D , a higher value indicates better performance), the goal of adapters is to find $g_{\Delta\Theta}$ such that:

$$\phi(f_{\Theta}(X)) \approx \phi(f_{\Theta} \circ g_{\Delta\Theta}(X)) \quad \phi_D(f_{\Theta}(D)) \leq \phi_D(f_{\Theta} \circ g_{\Delta\Theta}(D))$$

Despite most empirical studies on cross-lingual or multi-task learning, some recent works explore unsupervised domain adaptation particularly using adapters. Unsupervised Domain Adaptation (UDA) using adapters has been explored in recent work, aiming to enhance the cross-lingual or multi-task learning capabilities of pre-trained models. The first attempt [287] targeted multi-domain adaptation with a two-step strategy: domain-fusion training with Masked Language Model (MLM) loss on a mixed corpus, followed by task fine-tuning with a task-specific loss on the domain corpus. Subsequently, UDAdapter was introduced, which also adopted the two-step training and fine-tuning approach, but segregated this into two adapter modules: a domain adapter and a task adapter. The domain adapter first learned domain-invariant representations, which were then concatenated with the task adapter whose parameters were frozen [145]. This was achieved using the architecture defined in AdapterFusion [182]. AdapterSoup further improved adaptation efficiency by adopting a weight-average of domain adapters only during the testing phase [42]. To select domain adapters, three strategies were explored: exhaustive combination, text clustering, and semantic similarity.

Though these works focused on domain specialization, they were evaluated on pre-trained language models like GPT-2 [42, 145, 287], indicating potential applicability to larger language models. To address this, LLaMA-adapter was designed for efficient adaptation on Large Language Models with Adapters (LLaMAs) using self-instruct demonstrations. The adapter architecture incorporated a zero-init attention mechanism, and the domain specialization capability was tested on instruction-following and multi-modal reasoning tasks [286].

As the application of adapters expands, several techniques, while not explicitly claimed as effective for domain specialization, have either demonstrated potential by offering favorable performance on downstream tasks or served as integrated components in existing frameworks for domain specialization. Hence, adapters are usually classified based on their architectures into neural adapters and low-rank adapters. With the objective of facilitating user-friendly implementation, a growing body of work is dedicated to building comprehensive frameworks of different adapters [84, 183]. Certain studies have also shown that adapter integration can yield superior performance across a variety of downstream tasks.

Neural adapters. We call adapters with neural network architectures neural adapters. In their original design, [79] uses a composition of down-projection, GeLU non-linearity [76] and up-projection with the feed-forward layers as the backbone. Later [16] simplifies the architecture to a single hidden-layer feed-forward network and demonstrates the effectiveness on domain adaptation. The adapter modules are inserted after the multi-head attention and feed-forward layers in the transformer. These adapters have been named as bottleneck adapters or serial adapters. We use the latter throughout this paper when referring to [79].

The development of neural adapters naturally takes inspiration from neural network architecture design, such as ResNet, autoencoder, attention mechanism, etc. The adapters used in [182] have an additional residual connection. Soon after, [184] proposes MAD-X framework with invertible adapters, which are inserted adjacent to input and inverted to be fed into the output embeddings. At the high-level, invertible adapters can be considered a mimic of autoencoders. Tiny-attention adapter [290] explores the effectiveness of adapters using attention with tiny per-head dimensionality. Till now, most proposed architectures apply fully-connected layers for down-projection and up-projection in adapters. However, Compacters [100] considers parameterized hypercomplex multiplication layers [281] as an alternative, which has a similar form as a fully-connected layer, but learns a sum of Kronecker products. The main advantage is parameter efficiency. Another way of achieving this is inspired by network pruning, as proposed by SparseAdapter[75] to further reduce the training parameters by pruning at initialization. Note that SparseAdapter is a generic technique applicable to neural adapters. Congregating adapters via insertion can be considered as adaptation *inside* the language models, an alternative is adaptation *outside* the language models. K -adapters [248] proposes to train multiple adapters individually on various knowledge domains, then inject the learned knowledge with language models by concatenation. Recently, Sung et al. [230] raises a concern on the high training memory required because the backpropagation flows through the language model with inserted adapters in entirety. They further propose ladder side-tuning, which only adds small modules on the side of the language model connected to the language model backbone via shortcuts. Both techniques use MLP for demonstration, but keep flexible with different adapter architectures.

Low-rank adapters. Low-rank adaptation (LoRA) [82] is inspired by the observation that large language models reside on an intrinsic subspace [3], where model parameters are efficiently updated. Therefore, learning in this subspace significantly reduces the amount of parameters. LoRA modules implant learnable SVD blocks as the subspace with a low matrix rank $r \ll d$, where d is the dimension of input data. The matrices are added in parallel to the pre-trained weights, thus keeping them frozen during the fine-tuning. Notably, LoRA shows superiority in further reducing the number of trained parameters and introducing no latency during inference.

A follow-up work on this line is DyLora [238], which addresses two issues of LoRA using dynamic search: fixed block size and exhaustive search on the optimal rank. Recently, another concern of LoRA was raised that low-rank modules have limited representation power, and further resolved by the Kronecker adapter (KronA) [55]. The essence is to substitute the SVD modules with a Kronecker product module with two matrices of smaller sizes. Despite not many follow-ups on the low-rank adapters, LoRA modules are included in various integrated adaptation frameworks [74, 84, 148, 254] as an important building block. More details on these frameworks follow below.

Integrated adapter framework. With the flourishing results on effective adapters as introduced above, it is a natural extension to incorporate several adapters of various families to boost their performance. AdapterFusion [182] employs a straightforward idea: train multiple adapters on different tasks and combine the learned embeddings from each adapter with a fusion layer. UniPELT [148] proposes to activate different combinations of methods that best suit the current data or task setup via a gating mechanism. The sub-modules included serial adapter [79], LoRA [82], Prefix-tuning [122]

and Bitfit [278]. Orthogonal to UniPELT, AdaMix [254] stacks multiple adapters of the same type, but avoids more computational cost by training the activation with stochastic routing. AdaMix can be regarded as a general technique that applies to any adapter, despite their implementation on only serial adapters and LoRA.

The idea of learning a routing function on an inventory of adapters further inspires follow-up works. In the context of multi-task learning, Polytropon [187] jointly learns an inventory of adapters and a routing function to re-combine the fine-tuned adapters of various sizes shared among different tasks. Variants of this scheme are further studied by [29], including the replacement of the routing function with weights averaging, or multi-head routing function to achieve better expressivity. On the implementation-oriented aspect, AdapterHub [183] is the most comprehensive and easy-to-use library integrating all mainstream adapters. The only downside, however, is the absence of support of large language models. Recently, LLM-adapters [84] introduces a framework including open-access large language models such as LLaMA, OPT, GPT-J, etc. It subsumes four adapters as basic components (Serial adapter [79], MAD-X [184], Parallel adapter [275] and LoRA [82]) and remains extensible to new modules. The study of domain specialization further explores mathematical reasoning.

5.1.2 Open Challenges. The adapter’s wide applications stem from its modular compatibility with language models, flexible design for integration, and efficient domain-specific data fine-tuning, advancing the adapter-based fine-tuning paradigm. However, these methods have drawbacks. Firstly, the performance of inserted modules can be sensitive to architectural design and size across different tasks and domains, risking insufficient representational power or overfitting on limited data. Secondly, additional modules enlarge the model size, imposing new resource demands and possibly extending inference time. Lastly, as Sung et al. note, the training memory needed by adapter-based methods remains substantial as backpropagation involves the entire model even when previous parameters are frozen [230]. Given these discussions, we outline the open challenges in applying adapters to LLMs for domain specialization:

- (1) *Stability and universality:* The performance of adapters can be subject to various architecture or hyper-parameters applied even on pre-trained language models (PTLMs), thus imposes a question mark on the stability and universality. This concern further extends to LLMs. A deeper understanding on how different adapters match with different task settings would be a significant boost to broader applications of adapters.
- (2) *Computational resources:* Adapters have shown remarkable results with a million-size of parameters on PTLMs. However, it remains unproven if they are enough for LLMs. If more adapter modules (more parameters) are required, then the issue of computational cost can be raised again. Another ideal spot on this issue is to reduce the training memory with novel architecture design or fine-tuning strategy.

5.2 Task-oriented Fine-tuning

Despite these incredible capabilities of LLMs trained on large text corpus, fundamentally improving the model performances beyond few-shot examples and auxiliary adapters still requires updating the inner parameters of LLMs on an extensive amount of high-quality domain-specific datasets. However, fine-tuning an LLM on any (domain) specific tasks poses two challenges: 1) updating LLM’s global knowledge may destroy the in-context learning ability due to reasons including but not limited to overfitting, catastrophic forgetting, and task-specific biases [255]. 2) fine-tuning LLMs is computationally expensive due to the vast parameter space and the deep model architecture. In this section, we review recent techniques on how to update the global knowledge of LLMs, which can be primarily categorized into two areas: **Instruction-based Fine-tuning** and **Partial Knowledge Update** to address both challenges, respectively.




Fig. 8. The overview of fine-tuning an LLM with explicit instructions across various domains and datasets. Particularly, the LLM is fine-tuned on a collection of tasks (e.g., commonsense reasoning, information extraction, etc.) with detailed instructions, and the fine-tuned LLM is expected to obtain problem-solving skills.

5.2.1 Instruction-based Knowledge Update. Instruction-based Knowledge Update refers to updating an LLM’s parametric knowledge by fine-tuning LLMs on a diverse set of tasks with explicit instructions or prompts, which is conceptually the same as Instruct Learning introduced in [174]. An illustration of fine-tuning an LLM with instructions is provided in Figure 8, where an LLM is fine-tuned on a collection of tasks across the whole NLP application domain, and the LLM is deployed on the held-out and unseen tasks. Wei et al. [258] provided the very first attempt to fine-tune LLMs based on a collection of datasets described via instructions. Empirically, effective instructions can substantially improve zero-shot performance on unseen tasks. The instruction-tuned language model FLAN is fine-tuned on a 137B LLM over 60 NLP datasets using natural language instruction templates. The study shows that FLAN outperforms its unmodified counterpart and even surpasses both zero-shot and few-shot 175B GPT-3 on most unseen tasks. Subsequently, in recent works by Chung et al. [43], Menick et al. [154], explicit instructions have been employed to fine-tune LLMs, with emphasis placed on (1) expanding the number of tasks, (2) enlarging the model’s size, and (3) fine-tuning on chain-of-thought data. As a result, the fine-tuned LLM attains state-of-the-art performance on numerous benchmarks in both zero-shot and few-shot NLP tasks.

Fine-tuning with Human Instructions. Fine-tuning with human instructions aims to guide LLMs towards generating safer, truthful, less toxic content in line with user intentions. Most LLMs utilize autoregressive approaches, making the generated content largely influenced by the training corpus distribution and less controllable. Reinforcement learning from human feedback (RLHF) is a notable technique for aligning LLM content with human needs [41]. In RLHF: 1) LLMs create multiple content options for a prompt, ranked by humans for quality, relevance, and desired output alignment; 2) an external reward model assigns scores to content based on rankings, capturing evaluator preferences; 3) model policy is updated using reinforcement learning techniques to maximize expected reward, fine-tuning the model to better align with human preferences; 4) this process of content generation, ranking, reward modeling, and policy optimization repeats in iterations, with the model continually learning from human feedback. Existing methods successfully apply RLHF to fine-tune LLMs on complex reasoning tasks using human instructions [181, 246].

Potential Limitations of Instruction-based Knowledge Update. Knowledge updates based on explicit instructions tend to perform well on Natural Language Understanding tasks but are limited to simpler instructions and struggle with tasks diverging from evaluation sets. Improving adaptability to diverse tasks often incurs catastrophic forgetting. A crucial question is extending model knowledge and abilities without causing such forgetting. Recently, Huang et al. proposed a method that uses a pre-trained LLM to generate high-confidence, rationale-augmented answers for unlabeled questions, improving general reasoning without ground truth labels or explicit instructions [85]. Additionally, Scialom et al. are

expanding LLM knowledge and abilities without forgetting previous skills by fine-tuning LLMs across various tasks and introducing an approach to counter catastrophic forgetting with Continual Learning via Rehearsal [216, 220].

5.2.2 Partial Knowledge Update. Other than leveraging task-specific instructions to fine-tune LLMs, a number of approaches emerge to conduct LLM fine-tuning by updating/editing a part of LLM parameters that link to specific knowledge without leveraging external guidance. Suppose $f_{\Theta}(\cdot)$ denotes the function of LLM parametrized with the set of parameters Θ and $\theta \in \Theta$ is the single parameter in Θ . Updating the inner knowledge of $f_{\Theta}(\cdot)$ based on a collection of training data D is denoted as:

$$\tilde{\Theta} = \Theta + \nabla f_{\Theta}(D) \odot T, \quad T^{(i)} = \begin{cases} 1, & \text{if } \theta^{(i)} \in \Theta_T \\ 0, & \text{if } \theta^{(i)} \notin \Theta_T \end{cases} \quad (1)$$

where T denotes the mask vector and $T^{(i)} \in T$ denote the i -th element of T . The mask controls the amount of LLM's inner knowledge to be updated in each fine-tuning iteration, where we use $\Theta_T \subseteq \Theta$ to denote the parameters that need to be updated in Θ . In the conventional setting of fine-tuning pre-trained language models [80, 205, 295], $|\Theta| = |\Theta_T|$. However, updating all of the parameters is computationally prohibited and resource-consuming in the context of LLM. Empirically, $|\Theta| \gg |\Theta_T|$, which refers to the modification of only a small number of parameters. Existing parameter-efficient knowledge update can be categorized into three streams: i.e., **Knowledge Editing** aims at directly locating and updating a small subset of parameters in an LLM; **Gradient Masking** aims at masking out the gradients of non-relative parameters during the fine-tuning; and **Knowledge Distillation** focuses on obtaining a child model with domain-specific knowledge from LLMs.

Knowledge Editing. Recent research has seen success in updating LLMs with new memories to replace outdated information or add specialized domain knowledge. For instance, improving the ability to update an outdated prediction like "Boris Johnson is Prime Minister of the UK" can enhance an LLM's reliability and generalization. Various methods have been proposed to locate and edit an LLM's parametric knowledge [45, 49, 77, 151, 152]. De Cao et al. proposed a hyper-network trained to update LLM parameters with a single fact needing modification, avoiding fine-tuning to prevent performance degeneration [49]. However, later works found that hyper-network-based editing falters as the LLM scales up, proposing retrieval-based methods to store edits in explicit memory and reason over them to adjust LLM predictions [159, 160]. Other methods focus on localizing and understanding LLM internal mechanisms. Notable works identify crucial neuron activations for LLM factual predictions through attention mechanisms and causal interventions, successfully updating domain facts [45, 151, 152]. A recent method is proposed learning a map from textual queries to fact encodings in an LLM's internal representation, using these encodings as knowledge editors and probes [77].

Gradient Masking. Gradient masking is a technique used to selectively update specific parts of an LLM during the fine-tuning process. The main goal is to reduce computational overhead and potentially mitigate issues such as catastrophic forgetting or overfitting, particularly when adapting pre-trained models to smaller or specialized datasets. Gradient masking involves modifying the gradients during back-propagation by applying a masking function (Equation (1)). This function determines which parts of the model will be updated, effectively *masking* the gradients for certain parameters and keeping them unchanged. The choice of parameters to mask can be based on various criteria, such as their relevance to the task, importance in the model, or contribution to the overall loss.

Earlier attempts [96, 278] have been made to efficiently fine-tune relatively small language models by utilizing various regularization techniques, their methods cannot easily adapt to fine-tuning LLMs. This is primarily due to the

substantially larger amounts of data and computational resources required to train LLMs effectively, which can be several orders of magnitude more than what is needed for smaller language models. To add gradient masks to LLMs, CHILD-TUNING [267] utilizes the downstream task data to detect the most task-related parameters as the child network and freezes the parameters in non-child network to their pre-trained weights. Moreover, Zhang et al. [284] propose a Dynamic Parameter Selection algorithm for efficiently fine-tuning LLMs, which adaptively selects a more promising sub-network to perform staging updates based on gradients of back-propagation, which brings great improvement in domain-specific downstream tasks under low-resource scenarios.

Knowledge Distillation. While most works on LLM self-knowledge update focus on task-specific instructions and parameter efficiency, a promising area of research explores distilling domain-specific knowledge from LLMs into smaller networks to reduce inference latency and enhance domain-specific task solving ability. Muhamed et al. compressed a 1.5 billion parameter LLM to a 70 million parameter model for Click-through-rate prediction, introducing twin-structured BERT-like encoders and a fusion layer for a cross-architecture distillation from a single LLM, resulting in superior performance in both online and offline settings [163]. Similarly, [13, 149, 243] employ a knowledge distillation module for LLM fine-tuning, achieving faster convergence and better resource utilization. This module leverages pre-trained parameters for quick convergence and trains a small subset of parameters to address model over-parameterization. Furthermore, [81, 222] distill the step-by-step chain-of-thought reasoning abilities of larger models into smaller models.

5.2.3 Open Challenges. Fine-tuning LLMs with the latest data ensures that they provide relevant and accurate information, especially in domains where rapid changes occur, such as technology, medicine, and current events. Furthermore, we have observed different applications or users may have unique requirements or preferences. However, fine-tuning the large-scale LLMs also poses several open challenges:

- (1) *Compliance with regulations:* In most cases, updating and fine-tuning LLMs are necessary to ensure compliance with specific regulations or guidelines, such as data protection laws or industry-specific requirements. The so-called *LLM alignment* can be accomplished during the fine-tuning phase.
- (2) *Computational resources:* Fine-tuning or updating inner knowledge of LLMs necessitates access to high-performance GPUs or specialized hardware, which can be expensive and difficult to obtain, particularly for individual researchers or smaller organizations. Pursuing fine-tuning efficiency is still a practical yet essential problem.

6 APPLICATIONS OF LLM DOMAIN SPECIALIZATION

In this survey paper, we explore the applications of LLMs across a range of domain-specific tasks in social sciences (e.g., education, finance, law), natural sciences (e.g., biomedicine, earth science), and applied sciences (e.g., human-computer interaction, software engineering, and cyber security). To achieve domain specialization for LLMs in these diverse fields, readers can employ various techniques, such as external augmentation, instruction crafting, and knowledge update. These approaches can help tailor LLMs to specific tasks and challenges in each domain, enabling more accurate, relevant, and effective applications. Although each domain has its unique challenges and requirements, several common applications of specialized LLMs are shared across these fields:

- *Advanced information extraction:* They can identify entities, relationships, and events from domain-specific texts, such as recognizing genes in biomedical literature or detecting legal clauses in contracts.
- *Text generation and summarization:* They can generate high-quality, domain-specific content and create accurate summaries of complex domain-specific texts.

- *Data-driven predictions and recommendations*: They can analyze domain-specific data for forecasting and providing recommendations, like predicting financial trends or suggesting personalized medical treatment plans.
- *Conversational agents and expert systems*: They can be incorporated into conversational agents or expert systems for domain-specific guidance, such as virtual tutors or legal chatbots.
- *Automated code generation and analysis*: In software engineering, they can generate or analyze code, identify bugs, or suggest improvements based on natural language descriptions.

In this section, we dive deep to review existing techniques on specializing LLM in domain-specific tasks and discuss potential open challenges in detail.

6.1 Biomedicine

Language models are becoming increasingly useful in the field of biology, from fundamental biomedical research to clinical healthcare support. At the fundamental biomedicine science level, LLMs can be trained on vast amounts of domain-specializing data (e.g., genomic and proteomic) to analyze and predict biological functions, disease mechanisms, and drug discovery. LLMs can also aid in predicting protein structures and interactions, which are critical for understanding cellular processes and designing new drugs. At the clinical healthcare support level, pre-trained or medical corpus fine-tuned LLMs can be used for the natural language processing of medical records to identify patterns, make diagnoses, and provide personalized treatment recommendations. Also, LLMs can assist in medical image analysis in a multi-modality learning way, such as identifying specific features in X-rays or MRI scans. Overall, LLMs offer tremendous potential for advancing biology research and improving healthcare outcomes.

6.1.1 Fundamental Biomedicine Science. Recent advancements in LLMs have provided promising results in fundamental biomedical research [142]. These models allow for the integration of diverse data sources, such as molecular structures, genomics, proteomics, and metabolic pathways, enabling a more comprehensive understanding of biological systems. Several application-specific LLMs have been developed for molecular and biological sciences. For instance, MoLFormer [207] is a large-scale molecular SMILES transformer model with relative position embedding that enables the encoding of spatial information in molecules. Nucleotide Transformer [47] is a foundation model pre-trained on DNA sequences for accurate molecular phenotype prediction, essentially treating single amino acids or stretches of k-memorized DNA as words in a vocabulary and using LLMs to learn from the biological sequences. ESM-2 and ESMFold are transformer protein language models included in the Evolutionary Scale Modeling (ESM) family [204], which can generate accurate structure predictions directly from a protein’s sequence and outperform other single-sequence protein language models. ProtGPT2 [59] is a language model trained on the protein space that generates de novo protein sequences following the natural principles. ProGen [141] is a deep-learning language model that can generate protein sequences with predictable functions across large protein families. Artificial proteins fine-tuned to distinct lysozyme families using ProGen have shown similar catalytic efficiencies as natural proteins with low sequence identity.

6.1.2 Clinical Healthcare Support. LLMs have greatly assisted in medical text and electronic health records understanding, leading to a range of potential benefits for healthcare [7, 92, 161, 167, 198, 199, 297]. Recently, there have been several major domain-specific LLMs developed for various NLP tasks in the biomedical domains. BioGPT [136] is a domain-specific generative Transformer language model pre-trained on large-scale biomedical literature, which achieves impressive results on most biomedical NLP tasks spanning document classification, relation extraction, question answering, etc. BioMedLM [1] is another state-of-the-art GPT-style language model trained on biomedical literature

from The Pile¹, capable of achieving high accuracy on various biomedical NLP tasks, including MedQA, and possesses language generation capabilities, developed for research purposes to advance the development of biomedical NLP applications and responsible training and utilization of domain-specific language models. [234] explore the utilization of LLM synthetic data for clinical text mining. Med-PaLM [226] also generates accurate, helpful long-form answers to consumer health questions, as judged by panels of physicians and users, which helps in treatment recommendations to suggest customized medications to patients based on their unique health circumstances. This can assist physicians in selecting the most efficient drugs for their patients, ultimately enhancing treatment outcomes. GatorTron [274] are models from scratch using the BERT architecture of different sizes on a large corpus consisting of over 90 billion words from clinical narratives, and scientific literature. [272] further developed a GatorTro model from scratch on learning unstructured electronic health records. Besides, LLMs can also directly help to anticipate the potential illnesses that a patient may develop based on their past medical history. For example, Med-BERT [200] produces pre-trained contextualized embeddings on large-scale structured electronic health records for disease prediction. ChatCAD [249] proposes a method for integrating LLMs into medical-image Computer-Aided Diagnosis (CAD) networks with the goal of creating a more user-friendly system for patients and improving CAD performance. LLM also brings opportunities and challenges for domain educators or policymakers [60, 110, 111, 223, 250]. On public health informatics, the use of LLMs can also help in analyzing the conversations about the impact of the COVID-19 pandemic on societies worldwide [210], which are taking place on social media platforms such as Twitter [32].

6.1.3 Future Research Opportunities. In biomedicine domain, the utilization of LLMs faces significant challenges due to the scarcity of high-quality labeled data. The complexity and heterogeneity of biomedical data across different tasks make proper annotation time-consuming and expensive, requiring domain-specific expertise [113]. LLMs trained on specific biomedical datasets may lack generalizability to diverse patient populations and clinical contexts. Hallucination, the generation of inaccurate or incorrect information, is a major concern with LLMs in biomedicine, which can lead to harmful errors in real-world applications [12, 18]. For instance, false-positive relationships between drugs and diseases in biomedical literature-trained LLMs can result in incorrect treatment recommendations or adverse effects [112]. Despite these challenges, the use of LLMs in biomedicine offers promising opportunities for research and clinical practice. Recent studies demonstrate that LLMs like GPT-3 can effectively capture the semantics and context of biomedical language, facilitating more accurate representation of biomedical knowledge [4, 168, 226]. LLMs enable integration and analysis of diverse biomedical data sources, supporting the development of precise medical ontologies and taxonomies [133]. They also enable the identification of new phenotypic and genotypic correlations and the discovery of novel drug targets. To mitigate the risk of generating inaccurate information, various approaches can be employed, such as utilizing curated datasets, implementing sampling-based self-check methods [147], conducting causality analysis [10], and incorporating external knowledge to enhance LLMs [180]. By addressing these challenges and leveraging the strengths of LLMs, the use of LLMs in biomedicine can revolutionize healthcare delivery and improve patient outcomes.

6.2 Earth Science

Earth science is an interdisciplinary domain focused on examining the interactions between physical and human systems across diverse spatial and temporal scales. This field incorporates methods from Earth observation, information science, spatial analysis, complexity theory, and simulation modeling to investigate phenomena like climate change, land-use change, natural disasters, environmental development, and urbanization. Spatial information is vital to Earth

¹<https://pile.eleuther.ai/>

science, and geographic information science tools are invaluable for a wide range of interdisciplinary studies involving spatial data. Large language models like ChatGPT can act as question-answer systems, assisting those interested in Earth Science to gain pertinent knowledge, such as recommending the optimal earth observation dataset for specific research purposes, offering code examples like Google Earth Engine code for processing satellite data, providing high-quality responses to environmental-related questions [294], developing innovative idea [21], generating climate scenario [21]. LLMs can also be tailored to various Earth Science-related downstream tasks through methods such as fine-tuning, few-shot, or even zero-shot learning.

6.2.1 Geospatial Semantic Tasks. A study [143] evaluated GPT-2 and GPT-3 on two geospatial semantic tasks: toponym identification and location description recognition. The former task involves detecting named locations within a given text snippet, and the latter entails identifying more detailed location descriptions, like home addresses, highways, roads, and administrative regions, from a text snippet such as tweets. The findings suggest that task-agnostic LLMs can surpass task-specific, fully-supervised models in both tasks, with significant improvement in few-shot learning scenarios. Moreover, [144] broadened this comparison by incorporating four differently-sized pre-trained GPT-2 models, the latest GPT-3, InstructGPT, and ChatGPT, as well as various supervised, task-specific benchmarks for the same two representative geospatial semantic tasks. The outcomes showed that LLMs, with only a limited number of few-shot samples, can outperform fully supervised, task-specific models in well-defined geospatial semantic tasks. This underscores the potential of LLMs to considerably lessen the reliance on custom architectures or extensive labeled datasets for geospatial semantic tasks. However, further research is necessary to determine the most effective ways to create appropriate prompts for guiding LLMs in specific geospatial semantic tasks.

6.2.2 Spatial Classification and Time-Series Forecasting. LLMs are increasingly used for various geospatial artificial intelligence tasks, including classification and time series forecasting, through a prompt-based learning approach. In a study by Mai et al. [144], the frequencies of different types of Point of Interest (POIs) within an urban region were transformed into a textual paragraph. The researchers employed LLMs to predict the urban function of the region based on this paragraph. The study compared GPT2, GPT-3, and ChatGPT models under zero-shot and one-shot settings with two supervised learning neural networks. The results indicated that ChatGPT exhibited superior precision and recall compared to other GPT models. Another example is PromptCast [268], a prompt-based time series forecasting model. The authors conducted forecasting on three real-world datasets: city-temperature data, safe graph human mobility, and electricity consumption load data. The findings in both single- and multi-step forecasting settings demonstrated the potential of prompt-based time series forecasting with language generation models. Moreover, PromptCast exhibited better generalization ability in the zero-shot setting compared to conventional numerical-based forecasting. Additionally, Mai et al. [144] utilized LLMs to forecast the number of deaths due to dementia in different geographic regions, including state and county levels.

6.2.3 Text-Guided Earth Observation Image Processing. LLMs enable multimodal analysis, integrating text, audio, images, and videos with other foundation models like vision models. In a study by Mai et al. [144], VLFMs (OpenCLIP, BLIP, and OpenFlamingo-9B) were employed for an urban perception task using a street view image noise intensity dataset. While most VLFMs faced challenges in connecting visual features to nuanced semantics and concepts in urban studies, BLIP demonstrated balanced and reasonable predictions, comparable to fine-tuned CNN models. VLFMs can comprehend certain characteristics of urban neighborhoods from visual inputs, but their language-only abilities are not as strong as LLMs. Zhang et al. [285] developed Text2Seg, a pipeline that facilitates remote sensing image semantic

segmentation tasks using text prompts. The pipeline incorporates three methods: Grounding DINO generates bounding boxes from a text prompt, which are used in SAM to create a segmentation map; CLIP Surgery produces a heat map from a text prompt, sampled to create point prompts for SAM to generate segmentation masks; and SAM generates multiple segmentation maps, compared to the semantic similarity of the text prompt using CLIP. The initial assessment highlighted the benefits of integrating multiple foundation models in a single pipeline [144, 285].

6.2.4 Future Research Opportunities. Although LLMs have shown promise in some language-only geospatial tasks as zero-shot or few-shot learners, they have also been criticized for generating inaccurate and misleading results. This issue is particularly problematic in a geographic context, where generating geographic faithful results is crucial for nearly all GeoAI tasks [144]. LLMs face challenges in more complex geospatial semantics tasks, such as toponym resolution/geoparsing and geographic question answering, as they are unable to conduct implicit spatial reasoning grounded in the real world. Additionally, integrating LLMs with other foundation models to perform tasks involving multiple data modalities remains a significant challenge in this field.

6.3 Finance

Utilization of NLP technologies in the financial industry, especially the fintech domain, is wide and growing [266] [62]. Common NLP tasks in the trading and investment management domain include sentiment analysis [227], question answering with generative language model [256], named entity recognition (NER) [6], text classification [291].

6.3.1 Sentiment Analysis Task in Finance. The value of financial instruments is often influenced by public information about the underlying entities. Various forms of financial texts, such as news articles, analyst reports, and SEC filings, are all potential sources of new information that could affect market prices. As a result, the use of NLP techniques for sentiment analysis to predict price movements and returns has become increasingly popular in recent years [40].

FinBERT, a fine-tuned BERT model for finance, has emerged as a useful tool for financial sentiment analysis. While the general BERT model has proven to be more effective than traditional dictionary-based approaches for financial sentiment analysis [114], FinBERT has been particularly useful for addressing the unique challenges posed by the specialist language used in financial contexts [8]. The FinBERT model was pre-trained on a financial corpus and then fine-tuned specifically for sentiment analysis tasks, yielding significantly better results than previous state-of-the-art approaches. The ChatGPT model, a newer entrant to the field, has also shown promising results, outperforming BERT and earlier models such as GPT-1 and GPT-2 in studies on predicting stock prices using sentiment analysis of news headlines [132]. It has also been utilized to assess news outlet credibility [271], adding another dimension to its applications in the financial sector. Just like BERT, ChatGPT could encounter challenges when dealing with the specialist language used in finance, indicating a substantial opportunity for fine-tuning this model to make it a dominant tool for financial sentiment analysis tasks.

6.3.2 Financial Data Generation. Generative language models are found to be particularly helpful in sales and client engagement processes in the finance industry, either in consumer finance, insurance, or investment finance [194]. For example, financial institutions have deployed customer service chatbots to provide instant responses to customer inquiries. The Chatbots have mostly been developed with the general BERT model, seq2seq, word2vec or machine learning-based classification models with historical question and answer data [264]. For financial products salespersons, generative language models are used widely to automate sales emails, build sales call agendas and summarize sales

notes for client insights and discover product trends. With arising LLM, fine-tuned LLM with historical client inquiries data or domain financial contexts will be promising enhancements to current generative tasks in finance.

6.3.3 Future Research Opportunities. The release of GPT-3 sparked curiosity around its potential uses in various specific fields, such as tax consultancy [24] and investment portfolio construction [170]. Initial assessments indicated that, while GPT-3 performed better than random guessing, it still fell short of human performance in these tasks. However, the debut of Bloomberg’s domain-specific LLM, BloombergGPT, demonstrated significant improvements in performance on financial tasks without compromising its performance on general LLM benchmarks [262]. This has fueled further interest in domain-specific LLMs in the finance field, prompting predictions of future development of models with similar capabilities, whether through training on specific financial data or fine-tuning existing models. Despite these promising results, the adoption of LLMs in finance is not without its challenges. Notably, data security and intellectual property concerns are of significant importance in this sector. Companies are rightfully apprehensive about the possibility of their proprietary data being leaked during the model training process. Furthermore, questions arise around the ownership of intellectual property when GPT-generated code is incorporated into commercial software products. As the field progresses, these concerns will need to be addressed to facilitate the wider application of LLMs in the financial sector.

6.4 Law

General PLMs have shown great promise in many NLP tasks, including in the legal domain. However, due to the intricate nature of legal language, these generic models might fall short [31]. To improve accuracy in legal applications, Chalkidis et al. proposed LEGAL-BERT, trained specifically on legal corpora, resulting in enhanced legal text classification performance [31]. Prasad et al. further showed that fine-tuning LEGAL-BERT can improve its effectiveness even in unfamiliar legal domains [189]. Despite these improvements, Valvoda et al. cautioned that LEGAL-BERT might produce asymmetric predictions in legal outcomes, which could influence real-life legal decisions [239]. These findings have driven NLP researchers to develop specialized language models such as LLMs for nuanced fields like law. Current efforts can be classified into two main categories: one focusing on LLM domain adaptation, such as legal prompting, and the other examining the transformative potential of LLMs in legal applications through empirical evaluations.

6.4.1 Legal Prompting. LLMs, like GPT-3, have been explored for their potential in handling legal tasks, including statutory reasoning and legal judgment prediction. Legal prompting has been employed to guide and assist LLMs in these tasks, given its significance in the legal domain. For instance, Blair-Stanek et al. investigated the performance of GPT-3 on statutory reasoning using different instruction methods, including zero-shot prompting, few-shot prompting, and chain-of-thought prompting [23]. The study revealed that chain-of-thought prompting achieved the highest accuracy, although legal reasoning techniques surpassed this by further enhancing accuracy [23]. The researchers utilized templates and schemas to guide the reasoning process of GPT-3, resulting in improved accuracy in legal reasoning tasks [23]. Similar performance improvement was observed for legal judgment prediction. Yu et al. [277] used chain-of-thought prompting to guide the GPT-3 model and demonstrated improved performance on this task. Specifically, they provided GPT-3 with legal case summaries and asked it to predict the outcome of the case. The authors found that the model could accurately predict the outcome with the help of legal prompting. Trautmann et al. [236] demonstrated that zero-shot legal prompt engineering performed better than all baselines, but fall short compared with supervised approaches.

6.4.2 Legal Assessment. The development of LLMs has attracted the attention of legal scholars who have explored their potential to revolutionize the legal field. One potential application of LLMs is to assist in legal education and preparation for law exams. Bommarito et al. [25], Choi et al. [38], and Katz et al. [102] have conducted case studies to evaluate the performance of LLMs on law exams with the help of legal prompting. The results indicate that LLMs can achieve promising performance with the assistance of legal prompting. However, it is important to note that LLMs are not expected to replace lawyers in the legal profession. One limitation of LLMs in the legal domain is their difficulty in identifying recent case law due to outdated data sources. Macey et al. [138, 139] and Iu et al. [88] discussed this issue and suggested that LLMs need to be trained on up-to-date legal databases to improve their performance in this area.

6.4.3 Future Research Opportunities. The use of LLMs in the legal domain has shown promise, but further research is needed to fully understand their potential and limitations. One concern is the generation of human-like texts that resemble copyrighted works, raising intellectual property issues [229]. Future work could focus on developing LLM-based tools to detect copyright infringement by leveraging external copyright databases. Data privacy is another challenge, as LLMs are trained on extensive datasets that may contain personal and sensitive information. This raises concerns about unintentional disclosure or reconstruction of protected data, requiring compliance with data privacy legislation like the GDPR [203]. Potential solutions include providing prompts to remind LLMs to keep protected data private or fine-tuning them using external datasets without private information. Bias and fairness are also concerns in LLMs, as they may learn and perpetuate biases from training data. This can affect decision-making in legal contexts, such as legal argument development or case outcome prediction. Addressing this issue may involve prompt-based guidances to correct biases or referencing external unbiased case databases. These areas warrant further investigation to ensure the responsible and ethical use of LLMs in the legal domain.

6.5 Education

LLMs revolutionize education in various aspects, including healthcare research [209], medical education [78, 105, 171], library management [135, 241], journal and media education [177], engineering education [190], academia [87, 137], autodidactic education [61], and higher education [208, 282]. The introduction of LLMs benefits both students and teachers [14, 101]: for students, the utilization of LLMs facilitates the training of question-asking skills [2], engages students in solving real-world problems [146, 280], enhances learning pharmacology and algebra as a self-studying instrument [169, 175]. For teachers or faculties, LLMs have great potentials to assist in teaching skills such as mechanical writing [20] and reducing teaching workloads such as the quality assessment of short answer questions [162] and course design [71]. Despite the potential benefits of LLMs, many existing works demonstrate the limitations of LLMs such as the difficulty in understanding advanced subjects [63] and failure to achieve difficult problems [218]. Some preliminary investigations utilized prompting techniques to improve the performance of LLMs [127]. This will inspire more future works to explore domain specification techniques for education, which is still an open research area. Existing research works in this area can be classified into two categories: assessing LLMs' performance on various education tasks and addressing concerns raised by LLMs' introduction to educational ethics.

6.5.1 Capability Assessment. A significant area of research examines the effectiveness of LLMs in various subjects, aiming to evaluate their potential as learning assistants. The debate over the capabilities of LLMs on medical exams remains a hot topic in the field. While some case studies suggest that LLMs performed worse than students in medical exams [60, 86, 250], other experiments indicate their promising performance on medical exams, either through the chain of thought prompting [127] or without additional fine-tuning [66, 111, 150]. A similar controversy also exists in

mathematics, where the performance of ChatGPT varies considerably depending on the problem requirement [218, 296]. Furthermore, some studies have shown that the mathematical capabilities of ChatGPT fall below those of average math graduate students [63]. In contrast, Kortemeyer, Bordt, and Luxburg found that ChatGPT can pass undergraduate-level exams in physics and computer science [26, 110]. Bommarito et al. have also demonstrated ChatGPT’s improved performance on the CPA examination [24].

In contrast to most assessments on the exam performance of LLMs, Tack and Piech investigated the pedagogical ability of GPT-3 and discovered that it was quantifiably worse than real teachers on several pedagogical dimensions [233]. Despite these mixed results, the potential of LLMs to assist students in their learning remains a significant area of research. By providing instant feedback and personalized learning experiences, LLMs can help students identify their strengths and weaknesses and gain a better understanding of the material. Moreover, they can aid teachers in delivering personalized instruction, developing course content, and reducing the burden of administrative tasks.

6.5.2 Educational Ethics. LLMs offer many possibilities for enhancing students’ learning experiences, but they also present ethical challenges and risks. One significant risk is scientific misconduct, where students could use LLMs such as GPT-3 to generate convincing and seemingly original texts without additional learning or effort [50]. This potential for academic dishonesty could result in serious repercussions for both students and educational institutions. Additionally, students might use LLMs like ChatGPT to author their essays instead of composing them independently, undermining the core principles of education [107]. ChatGPT’s capacity to think critically and produce highly realistic texts with minimal input further exacerbates this threat to educational integrity [232]. Moreover, Khalil and Er noted that ChatGPT can create sophisticated, original texts that remain undetected by prominent plagiarism detection tools, heightening the risk of academic misconduct [103]. Several papers propose potentially effective strategies to combat plagiarism and dishonesty. For instance, designing open-ended assessments that encourage originality and creativity can promote academic honesty [44]. Incorporating oral presentations, group projects, and hands-on activities requiring students to demonstrate their knowledge and skills in a more engaging, interactive way can also help foster honesty and reduce reliance on LLMs [107]. By deploying such strategies, educators can cultivate an environment that encourages originality, creativity, and academic integrity, while still leveraging the advantages of LLMs.

6.5.3 Future Research Opportunities and Challenges. Specializing LLMs in the education domain presents numerous opportunities for future research and innovation. From the perspective of *external resource augmentation*, LLMs can be utilized to search for customized educational content and resources tailored to individual learners’ needs, preferences, and abilities, facilitating more effective and personalized learning experiences. From the perspective of *prompt crafting*, LLMs can be tailored to generate step-by-step explanations for the user provided question. In terms of *model fine-tuning*, a specialized LLMs can be employed as an automatic homework grader of one specific subject. Additionally, LLMs can also be fine-tuned to detect potential plagiarism and academic misconduct.

6.6 Software Engineering

With the development of AI, the pair programming paradigm led to coin the term “AI pair programmer”. And the LLM’s capability to respond with different types of prompt crafting brought huge possibilities for using LLM as a compelling and promising component for software developers. Moreover, pre-trained LLMs can outperform the non-pre-trained model whether it is code understanding or code generation tasks [279]. The ongoing research can be diversified into two broader perspectives. Whether one group is trying to use a generalized LLM and the other group is more focused on using the LLM trained on code (LLMC).

6.6.1 Generalized LLM. An early study on GPT-3 proved the hypothesis that AI-generated code can reduce the software production time [166]. But still, generalized LLM's performance is not promising in software development. For example, in the study [91], ChatGPT faced difficulties with low-resource or distant language and was able to answer only 37.5% correctly. Despite this, dialogue-based LLMs can still be a promising aid in solving programming exercises in the life sciences, particularly in the field of bioinformatics [186]. With expert-drafted codebooks, GPT-3 performance can be improved substantially compared to expert-coded result [265]. Another way to boost performance is using multiple LLMs to distribute the workloads. By following a self-collaborating framework, three ChatGPT roles (i.e., analyst, coder, and tester) can improve the traditional coding generation of LLMs [52].

6.6.2 LLM Trained on Code. Large language models trained on code (LLMC) are the specialized versions of LLM trained on code, such as GPT-J-6B [245], GPT-Neo [22], Codex (can be applied to almost all programming tasks) [33], Alphacode (problems that require deeper reasoning) [123]. There are many studies on LLM-based code completion frameworks focusing on cost reduction [289], security implication [178], code explanation [212] or potential extension in different domains (for example, reverse engineering [179], programming syntax error [185] etc.). By utilizing the Codex, the model CodexDB can process the SQL queries. They utilized the DB catalog and text-to-SQL methods to make the SQL query from the natural language instruction [237]. The Codex can also be used to generate documentation-specific code examples by combining contexts from multiple sources (source code, documentation, and error logs from the compiler) [104]. Again, LLMCs are not limited to coding only. In a task where structure follows the code-like pattern (meaning representation [221]), LLMCs can follow the same performance as LLM.

6.6.3 Future Research Opportunities. The LLM's semantic understanding from the given prompt is not identical to the human level. It can lead to producing different codes from logically the same but semantically different prompts [124]. Moreover, program synthesis also needs to make sure of the performance, security, licensing attribution, and multi-modal specifications [90]. On the other hand, LLM updates can lead to producing inconsistent code for the same prompt [57]. So, although promising, the adaptation of LLM still requires further advancement before it can be dependably used for software development in various contexts including code summarization, code repair, code translation, code generation, code search, program understanding, program debugging (i.e., detection, localization, and repair), code explanations, etc.

7 CONCLUSION

In conclusion, the rapid advancement of large language models has sparked significant interest in harnessing their potential to tackle domain-specific tasks in various natural, social, and applied science fields. However, several challenges, such as limited domain-specific expertise, knowledge elicitation, and model complexity, hinder the direct application of LLMs in these domains. This survey systematically categorizes and summarizes existing domain specialization techniques based on their access level to the LLM, along with a comprehensive taxonomy of application domains that can benefit from specialized LLMs.

By offering a detailed analysis of the advantages, disadvantages, and relationships among different techniques and domains, this survey aims to assist domain experts in identifying suitable techniques for their target problem settings, while also providing data scientists with a clear understanding of the practical significance and open challenges in various application domains. Moreover, the paper highlights the current status of research in this area, shedding light on future trends and potential avenues for interdisciplinary collaboration. As the field of LLM domain specialization

continues to evolve, this survey serves as a valuable resource for researchers and practitioners, fostering further advancements and innovations in the application of artificial intelligence across diverse domains.

REFERENCES

- [1] [n. d.]. <https://crfm.stanford.edu/2022/12/15/pubmedgpt.html>. <https://crfm.stanford.edu/2022/12/15/pubmedgpt.html>.
- [2] Rania Abdelghani, Yen-Hsiang Wang, Xingdi Yuan, Tong Wang, Hélène Sauzéon, and Pierre-Yves Oudeyer. 2022. GPT-3-driven pedagogical agents for training children’s curious question-asking skills. *arXiv preprint arXiv:2211.14228* (2022).
- [3] Armen Aghajanyan, Luke Zettlemoyer, and Sonal Gupta. 2020. Intrinsic dimensionality explains the effectiveness of language model fine-tuning. *arXiv preprint arXiv:2012.13255* (2020).
- [4] Monica Agrawal, Stefan Hegselmann, Hunter Lang, Yoon Kim, and David Sontag. 2022. Large language models are few-shot clinical information extractors. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*. 1998–2022.
- [5] Zeyuan Allen-Zhu, Yuanzhi Li, and Zhao Song. 2019. A convergence theory for deep learning via over-parameterization. In *International Conference on Machine Learning*. PMLR, 242–252.
- [6] Julio Cesar Salinas Alvarado, Karin Verspoor, and Timothy Baldwin. 2015. Domain adaption of named entity recognition to support credit risk assessment. In *Proceedings of the Australasian Language Technology Association Workshop 2015*. 84–90.
- [7] Fares Antaki, Samir Touma, Daniel Milad, Jonathan El-Khoury, and Renaud Duval. 2023. Evaluating the performance of chatgpt in ophthalmology: An analysis of its successes and shortcomings. *medRxiv* (2023), 2023–01.
- [8] Dogu Araci. 2019. Finbert: Financial sentiment analysis with pre-trained language models. *arXiv preprint arXiv:1908.10063* (2019).
- [9] Simran Arora, Avani Narayan, Mayee F Chen, Laurel Orr, Neel Guha, Kush Bhatia, Ines Chami, and Christopher Re. 2023. Ask Me Anything: A simple strategy for prompting language models. In *The Eleventh International Conference on Learning Representations*.
- [10] Vahan Arsenyan and Davit Shahnazaryan. 2023. Large Language Models for Biomedical Causal Graph Construction. *arXiv preprint arXiv:2301.12473* (2023).
- [11] Akari Asai, Mohammadreza Salehi, Matthew E Peters, and Hannaneh Hajishirzi. 2022. Attempt: Parameter-efficient multi-task tuning via attentional mixtures of soft prompts. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*. 6655–6672.
- [12] Amos Azaria and Tom Mitchell. 2023. The Internal State of an LLM Knows When its Lying. *arXiv preprint arXiv:2304.13734* (2023).
- [13] Zhangir Azerbayev, Ansong Ni, Hailey Schoelkopf, and Dragomir Radev. 2022. Explicit Knowledge Transfer for Weakly-Supervised Code Generation. *arXiv preprint arXiv:2211.16740* (2022).
- [14] David Baidoo-Anu and Leticia Owusu Ansah. 2023. Education in the era of generative artificial intelligence (AI): Understanding the potential benefits of ChatGPT in promoting teaching and learning. *Available at SSRN 4337484* (2023).
- [15] Yejin Bang, Samuel Cahyawijaya, Nayeon Lee, Wenliang Dai, Dan Su, Bryan Wilie, Holy Lovenia, Ziwei Ji, Tiezheng Yu, Willy Chung, et al. 2023. A multitask, multilingual, multimodal evaluation of chatgpt on reasoning, hallucination, and interactivity. *arXiv preprint arXiv:2302.04023* (2023).
- [16] Ankur Bapna, Naveen Arivazhagan, and Orhan Firat. 2019. Simple, scalable adaptation for neural machine translation. *arXiv preprint arXiv:1909.08478* (2019).
- [17] Eyal Ben-David, Nadav Oved, and Roi Reichart. 2022. PADA: Example-based Prompt Learning for on-the-fly Adaptation to Unseen Domains. *Transactions of the Association for Computational Linguistics* 10 (2022), 414–433.
- [18] Gernot Beutel, Eline Geerits, and Jan T Kielstein. 2023. Artificial hallucination: GPT on LSD? *Critical Care* 27, 1 (2023), 1–3.
- [19] Rishabh Bhardwaj, Amrita Saha, and Steven C. H. Hoi. 2022. Vector-Quantized Input-Contextualized Soft Prompts for Natural Language Understanding. In *Conference on Empirical Methods in Natural Language Processing*.
- [20] Lea Bishop. 2023. A computer wrote this paper: What chatgpt means for education, research, and writing. *Research, and Writing (January 26, 2023)* (2023).
- [21] Som S Biswas. 2023. Potential Use of Chat GPT in Global Warming. *Annals of Biomedical Engineering* (2023), 1–2.
- [22] Sid Black, Leo Gao, Phil Wang, Connor Leahy, and Stella Biderman. 2021. Gpt-neo: Large scale autoregressive language modeling with mesh-tensorflow. *If you use this software, please cite it using these metadata* 58 (2021).
- [23] Andrew Blair-Stanek, Nils Holzenberger, and Benjamin Van Durme. 2023. Can GPT-3 perform statutory reasoning? *arXiv preprint arXiv:2302.06100* (2023).
- [24] Jillian Bommarito, Michael Bommarito, Daniel Martin Katz, and Jessica Katz. 2023. GPT as Knowledge Worker: A Zero-Shot Evaluation of (AI) CPA Capabilities. *arXiv preprint arXiv:2301.04408* (2023).
- [25] Michael Bommarito II and Daniel Martin Katz. 2022. GPT Takes the Bar Exam. *arXiv preprint arXiv:2212.14402* (2022).
- [26] Sebastian Bordt and Ulrike von Luxburg. 2023. Chatgpt participates in a computer science exam. *arXiv preprint arXiv:2303.09461* (2023).
- [27] Sebastian Borgeaud, Arthur Mensch, Jordan Hoffmann, Trevor Cai, Eliza Rutherford, Katie Millican, George Bm Van Den Driessche, Jean-Baptiste Lespiau, Bogdan Damoc, Aidan Clark, et al. 2022. Improving language models by retrieving from trillions of tokens. In *International conference on machine learning*. 2206–2240.
- [28] Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. 2020. Language models are few-shot learners. *Advances in neural information processing systems* 33 (2020), 1877–1901.

- [29] Lucas Caccia, Edoardo Ponti, Lucas Liu, Matheus Pereira, Nicolas Le Roux, and Alessandro Sordani. 2022. Multi-Head Adapter Routing for Data-Efficient Fine-Tuning. *arXiv preprint arXiv:2211.03831* (2022).
- [30] Yihan Cao, Siyu Li, Yixin Liu, Zhiling Yan, Yutong Dai, Philip S Yu, and Lichao Sun. 2023. A comprehensive survey of ai-generated content (aigc): A history of generative ai from gan to chatgpt. *arXiv preprint arXiv:2303.04226* (2023).
- [31] Ilias Chalkidis, Manos Fergadiotis, Prodromos Malakasiotis, Nikolaos Aletras, and Ion Androutsopoulos. 2020. LEGAL-BERT: The Muppets straight out of Law School. In *Findings of the Association for Computational Linguistics: EMNLP 2020*. 2898–2904.
- [32] Emily Chen, Kristina Lerman, Emilio Ferrara, et al. 2020. Tracking social media discourse about the covid-19 pandemic: Development of a public coronavirus twitter data set. *JMIR public health and surveillance* 6, 2 (2020), e19273.
- [33] Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. 2021. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374* (2021).
- [34] Sanyuan Chen, Yutai Hou, Yiming Cui, Wanxiang Che, Ting Liu, and Xiangzhan Yu. 2020. Recall and learn: Fine-tuning deep pretrained language models with less forgetting. *arXiv preprint arXiv:2004.12651* (2020).
- [35] Wenhui Chen, Xueguang Ma, Xinyi Wang, and William W. Cohen. 2022. Program of Thoughts Prompting: Disentangling Computation from Reasoning for Numerical Reasoning Tasks. *arXiv preprint arXiv:2211.12588* (2022).
- [36] Xiang Chen, Ningyu Zhang, Ningyu Zhang, Xin Xie, Shumin Deng, Yunzhi Yao, Chuanqi Tan, Fei Huang, Luo Si, and Huajun Chen. 2021. KnowPrompt: Knowledge-aware Prompt-tuning with Synergistic Optimization for Relation Extraction. *Proceedings of the ACM Web Conference 2022* (2021).
- [37] Zhoujun Cheng, Tianbao Xie, Peng Shi, Chengzu Li, Rahul Nadkarni, Yushi Hu, Caiming Xiong, Dragomir Radev, Mari Ostendorf, Luke Zettlemoyer, Noah A. Smith, and Tao Yu. 2023. Binding Language Models in Symbolic Languages. In *The Eleventh International Conference on Learning Representations*.
- [38] Jonathan H Choi, Kristin E Hickman, Amy Monahan, and Daniel Schwarcz. 2023. ChatGPT goes to law school. *Available at SSRN* (2023).
- [39] Aakanksha Chowdhery, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, Hyung Won Chung, Charles Sutton, Sebastian Gehrmann, et al. 2022. Palm: Scaling language modeling with pathways. *arXiv preprint arXiv:2204.02311* (2022).
- [40] Spandan Ghose Chowdhury, Soham Routh, and Satyajit Chakrabarti. 2014. News analytics and sentiment analysis to predict stock price trends. *International Journal of Computer Science and Information Technologies* 5, 3 (2014), 3595–3604.
- [41] Paul F Christiano, Jan Leike, Tom Brown, Miljan Martic, Shane Legg, and Dario Amodei. 2017. Deep reinforcement learning from human preferences. *Advances in neural information processing systems* 30 (2017).
- [42] Alexandra Chronopoulou, Matthew E Peters, Alexander Fraser, and Jesse Dodge. 2023. AdapterSoup: Weight Averaging to Improve Generalization of Pretrained Language Models. *arXiv preprint arXiv:2302.07027* (2023).
- [43] Hyung Won Chung, Le Hou, Shayne Longpre, Barret Zoph, Yi Tay, William Fedus, Eric Li, Xuezhi Wang, Mostafa Dehghani, Siddhartha Brahma, et al. 2022. Scaling instruction-finetuned language models. *arXiv preprint arXiv:2210.11416* (2022).
- [44] Debby RE Cotton, Peter A Cotton, and J Reuben Shipway. 2023. Chatting and cheating: Ensuring academic integrity in the era of ChatGPT. *Innovations in Education and Teaching International* (2023), 1–12.
- [45] Damai Dai, Li Dong, Yaru Hao, Zhifang Sui, Baobao Chang, and Furu Wei. 2021. Knowledge neurons in pretrained transformers. *arXiv preprint arXiv:2104.08696* (2021).
- [46] Zhuyun Dai, Vincent Y Zhao, Ji Ma, Yi Luan, Jianmo Ni, Jing Lu, Anton Bakalov, Kelvin Guu, Keith B Hall, and Ming-Wei Chang. 2022. Promptagator: Few-shot dense retrieval from 8 examples. *arXiv preprint arXiv:2209.11755* (2022).
- [47] Hugo Dalla-Torre, Liam Gonzalez, Javier Mendoza-Revilla, Nicolas Lopez Carranza, Adam Henryk Grzywaczewski, Francesco Oteri, Christian Dallago, Evan Trop, Hassan Sirelkhatim, Guillaume Richard, et al. 2023. The Nucleotide Transformer: Building and Evaluating Robust Foundation Models for Human Genomics. *bioRxiv* (2023), 2023–01.
- [48] Ishita Dasgupta, Christine Kaeser-Chen, Kenneth Marino, Arun Ahuja, Sheila Babayan, Felix Hill, and Rob Fergus. 2023. Collaborating with language models for embodied reasoning. In *Second Workshop on Language and Reinforcement Learning*.
- [49] Nicola De Cao, Wilker Aziz, and Ivan Titov. 2021. Editing Factual Knowledge in Language Models. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*.
- [50] Nassim Dehouche. 2021. Plagiarism in the age of massive Generative Pre-trained Transformers (GPT-3). *Ethics in Science and Environmental Politics* 21 (2021), 17–23.
- [51] Ning Ding, Yujia Qin, Guang Yang, Fuchao Wei, Zonghan Yang, Yusheng Su, Shengding Hu, Yulin Chen, Chi-Min Chan, Weize Chen, et al. 2022. Delta tuning: A comprehensive study of parameter efficient methods for pre-trained language models. *arXiv* (2022).
- [52] Yihong Dong, Xue Jiang, Zhi Jin, and Ge Li. 2023. Self-collaboration Code Generation via ChatGPT. *arXiv preprint arXiv:2304.07590* (2023).
- [53] Andrew Drozdov, Nathanael Schärli, Ekin Akyürek, Nathan Scales, Xinying Song, Xinyun Chen, Olivier Bousquet, and Denny Zhou. 2023. Compositional Semantic Parsing with Large Language Models. In *The Eleventh International Conference on Learning Representations*.
- [54] Dheeru Dua, Shivanshu Gupta, Sameer Singh, and Matt Gardner. 2022. Successive Prompting for Decomposing Complex Questions. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*.
- [55] Ali Edalati, Marzieh Tahaei, Ivan Kobyzev, Vahid Partovi Nia, James J Clark, and Mehdi Rezagholizadeh. 2022. KronA: Parameter Efficient Tuning with Kronecker Adapter. *arXiv preprint arXiv:2212.10650* (2022).

- [56] Mohammad Fahes, Tuan-Hung Vu, Andrei Bursuc, Patrick Pérez, and Raoul de Charette. 2022. P{\O}DA: Prompt-driven Zero-shot Domain Adaptation. *arXiv preprint arXiv:2212.03241* (2022).
- [57] Sarah Fakhoury, Saikat Chakraborty, Madan Musuvathi, and Shuvendu K Lahiri. 2023. Towards Generating Functionally Correct Code Edits from Natural Language Issue Descriptions. *arXiv preprint arXiv:2304.03816* (2023).
- [58] Yang Feng, Shiyue Zhang, Andi Zhang, Dong Wang, and Andrew Abel. 2017. Memory-augmented Neural Machine Translation. In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*.
- [59] Noelia Ferruz, Steffen Schmidt, and Birte Höcker. 2022. ProtGPT2 is a deep unsupervised language model for protein design. *Nature communications* 13, 1 (2022), 4348.
- [60] Nino Fijačko, Lucija Gosak, Gregor Štiglic, Christopher T Picard, and Matthew John Douma. 2023. Can ChatGPT pass the life support exams without entering the American heart association course? *Resuscitation* 185 (2023).
- [61] Mehmet Firat. 2023. How chat GPT can transform autodidactic experiences and open education. *Department of Distance Education, Open Education Faculty, Anadolu Unive* (2023).
- [62] Ingrid E Fisher, Margaret R Garnsey, and Mark E Hughes. 2016. Natural language processing in accounting, auditing and finance: A synthesis of the literature with a roadmap for future research. *Intelligent Systems in Accounting, Finance and Management* 23, 3 (2016), 157–214.
- [63] Simon Frieder, Luca Pinchetti, Ryan-Rhys Griffiths, Tommaso Salvatori, Thomas Lukasiewicz, Philipp Christian Petersen, Alexis Chevalier, and Julius Berner. 2023. Mathematical capabilities of chatgpt. *arXiv preprint arXiv:2301.13867* (2023).
- [64] Luyu Gao, Aman Madaan, Shuyan Zhou, Uri Alon, Pengfei Liu, Yiming Yang, Jamie Callan, and Graham Neubig. 2022. PAL: Program-aided Language Models. *arXiv preprint arXiv:2211.10435* (2022).
- [65] Chunjiang Ge, Rui Huang, Mixue Xie, Zihang Lai, Shiji Song, Shuang Li, and Gao Huang. 2022. Domain Adaptation via Prompt Learning. *arXiv preprint arXiv:2202.06687* (2022).
- [66] Aidan Gilson, Conrad W Safranek, Thomas Huang, Vimig Socrates, Ling Chi, Richard Andrew Taylor, David Chartash, et al. 2023. How does CHATGPT perform on the United States Medical Licensing Examination? the implications of large language models for medical education and knowledge assessment. *JMIR Medical Education* 9, 1 (2023), e45312.
- [67] Edouard Grave, Armand Joulin, and Nicolas Usunier. 2017. Improving neural language models with a continuous cache. In *International Conference on Learning Representations*.
- [68] Yuxian Gu, Xu Han, Zhiyuan Liu, and Minlie Huang. 2022. PPT: Pre-trained Prompt Tuning for Few-shot Learning. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, Dublin, Ireland, 8410–8423.
- [69] Xu Guo, Boyang Li, and Han Yu. 2022. Improving the Sample Efficiency of Prompt Tuning with Domain Adaptation. In *Findings of the Association for Computational Linguistics: EMNLP 2022*. Association for Computational Linguistics, Abu Dhabi, United Arab Emirates.
- [70] Xu Guo and Han Yu. 2022. On the Domain Adaptation and Generalization of Pretrained Language Models: A Survey. *arXiv* (2022).
- [71] Puneet Gupta, Swati Raturi, and P Venkateswarlu. 2023. Chatgpt for designing course outlines: A boon or bane to modern technology. *Available at SSRN 4386113* (2023).
- [72] Karen Hambardzumyan, Hrant Khachatryan, and Jonathan May. 2021. WARP: Word-level Adversarial ReProgramming. In *Annual Meeting of the Association for Computational Linguistics*.
- [73] Hangfeng He, Hongming Zhang, and Dan Roth. 2022. Rethinking with Retrieval: Faithful Large Language Model Inference. *arXiv preprint arXiv:2301.00303* (2022).
- [74] Junxian He, Chunting Zhou, Xuezhe Ma, Taylor Berg-Kirkpatrick, and Graham Neubig. 2021. Towards a unified view of parameter-efficient transfer learning. *arXiv preprint arXiv:2110.04366* (2021).
- [75] Shwai He, Liang Ding, Daize Dong, Miao Zhang, and Dacheng Tao. 2022. Sparseadapter: An easy approach for improving the parameter-efficiency of adapters. *arXiv preprint arXiv:2210.04284* (2022).
- [76] Dan Hendrycks and Kevin Gimpel. 2016. Gaussian error linear units (gelus). *arXiv preprint arXiv:1606.08415* (2016).
- [77] Evan Hernandez, Belinda Z Li, and Jacob Andreas. 2023. Measuring and Manipulating Knowledge Representations in Language Models. *arXiv preprint arXiv:2304.00740* (2023).
- [78] Urfa Khairatun Hisan and Muhammad Miftahul Amri. 2023. ChatGPT and Medical Education: A Double-Edged Sword. *Journal of Pedagogy and Education Science* 2, 01 (2023), 71–89.
- [79] Neil Houlsby, Andrei Giurgiu, Stanislaw Jastrzebski, Bruna Morrone, Quentin De Laroussilhe, Andrea Gesmundo, Mona Attariyan, and Sylvain Gelly. 2019. Parameter-efficient transfer learning for NLP. In *International Conference on Machine Learning*. PMLR, 2790–2799.
- [80] Jeremy Howard and Sebastian Ruder. 2018. Universal language model fine-tuning for text classification. *arXiv preprint arXiv:1801.06146* (2018).
- [81] Cheng-Yu Hsieh, Chun-Liang Li, Chih-Kuan Yeh, Hootan Nakhost, Yasuhisa Fujii, Alexander Ratner, Ranjay Krishna, Chen-Yu Lee, and Tomas Pfister. 2023. Distilling Step-by-Step! Outperforming Larger Language Models with Less Training Data and Smaller Model Sizes. *arXiv preprint arXiv:2305.02301* (2023).
- [82] Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2021. Lora: Low-rank adaptation of large language models. *arXiv preprint arXiv:2106.09685* (2021).
- [83] Shengding Hu, Ning Ding, Huadong Wang, Zhiyuan Liu, Jingang Wang, Juanzi Li, Wei Wu, and Maosong Sun. 2022. Knowledgeable Prompt-tuning: Incorporating Knowledge into Prompt Verbalizer for Text Classification. In *Proceedings of the 60th Annual Meeting of the Association for*

- Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, Dublin, Ireland.
- [84] Zhiqiang Hu, Yihuai Lan, Lei Wang, Wanyu Xu, Ee-Peng Lim, Roy Ka-Wei Lee, Lidong Bing, and Soujanya Poria. 2023. LLM-Adapters: An Adapter Family for Parameter-Efficient Fine-Tuning of Large Language Models. *arXiv preprint arXiv:2304.01933* (2023).
 - [85] Jiaxin Huang, Shixiang Shane Gu, Le Hou, Yuxin Wu, Xuezhi Wang, Hongkun Yu, and Jiawei Han. 2022. Large language models can self-improve. *arXiv preprint arXiv:2210.11610* (2022).
 - [86] Sun Huh. 2023. Are ChatGPT's knowledge and interpretation ability comparable to those of medical students in Korea for taking a parasitology examination?: a descriptive study. *Journal of Educational Evaluation for Health Professions* 20 (2023), 1.
 - [87] Iyolita Islam and Muhammad Nazrul Islam. 2023. Opportunities and Challenges of ChatGPT in Academia: A Conceptual Analysis. *Authorea Preprints* (2023).
 - [88] Kwan Yuen Iu and Vanessa Man-Yi Wong. 2023. ChatGPT by OpenAI: The End of Litigation Lawyers? *Available at SSRN* (2023).
 - [89] Gautier Izacard, Patrick Lewis, Maria Lomeli, Lucas Hosseini, Fabio Petroni, Timo Schick, Jane Dwivedi-Yu, Armand Joulin, Sebastian Riedel, and Edouard Grave. 2022. Few-shot learning with retrieval augmented language models. *arXiv preprint arXiv:2208.03299* (2022).
 - [90] Naman Jain, Skanda Vaidyanath, Arun Iyer, Nagarajan Natarajan, Suresh Parthasarathy, Sriram Rajamani, and Rahul Sharma. 2022. Jigsaw: Large language models meet program synthesis. In *Proceedings of the 44th International Conference on Software Engineering*. 1219–1231.
 - [91] Sajed Jalil, Suzzana Rafi, Thomas D LaToza, Kevin Moran, and Wing Lam. 2023. Chatgpt and software testing education: Promises & perils. *arXiv preprint arXiv:2302.03287* (2023).
 - [92] Katharina Jeblick, Balthasar Schachtner, Jakob Dextl, Andreas Mittermeier, Anna Theresa Stüber, Johanna Topalis, Tobias Weber, Philipp Wesp, Bastian Sabel, Jens Ricke, et al. 2022. ChatGPT Makes Medicine Easy to Swallow: An Exploratory Case Study on Simplified Radiology Reports. *arXiv preprint arXiv:2212.14882* (2022).
 - [93] Ziwei Ji, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Ye Jin Bang, Andrea Madotto, and Pascale Fung. 2023. Survey of hallucination in natural language generation. *Comput. Surveys* 55, 12 (2023), 1–38.
 - [94] Chen Jia and Yue Zhang. 2022. Prompt-based Distribution Alignment for Domain Generalization in Text Classification. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*. 10147–10157.
 - [95] Albert Qiaochu Jiang, Sean Welleck, Jin Peng Zhou, Timothee Lacroix, Jiacheng Liu, Wenda Li, Mateja Jamnik, Guillaume Lample, and Yuhuai Wu. 2023. Draft, Sketch, and Prove: Guiding Formal Theorem Provers with Informal Proofs. In *The Eleventh International Conference on Learning Representations*.
 - [96] Haoming Jiang, Pengcheng He, Weizhu Chen, Xiaodong Liu, Jianfeng Gao, and Tuo Zhao. 2019. Smart: Robust and efficient fine-tuning for pre-trained natural language models through principled regularized optimization. *arXiv preprint arXiv:1911.03437* (2019).
 - [97] Feihu Jin, Jinliang Lu, Jiajun Zhang, and Chengqing Zong. 2022. Instance-aware prompt learning for language understanding and generation. *arXiv preprint arXiv:2201.07126* (2022).
 - [98] Qiao Jin, Yifan Yang, Qingyu Chen, and Zhiyong Lu. [n. d.]. GeneGPT: Augmenting Large Language Models with Domain Tools for Improved Access to Biomedical Information. *ArXiv* ([n. d.]).
 - [99] Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. 2020. Scaling laws for neural language models. *arXiv* (2020).
 - [100] Rabeeh Karimi Mahabadi, James Henderson, and Sebastian Ruder. 2021. Compacter: Efficient low-rank hypercomplex adapter layers. *Advances in Neural Information Processing Systems* 34 (2021), 1022–1035.
 - [101] Enkelejda Kasneci, Kathrin Seßler, Stefan Küchemann, Maria Bannert, Daryna Dementieva, Frank Fischer, Urs Gasser, Georg Groh, Stephan Günemann, Eyke Hüllermeier, et al. 2023. ChatGPT for good? On opportunities and challenges of large language models for education. *Learning and Individual Differences* 103 (2023), 102274.
 - [102] Daniel Martin Katz, Michael James Bommarito, Shang Gao, and Pablo Arredondo. 2023. Gpt-4 passes the bar exam. *Available at SSRN* 4389233 (2023).
 - [103] Mohammad Khalil and Erkan Er. 2023. Will ChatGPT get you caught? Rethinking of plagiarism detection. *arXiv preprint arXiv:2302.04335* (2023).
 - [104] Junaed Younus Khan and Gias Uddin. 2023. Combining Contexts from Multiple Sources for Documentation-Specific Code Example Generation. *arXiv preprint arXiv:2303.14542* (2023).
 - [105] Rehan Ahmed Khan, Masood Jawaid, Aymen Rehan Khan, and Madiha Sajjad. 2023. ChatGPT-Reshaping medical education and clinical management. *Pakistan Journal of Medical Sciences* 39, 2 (2023).
 - [106] Tushar Khot, Harsh Trivedi, Matthew Finlayson, Yao Fu, Kyle Richardson, Peter Clark, and Ashish Sabharwal. 2022. Decomposed prompting: A modular approach for solving complex tasks. *arXiv preprint arXiv:2210.02406* (2022).
 - [107] Michael R King and ChatGPT. 2023. A conversation on artificial intelligence, chatbots, and plagiarism in higher education. *Cellular and Molecular Bioengineering* 16, 1 (2023), 1–2.
 - [108] Takeshi Kojima, Shixiang Shane Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa. 2022. Large Language Models are Zero-Shot Reasoners. In *ICML 2022 Workshop on Knowledge Retrieval and Language Models*.
 - [109] Mojtaba Komeili, Kurt Shuster, and Jason Weston. 2022. Internet-Augmented Dialogue Generation. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 8460–8478.
 - [110] Gerd Kortemeyer. 2023. Could an Artificial-Intelligence agent pass an introductory physics course? *arXiv preprint arXiv:2301.12127* (2023).

- [111] Tiffany H Kung, Morgan Cheatham, Arielle Medenilla, Czarina Sillos, Lorie De Leon, Camille Elepaño, Maria Madriaga, Rimel Aggabao, Giezel Diaz-Candido, James Maningo, et al. 2023. Performance of ChatGPT on USMLE: Potential for AI-assisted medical education using large language models. *PLoS digital health* 2, 2 (2023), e0000198.
- [112] Peter Lee, Sebastien Bubeck, and Joseph Petro. 2023. Benefits, Limits, and Risks of GPT-4 as an AI Chatbot for Medicine. *New England Journal of Medicine* 388, 13 (2023), 1233–1239.
- [113] Eric Lehman, Evan Hernandez, Diwakar Mahajan, Jonas Wulff, Micah J Smith, Zachary Ziegler, Daniel Nadler, Peter Szolovits, Alistair Johnson, and Emily Alsentzer. 2023. Do We Still Need Clinical Language Models? *arXiv preprint arXiv:2302.08091* (2023).
- [114] Markus Leippold. 2023. Sentiment Spin: Attacking Financial Sentiment with GPT-3. *Available at SSRN* (2023).
- [115] Brian Lester, Rami Al-Rfou, and Noah Constant. 2021. The Power of Scale for Parameter-Efficient Prompt Tuning. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, Online and Punta Cana, Dominican Republic, 3045–3059. <https://doi.org/10.18653/v1/2021.emnlp-main.243>
- [116] Yoav Levine, Itay Dalmedigos, Ori Ram, Yoel Zeldes, Daniel Jannai, Dor Muhlgay, Yoni Osin, Opher Lieber, Barak Lenz, Shai Shalev-Shwartz, Amnon Shashua, Kevin Leyton-Brown, and Yoav Shoham. 2022. Standing on the Shoulders of Giant Frozen Language Models. *ArXiv abs/2204.10019* (2022).
- [117] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, et al. 2020. Retrieval-augmented generation for knowledge-intensive nlp tasks. *Advances in Neural Information Processing Systems* 33 (2020), 9459–9474.
- [118] Daliang Li, Ankit Singh Rawat, Manzil Zaheer, Xin Wang, Michal Lukasik, Andreas Veit, Felix Yu, and Sanjiv Kumar. 2022. Large Language Models with Controllable Working Memory. *arXiv preprint arXiv:2211.05110* (2022).
- [119] Guohao Li, Hasan Abed Al Kader Hammoud, Hani Itani, Dmitrii Khizbullin, and Bernard Ghanem. 2023. CAMEL: Communicative Agents for" Mind" Exploration of Large Scale Language Model Society. *arXiv preprint arXiv:2303.17760* (2023).
- [120] Haochen Li, Tong Mo, Hongcheng Fan, Jingkun Wang, Jiaxi Wang, Fuhao Zhang, and Weiping Li. 2022. KiPT: Knowledge-injected Prompt Tuning for Event Detection. In *International Conference on Computational Linguistics*.
- [121] Jinyang Li, Binyuan Hui, Ge Qu, Binhua Li, Jiaxi Yang, Bowen Li, Bailin Wang, Bowen Qin, Rongyu Cao, Ruiying Geng, et al. 2023. Can LLM Already Serve as A Database Interface? A Big Bench for Large-Scale Database Grounded Text-to-SQLs. *arXiv preprint arXiv:2305.03111* (2023).
- [122] Xiang Lisa Li and Percy Liang. 2021. Prefix-Tuning: Optimizing Continuous Prompts for Generation. *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)* abs/2101.00190 (2021).
- [123] Yujia Li, David Choi, Junyoung Chung, Nate Kushman, Julian Schrittwieser, Rémi Leblond, Tom Eccles, James Keeling, Felix Gimeno, Agustin Dal Lago, et al. 2022. Competition-level code generation with alphacode. *Science* 378, 6624 (2022), 1092–1097.
- [124] Zongjie Li, Chaozheng Wang, Zhibo Liu, Haoxuan Wang, Shuai Wang, and Cuiyun Gao. 2022. CCTEST: Testing and Repairing Code Completion Systems. *arXiv preprint arXiv:2208.08289* (2022).
- [125] Jacky Liang, Wenlong Huang, Fei Xia, Peng Xu, Karol Hausman, Pete Florence, Andy Zeng, et al. 2022. Code as Policies: Language Model Programs for Embodied Control. In *Workshop on Language and Robotics at CoRL 2022*.
- [126] Yaobo Liang, Chenfei Wu, Ting Song, Wenshan Wu, Yan Xia, Yu Liu, Yang Ou, Shuai Lu, Lei Ji, Shaoguang Mao, et al. 2023. TaskMatrix. AI: Completing Tasks by Connecting Foundation Models with Millions of APIs. *arXiv preprint arXiv:2303.16434* (2023).
- [127] Valentin Liévin, Christoffer Egeberg Hother, and Ole Winther. 2022. Can large language models reason about medical questions? *arXiv preprint arXiv:2207.08143* (2022).
- [128] Hongzhan Lin, Pengyao Yi, Jing Ma, Haiyun Jiang, Ziyang Luo, Shuming Shi, and Ruifang Liu. 2022. Zero-Shot Rumor Detection with Propagation Structure via Prompt Learning. *arXiv preprint arXiv:2212.01117* (2022).
- [129] Qi Liu, Dani Yogatama, and Phil Blunsom. 2022. Relational Memory-Augmented Language Models. *Transactions of the Association for Computational Linguistics* 10 (2022), 555–572.
- [130] Ruibo Liu, Jason Wei, Shixiang Shane Gu, Te-Yen Wu, Soroush Vosoughi, Claire Cui, Denny Zhou, and Andrew M Dai. 2022. Mind's Eye: Grounded Language Model Reasoning through Simulation. *arXiv preprint arXiv:2210.05359* (2022).
- [131] Xiangyang Liu, Tianxiang Sun, Xuanjing Huang, and Xipeng Qiu. 2022. Late Prompt Tuning: A Late Prompt Could Be Better Than Many Prompts. In *Findings of the Association for Computational Linguistics: EMNLP 2022*. Association for Computational Linguistics, Abu Dhabi, United Arab Emirates.
- [132] Alejandro Lopez-Lira and Yuehua Tang. 2023. Can ChatGPT Forecast Stock Price Movements? Return Predictability and Large Language Models. *arXiv preprint arXiv:2304.07619* (2023).
- [133] Jiaying Lu, Jiaming Shen, Bo Xiong, Wengjing Ma, Staab Steffen, and Carl Yang. 2023. HiPrompt: Few-Shot Biomedical Knowledge Fusion via Hierarchy-Oriented Prompting. In *46th International ACM SIGIR Conference on Research and Development in Information Retrieval - Short Paper*.
- [134] Yao Lu, Max Bartolo, Alastair Moore, Sebastian Riedel, and Pontus Stenetorp. 2022. Fantastically Ordered Prompts and Where to Find Them: Overcoming Few-Shot Prompt Order Sensitivity. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 8086–8098.
- [135] Brady D Lund and Ting Wang. 2023. Chatting about ChatGPT: how may AI and GPT impact academia and libraries? *Library Hi Tech News* (2023).

- [136] Renqian Luo, Liai Sun, Yingce Xia, Tao Qin, Sheng Zhang, Hoifung Poon, and Tie-Yan Liu. 2022. BioGPT: generative pre-trained transformer for biomedical text generation and mining. *Briefings in Bioinformatics* 23, 6 (2022).
- [137] Muneer M Alshater. 2022. Exploring the role of artificial intelligence in enhancing academic performance: A case study of ChatGPT. *Available at SSRN* (2022).
- [138] Rupert Macey-Dare. 2023. ChatGPT & Generative AI Systems as Quasi-Expert Legal Advice Lawyers-Case Study Considering Potential Appeal Against Conviction of Tom Hayes. *Available at SSRN 4342686* (2023).
- [139] Rupert Macey-Dare. 2023. How ChatGPT and Generative AI Systems will Revolutionize Legal Services and the Legal Profession. *Available at SSRN* (2023).
- [140] Aman Madaan, Shuyan Zhou, Uri Alon, Yiming Yang, and Graham Neubig. 2022. Language models of code are few-shot commonsense learners. *arXiv preprint arXiv:2210.07128* (2022).
- [141] Ali Madani, Ben Krause, Eric R Greene, Subu Subramanian, Benjamin P Mohr, James M Holton, Jose Luis Olmos Jr, Caiming Xiong, Zachary Z Sun, Richard Socher, et al. 2023. Large language models generate functional protein sequences across diverse families. *Nature Biotechnology* (2023), 1–8.
- [142] Babak Mahjour, Jillian Hoffstadt, and Tim Cernak. 2023. Designing Chemical Reaction Arrays using phactor and ChatGPT. (2023).
- [143] Gengchen Mai, Chris Cundy, Kristy Choi, Yingjie Hu, Ni Lao, and Stefano Ermon. 2022. Towards a foundation model for geospatial artificial intelligence (vision paper). In *Proceedings of the 30th International Conference on Advances in Geographic Information Systems*. 1–4.
- [144] Gengchen Mai, Weiming Huang, Jin Sun, Suhang Song, Deepak Mishra, Ninghao Liu, Song Gao, Tianming Liu, Gao Cong, Yingjie Hu, et al. 2023. On the opportunities and challenges of foundation models for geospatial artificial intelligence. *arXiv preprint arXiv:2304.06798* (2023).
- [145] Bhavitvya Malik, Abhinav Ramesh Kashyap, Min-Yen Kan, and Soujanya Poria. 2023. UDApter—Efficient Domain Adaptation Using Adapters. *arXiv preprint arXiv:2302.03194* (2023).
- [146] Kamil Malinka, Martin Perešini, Anton Firc, Ondřej Hujňák, and Filip Januš. 2023. On the educational impact of chatgpt: Is artificial intelligence ready to obtain a university degree? *arXiv preprint arXiv:2303.11146* (2023).
- [147] Potsawee Manakul, Adian Liusie, and Mark JF Gales. 2023. Selfcheckgpt: Zero-resource black-box hallucination detection for generative large language models. *arXiv preprint arXiv:2303.08896* (2023).
- [148] Yuning Mao, Lambert Mathias, Rui Hou, Amjad Almahairi, Hao Ma, Jiawei Han, Wen-tau Yih, and Madian Khabsa. 2021. Unipelt: A unified framework for parameter-efficient language model tuning. *arXiv preprint arXiv:2110.07577* (2021).
- [149] Raja Marjeh, Ilia Sucholutsky, Pol van Rijn, Nori Jacoby, and Thomas L Griffiths. 2023. What Language Reveals about Perception: Distilling Psychophysical Knowledge from Large Language Models. *arXiv preprint arXiv:2302.01308* (2023).
- [150] Amarachi B Mbakwe, Ismini Lourentzou, Leo Anthony Celi, Oren J Mechanic, and Alon Dagan. 2023. ChatGPT passing USMLE shines a spotlight on the flaws of medical education. , e0000205 pages.
- [151] Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. 2022. Locating and editing factual associations in gpt. *Advances in Neural Information Processing Systems* 35 (2022), 17359–17372.
- [152] Kevin Meng, Arnab Sen Sharma, Alex Andonian, Yonatan Belinkov, and David Bau. 2022. Mass-editing memory in a transformer. *arXiv preprint arXiv:2210.07229* (2022).
- [153] Jacob Menick, Maja Trebacz, Vladimir Mikulik, John Aslanides, Francis Song, Martin Chadwick, Mia Glaese, Susannah Young, Lucy Campbell-Gillingham, Geoffrey Irving, et al. 2022. Teaching language models to support answers with verified quotes. *arXiv preprint arXiv:2203.11147* (2022).
- [154] Jacob Menick, Maja Trebacz, Vladimir Mikulik, John Aslanides, Francis Song, Martin Chadwick, Mia Glaese, Susannah Young, Lucy Campbell-Gillingham, Geoffrey Irving, et al. 2022. Teaching language models to support answers with verified quotes. *arXiv preprint arXiv:2203.11147* (2022).
- [155] Stephen Merity, Caiming Xiong, James Bradbury, and Richard Socher. 2016. Pointer sentinel mixture models. In *International Conference on Learning Representations*.
- [156] Grégoire Mialon, Roberto Dessi, Maria Lomeli, Christoforos Nalmpantis, Ram Pasunuru, Roberta Raileanu, Baptiste Rozière, Timo Schick, Jane Dwivedi-Yu, Asli Celikyilmaz, et al. 2023. Augmented language models: a survey. *arXiv preprint arXiv:2302.07842* (2023).
- [157] Bonan Min, Hayley Ross, Elior Sulem, Amir Pouran Ben Veyseh, Thien Huu Nguyen, Oscar Sainz, Eneko Agirre, Ilana Heinz, and Dan Roth. 2021. Recent advances in natural language processing via large pre-trained language models: A survey. *arXiv* (2021).
- [158] Sewon Min, Xinxin Lyu, Ari Holtzman, Mikel Artetxe, Mike Lewis, Hannaneh Hajishirzi, and Luke Zettlemoyer. 2022. Rethinking the Role of Demonstrations: What Makes In-Context Learning Work? *arXiv preprint arXiv:2202.12837* (2022).
- [159] Eric Mitchell, Charles Lin, Antoine Bosselut, Chelsea Finn, and Christopher D Manning. 2021. Fast model editing at scale. *arXiv preprint arXiv:2110.11309* (2021).
- [160] Eric Mitchell, Charles Lin, Antoine Bosselut, Christopher D Manning, and Chelsea Finn. 2022. Memory-based model editing at scale. In *International Conference on Machine Learning*. 15817–15831.
- [161] Philip Moons and Liesbet Van Bulck. 2023. ChatGPT: Can artificial intelligence language models be of value for cardiovascular nurses and allied health professionals. *European Journal of Cardiovascular Nursing* (2023), zvad022–zvad022.
- [162] Steven Moore, Huy A Nguyen, Norman Bier, Tanvi Domadia, and John Stamper. 2022. Assessing the Quality of Student-Generated Short Answer Questions Using GPT-3. In *Educating for a New Future: Making Sense of Technology-Enhanced Learning Adoption: 17th European Conference on Technology Enhanced Learning, EC-TEL 2022, Toulouse, France, September 12–16, 2022, Proceedings*. Springer, 243–257.

- [163] Aashiq Muhamed, Iman Keivanloo, Sujan Perera, James Mracek, Yi Xu, Qingjun Cui, Santosh Rajagopalan, Belinda Zeng, and Trishul Chilimbi. 2021. CTR-BERT: Cost-effective knowledge distillation for billion-parameter teacher models. In *NeurIPS Efficient Natural Language and Speech Processing Workshop*.
- [164] Prasanth Murali, Ian Steenstra, Hye Sun Yun, Ameneh Shamekhi, and Timothy Bickmore. 2023. Improving Multiparty Interactions with a Robot Using Large Language Models. In *Extended Abstracts of the 2023 CHI Conference on Human Factors in Computing Systems*. 1–8.
- [165] Reiichiro Nakano, Jacob Hilton, Suchir Balaji, Jeff Wu, Long Ouyang, Christina Kim, Christopher Hesse, Shantanu Jain, Vineet Kosaraju, William Saunders, et al. 2021. Webgpt: Browser-assisted question-answering with human feedback. *arXiv preprint arXiv:2112.09332* (2021).
- [166] Aishwarya Narasimhan, Krishna Prasad Agara Venkatesha Rao, et al. 2021. CGEMs: A metric model for automatic code generation using GPT-3. *arXiv preprint arXiv:2108.10168* (2021).
- [167] Anthony J Nastasi, Katherine R Courtright, Scott D Halpern, and Gary E Weissman. 2023. Does ChatGPT Provide Appropriate and Equitable Medical Advice?: A Vignette-Based, Clinical Evaluation Across Care Contexts. *medRxiv* (2023), 2023–02.
- [168] Siddharth Nath, Abdullah Marie, Simon Ellershaw, Edward Korot, and Pearse A Keane. 2022. New meaning for NLP: the trials and tribulations of natural language processing with GPT-3 in ophthalmology. *British Journal of Ophthalmology* 106, 7 (2022), 889–892.
- [169] Saima Nisar and Muhammad Shahzad Aslam. 2023. Is ChatGPT a Good Tool for T&CM Students in Studying Pharmacology? *Available at SSRN 4324310* (2023).
- [170] Paweł Niszczoła and Sami Abbas. [n. d.]. GPT as a Financial Advisor. *Available at SSRN 4384861* ([n. d.]).
- [171] Siobhan O'Connor et al. 2022. Open artificial intelligence platforms in nursing education: Tools for academic progress or abuse? *Nurse Education in Practice* 66 (2022), 103537–103537.
- [172] OpenAI. [n. d.]. ChatGPT plugins. <https://openai.com/blog/chatgpt-plugins>. Accessed: 2023-04-05.
- [173] OpenAI. 2023. GPT-4 Technical Report. *arXiv:2303.08774* [cs.CL]
- [174] Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. 2022. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems* 35 (2022), 27730–27744.
- [175] Zachary A Pardos and Shreya Bhandari. 2023. Learning gain differences between ChatGPT and human tutor generated algebra hints. *arXiv preprint arXiv:2302.06871* (2023).
- [176] Joon Sung Park, Joseph C O'Brien, Carrie J Cai, Meredith Ringel Morris, Percy Liang, and Michael S Bernstein. 2023. Generative Agents: Interactive Simulacra of Human Behavior. *arXiv preprint arXiv:2304.03442* (2023).
- [177] John V Pavlik. 2023. Collaborating With ChatGPT: Considering the Implications of Generative Artificial Intelligence for Journalism and Media Education. *Journalism & Mass Communication Educator* (2023), 10776958221149577.
- [178] Hammond Pearce, Benjamin Tan, Baleegh Ahmad, Ramesh Karri, and Brendan Dolan-Gavitt. 2021. Can OpenAI Codex and Other Large Language Models Help Us Fix Security Bugs? *arXiv preprint arXiv:2112.02125* (2021).
- [179] Hammond Pearce, Benjamin Tan, Prashanth Krishnamurthy, Farshad Khorrami, Ramesh Karri, and Brendan Dolan-Gavitt. 2022. Pop Quiz! Can a Large Language Model Help With Reverse Engineering? *arXiv preprint arXiv:2202.01142* (2022).
- [180] Baolin Peng, Michel Galley, Pengcheng He, Hao Cheng, Yujia Xie, Yu Hu, Qiuyuan Huang, Lars Liden, Zhou Yu, Weizhu Chen, et al. 2023. Check your facts and try again: Improving large language models with external knowledge and automated feedback. *arXiv preprint arXiv:2302.12813* (2023).
- [181] Baolin Peng, Chunyuan Li, Pengcheng He, Michel Galley, and Jianfeng Gao. 2023. Instruction tuning with gpt-4. *arXiv preprint arXiv:2304.03277* (2023).
- [182] Jonas Pfeiffer, Aishwarya Kamath, Andreas Rücklé, Kyunghyun Cho, and Iryna Gurevych. 2020. AdapterFusion: Non-destructive task composition for transfer learning. *arXiv preprint arXiv:2005.00247* (2020).
- [183] Jonas Pfeiffer, Andreas Rücklé, Clifton Poth, Aishwarya Kamath, Ivan Vulić, Sebastian Ruder, Kyunghyun Cho, and Iryna Gurevych. 2020. Adapterhub: A framework for adapting transformers. *arXiv preprint arXiv:2007.07779* (2020).
- [184] Jonas Pfeiffer, Ivan Vulić, Iryna Gurevych, and Sebastian Ruder. 2020. Mad-x: An adapter-based framework for multi-task cross-lingual transfer. *arXiv preprint arXiv:2005.00052* (2020).
- [185] Tung Phung, José Cambronero, Sumit Gulwani, Tobias Kohn, Rupak Majumdar, Adish Singla, and Gustavo Soares. 2023. Generating High-Precision Feedback for Programming Syntax Errors using Large Language Models. *arXiv preprint arXiv:2302.04662* (2023).
- [186] Stephen R Piccolo, Paul Denny, Andrew Luxton-Reilly, Samuel Payne, and Perry G Ridge. 2023. Many bioinformatics programming tasks can be automated with ChatGPT. *arXiv preprint arXiv:2303.13528* (2023).
- [187] Edoardo M Ponti, Alessandro Sordani, and Siva Reddy. 2022. Combining modular skills in multitask learning. *arXiv preprint arXiv:2202.13914* (2022).
- [188] Mohammadreza Pourreza and Davood Rafiei. 2023. DIN-SQL: Decomposed In-Context Learning of Text-to-SQL with Self-Correction. *arXiv preprint arXiv:2304.11015* (2023).
- [189] Nishchal Prasad, Mohand Boughanem, and Taoufik Dkaki. 2022. Effect of Hierarchical Domain-specific Language Models and Attention in the Classification of Decisions for Legal Cases. In *Proceedings of the CIRCLE (Joint Conference of the Information Retrieval Communities in Europe), Samatan, Gers, France*. 4–7.
- [190] Junaid Qadir. 2022. Engineering Education in the Era of ChatGPT: Promise and Pitfalls of Generative AI for Education. (2022).

- [191] Chengwei Qin and Shafiq Joty. 2021. LFPT5: A unified framework for lifelong few-shot language learning based on prompt tuning of t5. *arXiv preprint arXiv:2110.07298* (2021).
- [192] Yujia Qin, Shengding Hu, Yankai Lin, Weize Chen, Ning Ding, Ganqu Cui, Zheni Zeng, Yufei Huang, Chaojun Xiao, Chi Han, et al. 2023. Tool Learning with Foundation Models. *arXiv preprint arXiv:2304.08354* (2023).
- [193] Xipeng Qiu, Tianxiang Sun, Yige Xu, Yunfan Shao, Ning Dai, and Xuanjing Huang. 2020. Pre-trained models for natural language processing: A survey. *Science China Technological Sciences* 63, 10 (2020), 1872–1897.
- [194] Jon TS Quah and YW Chua. 2019. Chatbot assisted marketing in financial service industry. In *Services Computing–SCC 2019: 16th International Conference, Held as Part of the Services Conference Federation, SCF 2019, San Diego, CA, USA, June 25–30, 2019, Proceedings 16*. Springer, 107–114.
- [195] Alec Radford, Karthik Narasimhan, Tim Salimans, Ilya Sutskever, et al. 2018. Improving language understanding by generative pre-training. (2018).
- [196] Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J Liu. 2020. Exploring the limits of transfer learning with a unified text-to-text transformer. *The Journal of Machine Learning Research* 21, 1 (2020), 5485–5551.
- [197] Ori Ram, Yoav Levine, Itay Dalmedigos, Dor Muhlgay, Amnon Shashua, Kevin Leyton-Brown, and Yoav Shoham. 2023. In-Context Retrieval-Augmented Language Models. *arXiv preprint arXiv:2302.00083* (2023).
- [198] Arya Rao, John Kim, Meghana Kamineni, Michael Pang, Winston Lie, and Marc D Succi. 2023. Evaluating ChatGPT as an adjunct for radiologic decision-making. *medRxiv* (2023), 2023–02.
- [199] Arya S Rao, Michael Pang, John Kim, Meghana Kamineni, Winston Lie, Anoop K Prasad, Adam Landman, Keith Dryer, and Marc D Succi. 2023. Assessing the Utility of ChatGPT Throughout the Entire Clinical Workflow. *medRxiv* (2023), 2023–02.
- [200] Laila Rasmy, Yang Xiang, Ziqian Xie, Cui Tao, and Degui Zhi. 2021. Med-BERT: pretrained contextualized embeddings on large-scale structured electronic health records for disease prediction. *NPJ digital medicine* 4, 1 (2021), 86.
- [201] Anastasia Razdaibiedina, Yuning Mao, Rui Hou, Madian Khabisa, Mike Lewis, and Amjad Almahairi. 2023. Progressive Prompts: Continual Learning for Language Models. In *The Eleventh International Conference on Learning Representations*.
- [202] Sylvestre-Alvise Rebuffi, Hakan Bilen, and Andrea Vedaldi. 2017. Learning multiple visual domains with residual adapters. *Advances in neural information processing systems* 30 (2017).
- [203] Protection Regulation. 2018. General data protection regulation. *Intouch* 25 (2018), 1–5.
- [204] Alexander Rives, Joshua Meier, Tom Sercu, Siddharth Goyal, Zeming Lin, Jason Liu, Demi Guo, Myle Ott, C Lawrence Zitnick, Jerry Ma, et al. 2021. Biological structure and function emerge from scaling unsupervised learning to 250 million protein sequences. *Proceedings of the National Academy of Sciences* 118, 15 (2021), e2016239118.
- [205] Adam Roberts, Colin Raffel, and Noam Shazeer. 2020. How much knowledge can you pack into the parameters of a language model? *arXiv preprint arXiv:2002.08910* (2020).
- [206] Joshua Robinson, Christopher Michael Rytting, and David Wingate. 2022. Leveraging Large Language Models for Multiple Choice Question Answering. *arXiv* (2022).
- [207] Jerret Ross, Brian Belgodere, Vijil Chenthamarakshan, Inkrit Padhi, Youssef Mroueh, and Payel Das. 2022. Large-scale chemical language representations capture molecular structure and properties. *Nature Machine Intelligence* 4, 12 (2022), 1256–1264.
- [208] Jürgen Rudolph, Samson Tan, and Shannon Tan. 2023. ChatGPT: Bullshit spewer or the end of traditional assessments in higher education? *Journal of Applied Learning and Teaching* 6, 1 (2023).
- [209] Malik Sallam. 2023. The utility of ChatGPT as an example of large language models in healthcare education, research and practice: Systematic review on the future perspectives and potential limitations. *medRxiv* (2023), 2023–02.
- [210] Malik Sallam, Nesreen A Salim, B Ala’a, Muna Barakat, Diaa Fayyad, Souheil Hallit, Harapan Harapan, Rabih Hallit, Azmi Mahafzah, and B Ala’a. 2023. ChatGPT Output Regarding Compulsory Vaccination and COVID-19 Vaccine Conspiracy: A Descriptive Study at the Outset of a Paradigm Shift in Online Search for Information. *Cureus* 15, 2 (2023).
- [211] Victor Sanh, Albert Webson, Colin Raffel, Stephen Bach, Lintang Sutawika, Zaid Alyafeai, Antoine Chaffin, Arnaud Stiegler, Arun Raja, Manan Dey, M Saiful Bari, Canwen Xu, Urmish Thakker, Shanya Sharma Sharma, Eliza Szczechla, Taewoon Kim, Gunjan Chhablani, Nihal Nayak, Debajyoti Datta, Jonathan Chang, Mike Tian-Jian Jiang, Han Wang, Matteo Manica, Sheng Shen, Zheng Xin Yong, Harshit Pandey, Rachel Bawden, Thomas Wang, Trishala Neeraj, Jos Rozen, Abheesh Sharma, Andrea Santilli, Thibault Fevry, Jason Alan Fries, Ryan Teehan, Teven Le Scao, Stella Biderman, Leo Gao, Thomas Wolf, and Alexander M Rush. 2022. Multitask Prompted Training Enables Zero-Shot Task Generalization. In *International Conference on Learning Representations*.
- [212] Sami Sarsa, Paul Denny, Arto Hellas, and Juho Leinonen. 2022. Automatic generation of programming exercises and code explanations using large language models. In *Proceedings of the 2022 ACM Conference on International Computing Education Research-Volume 1*. 27–43.
- [213] Timo Schick, Jane Dwivedi-Yu, Roberto Dessi, Roberta Raileanu, Maria Lomeli, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. 2023. Toolformer: Language models can teach themselves to use tools. *arXiv preprint arXiv:2302.04761* (2023).
- [214] Timo Schick and Hinrich Schütze. 2020. Exploiting cloze questions for few shot text classification and natural language inference. *arXiv preprint arXiv:2001.07676* (2020).
- [215] Dale Schuurmans. 2023. Memory Augmented Large Language Models are Computationally Universal. *arXiv preprint arXiv:2301.04589* (2023).
- [216] Thomas Scialom, Tuhin Chakrabarty, and Smaranda Muresan. 2022. Fine-tuned Language Models are Continual Learners. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*. 6107–6122.

- [217] Shohreh Shaghaghian, Luna Yue Feng, Borna Jafarpour, and Nicolai Pogrebnyakov. 2020. Customizing contextualized language models for legal document reviews. In *2020 IEEE International Conference on Big Data (Big Data)*. IEEE, 2139–2148.
- [218] Paulo Shakarian, Abhinav Koyyalamudi, Noel Ngu, and Lakshmihari Maredu. 2023. An Independent Evaluation of ChatGPT on Mathematical Word Problems (MWP). *arXiv preprint arXiv:2302.13814* (2023).
- [219] Yongliang Shen, Kaitao Song, Xu Tan, Dongsheng Li, Weiming Lu, and Yueting Zhuang. 2023. HuggingGPT: Solving AI Tasks with ChatGPT and its Friends in HuggingFace. *arXiv:2303.17580*
- [220] Hanul Shin, Jung Kwon Lee, Jaehong Kim, and Jiwon Kim. 2017. Continual learning with deep generative replay. *Advances in neural information processing systems* 30 (2017).
- [221] Richard Shin and Benjamin Van Durme. 2021. Few-shot semantic parsing with language models trained on code. *arXiv preprint arXiv:2112.08696* (2021).
- [222] Kumar Shridhar, Alessandro Stolfo, and Mrinmaya Sachan. 2022. Distilling multi-step reasoning capabilities of large language models into smaller models via semantic decompositions. *arXiv preprint arXiv:2212.00193* (2022).
- [223] Ridwan Islam Sifat. 2023. ChatGPT and the Future of Health Policy Analysis: Potential and Pitfalls of Using ChatGPT in Policymaking. *Annals of Biomedical Engineering* (2023), 1–3.
- [224] Devendra Singh, Siva Reddy, Will Hamilton, Chris Dyer, and Dani Yogatama. 2021. End-to-end training of multi-document reader and retriever for open-domain question answering. *Advances in Neural Information Processing Systems* 34 (2021), 25968–25981.
- [225] Ishika Singh, Valtis Blukis, Arsalan Mousavian, Ankit Goyal, Danfei Xu, Jonathan Tremblay, Dieter Fox, Jesse Thomason, and Animesh Garg. [n.d.]. ProgPrompt: Generating Situated Robot Task Plans using Large Language Models. In *Workshop on Language and Robotics at CoRL 2022*.
- [226] Karan Singhal, Shekoofeh Azizi, Tao Tu, S Sara Mahdavi, Jason Wei, Hyung Won Chung, Nathan Scales, Ajay Tanwani, Heather Cole-Lewis, Stephen Pfohl, et al. 2022. Large Language Models Encode Clinical Knowledge. *arXiv preprint arXiv:2212.13138* (2022).
- [227] Wataru Souma, Irena Vodenska, and Hideaki Aoyama. 2019. Enhanced news sentiment analysis using deep learning methods. *Journal of Computational Social Science* 2, 1 (2019), 33–46.
- [228] Yusheng Su, Xiaozhi Wang, Yujia Qin, Chi-Min Chan, Yankai Lin, Huadong Wang, Kaiyue Wen, Zhiyuan Liu, Peng Li, Juanzi Li, et al. 2022. On transferability of prompt tuning for natural language processing. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*. 3949–3969.
- [229] Zhongxiang Sun. 2023. A short survey of viewing large language models in legal aspect. *arXiv preprint arXiv:2303.09136* (2023).
- [230] Yi-Lin Sung, Jaemin Cho, and Mohit Bansal. 2022. Lst: Ladder side-tuning for parameter and memory efficient transfer learning. *arXiv preprint arXiv:2206.06522* (2022).
- [231] Didac Suris, Sachit Menon, and Carl Vondrick. 2023. Vipergpt: Visual inference via python execution for reasoning. *arXiv preprint arXiv:2303.08128* (2023).
- [232] Teo Susnjak. 2022. ChatGPT: The End of Online Exam Integrity? *arXiv preprint arXiv:2212.09292* (2022).
- [233] Anaïs Tack and Chris Piech. 2022. The AI Teacher test: Measuring the pedagogical ability of Blender and GPT-3 in educational dialogues. *arXiv preprint arXiv:2205.07540* (2022).
- [234] Ruixiang Tang, Xiaotian Han, Xiaoqian Jiang, and Xia Hu. 2023. Does synthetic data generation of llms help clinical text mining? *arXiv preprint arXiv:2303.04360* (2023).
- [235] Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, et al. 2023. Llama: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971* (2023).
- [236] Dietrich Trautmann, Alina Petrova, and Frank Schilder. 2022. Legal Prompt Engineering for Multilingual Legal Judgement Prediction. *arXiv preprint arXiv:2212.02199* (2022).
- [237] Immanuel Trummer. 2022. CodexDB: Generating Code for Processing SQL Queries using GPT-3 Codex. *arXiv preprint arXiv:2204.08941* (2022).
- [238] Mojtaba Valipour, Mehdi Rezagholizadeh, Ivan Kobyzev, and Ali Ghodsi. 2022. DyLoRA: Parameter Efficient Tuning of Pre-trained Models using Dynamic Search-Free Low-Rank Adaptation. *arXiv preprint arXiv:2210.07558* (2022).
- [239] Josef Valvoda, Ryan Cotterell, and Simone Teufel. 2023. On the Role of Negative Precedent in Legal Outcome Prediction. *Transactions of the Association for Computational Linguistics* 11 (01 2023), 34–48. https://doi.org/10.1162/tacl_a_00532 arXiv:https://direct.mit.edu/tacl/article-pdf/doi/10.1162/tacl_a_00532/2067838/tacl_a_00532.pdf
- [240] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Lukasz Kaiser, and Illia Polosukhin. 2017. Attention is all you need. *Advances in neural information processing systems* 30 (2017).
- [241] Manish Verma. 2023. Novel Study on AI-Based Chatbot (ChatGPT) Impacts on the Traditional Library Management. (2023).
- [242] Tu Vu, Brian Lester, Noah Constant, Rami Al-Rfou, and Daniel Matthew Cer. 2021. SPoT: Better Frozen Model Adaptation through Soft Prompt Transfer. *ArXiv abs/2110.07904* (2021).
- [243] Danilo Vucetic, Mohammadreza Tayarani, Maryam Ziaeeafard, James J Clark, Brett H Meyer, and Warren J Gross. 2022. Efficient Fine-Tuning of Compressed Language Models with Learners. *arXiv preprint arXiv:2208.02070* (2022).
- [244] Zhongwei Wan, Yichun Yin, Wei Zhang, Jiaxin Shi, Lifeng Shang, Guangyong Chen, Xin Jiang, and Qun Liu. 2022. G-MAP: General Memory-Augmented Pre-trained Language Model for Domain Tasks. In *2022 Conference on Empirical Methods in Natural Language Processing*.
- [245] Ben Wang and Aran Komatsuzaki. 2021. GPT-J-6B: A 6 billion parameter autoregressive language model.

- [246] Longyue Wang, Chenyang Lyu, Tianbo Ji, Zhirui Zhang, Dian Yu, Shuming Shi, and Zhaopeng Tu. 2023. Document-level machine translation with large language models. *arXiv preprint arXiv:2304.02210* (2023).
- [247] Ruijie Wang, Zheng Li, Dachun Sun, Shengzhong Liu, Jinning Li, Bing Yin, and Tarek Abdelzaher. 2022. Learning to sample and aggregate: Few-shot reasoning over temporal knowledge graphs. In *Advances in Neural Information Processing Systems*.
- [248] Ruize Wang, Duyu Tang, Nan Duan, Zhongyu Wei, Xuanjing Huang, Guihong Cao, Daxin Jiang, Ming Zhou, et al. 2020. K-adapter: Infusing knowledge into pre-trained models with adapters. *arXiv preprint arXiv:2002.01808* (2020).
- [249] Sheng Wang, Zihao Zhao, Xi Ouyang, Qian Wang, and Dinggang Shen. 2023. Chatcad: Interactive computer-aided diagnosis on medical image using large language models. *arXiv preprint arXiv:2302.07257* (2023).
- [250] Xinyi Wang, Zhenye Gong, Guoxin Wang, Jingdan Jia, Ying Xu, Jialu Zhao, Qingye Fan, Shaun Wu, Weiguo Hu, and Xiaoyang Li. 2023. ChatGPT Performs on the Chinese National Medical Licensing Examination. (2023).
- [251] Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, and Denny Zhou. 2022. Rationale-augmented ensembles in language models. *arXiv preprint arXiv:2207.00747* (2022).
- [252] Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc V Le, Ed H. Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. 2023. Self-Consistency Improves Chain of Thought Reasoning in Language Models. In *The Eleventh International Conference on Learning Representations*.
- [253] Yizhong Wang, Yeganeh Kordi, Swaroop Mishra, Alisa Liu, Noah A. Smith, Daniel Khashabi, and Hannaneh Hajishirzi. 2022. Self-Instruct: Aligning Language Model with Self Generated Instructions.
- [254] Yaqing Wang, Subhabrata Mukherjee, Xiaodong Liu, Jing Gao, Ahmed Hassan Awadallah, and Jianfeng Gao. 2022. Adamix: Mixture-of-adapter for parameter-efficient tuning of large language models. *arXiv preprint arXiv:2205.12410* (2022).
- [255] Yihan Wang, Si Si, Daliang Li, Michal Lukasik, Felix Yu, Cho-Jui Hsieh, Inderjit S Dhillon, and Sanjiv Kumar. 2022. Preserving In-Context Learning ability in Large Language Model Fine-tuning. *arXiv preprint arXiv:2211.00635* (2022).
- [256] Yanbo J Wang, Yuming Li, Hui Qin, Yuhang Guan, and Sheng Chen. 2022. A Novel DeBERTa-based Model for Financial Question Answering Task. *arXiv preprint arXiv:2207.05875* (2022).
- [257] Zihao Wang, Shaofei Cai, Anji Liu, Xiaojian Ma, and Yitao Liang. 2023. Describe, explain, plan and select: Interactive planning with large language models enables open-world multi-task agents. *arXiv preprint arXiv:2302.01560* (2023).
- [258] Jason Wei, Maarten Bosma, Vincent Zhao, Kelvin Guu, Adams Wei Yu, Brian Lester, Nan Du, Andrew M Dai, and Quoc V Le. 2022. Finetuned Language Models are Zero-Shot Learners. In *International Conference on Learning Representations*.
- [259] Jason Wei, Yi Tay, Rishi Bommasani, Colin Raffel, Barret Zoph, Sebastian Borgeaud, Dani Yogatama, Maarten Bosma, Denny Zhou, Donald Metzler, et al. 2022. Emergent abilities of large language models. *arXiv preprint arXiv:2206.07682* (2022).
- [260] Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed H Chi, Quoc V Le, Denny Zhou, et al. 2022. Chain-of-Thought Prompting Elicits Reasoning in Large Language Models. In *Advances in Neural Information Processing Systems*.
- [261] Stephen Wolfram. [n. d.]. ChatGPT Gets Its “Wolfram Superpowers”! <https://writings.stephenwolfram.com/2023/03/chatgpt-gets-its-wolfram-superpowers/>. Accessed: 2023-03-27.
- [262] Shijie Wu, Ozan Irsoy, Steven Lu, Vadim Dabravolski, Mark Dredze, Sebastian Gehrmann, Prabhanjan Kambadur, David Rosenberg, and Gideon Mann. 2023. Bloomberggpt: A large language model for finance. *arXiv preprint arXiv:2303.17564* (2023).
- [263] Zhuofeng Wu, Sinong Wang, Jiatao Gu, Rui Hou, Yuxiao Dong, V. G. Vinod Vydiswaran, and Hao Ma. 2022. IDPG: An Instance-Dependent Prompt Generation Method. In *North American Chapter of the Association for Computational Linguistics*.
- [264] Hana Demma Wube, Sintayehu Zekarias Esubalew, Firesew Fayiso Weldesellase, and Taye Girma Debelee. 2022. Text-based chatbot in financial sector: A systematic literature review. *Data Science in Finance and Economics* 2, 3 (2022), 232–259.
- [265] Ziang Xiao, Xingdi Yuan, Q Vera Liao, Rania Abdelghani, and Pierre-Yves Oudeyer. 2023. Supporting Qualitative Analysis with Large Language Models: Combining Codebook with GPT-3 for Deductive Coding. In *Companion Proceedings of the 28th International Conference on Intelligent User Interfaces*. 75–78.
- [266] Frank Z Xing, Erik Cambria, and Roy E Welsch. 2018. Natural language based financial forecasting: a survey. *Artificial Intelligence Review* 50, 1 (2018), 49–73.
- [267] Runxin Xu, Fuli Luo, Zhiyuan Zhang, Chuanqi Tan, Baobao Chang, Songfang Huang, and Fei Huang. 2021. Raise a child in large language model: Towards effective and generalizable fine-tuning. *arXiv preprint arXiv:2109.05687* (2021).
- [268] Hao Xue and Flora D Salim. [n. d.]. PromptCast: A New Prompt-based Learning Paradigm for Time Series Forecasting. ([n. d.]).
- [269] Jingfeng Yang, Haoming Jiang, Qingyu Yin, Danqing Zhang, Bing Yin, and Diyi Yang. 2022. SEZZERO: Few-shot Compositional Semantic Parsing with Sequential Prompts and Zero-shot Models. In *Findings of the Association for Computational Linguistics: NAACL 2022*.
- [270] Jingfeng Yang, Hongye Jin, Ruixiang Tang, Xiaotian Han, Qizhang Feng, Haoming Jiang, Bing Yin, and Xia Hu. 2023. Harnessing the Power of LLMs in Practice: A Survey on ChatGPT and Beyond. *arXiv* (2023).
- [271] Kai-Cheng Yang and Filippo Menczer. 2023. Large language models can rate news outlet credibility. *arXiv preprint arXiv:2304.00228* (2023).
- [272] Xi Yang, Aokun Chen, Nima PourNejatian, Hoo Chang Shin, Kaleb E Smith, Christopher Parisien, Colin Compas, Cheryl Martin, Anthony B Costa, Mona G Flores, et al. 2022. A large language model for electronic health records. *npj Digital Medicine* 5, 1 (2022), 194.
- [273] Xianjun Yang, Wei Cheng, Xujiang Zhao, Linda Petzold, and Haifeng Chen. 2023. Dynamic Prompting: A Unified Framework for Prompt Tuning. *ArXiv abs/2303.02909* (2023).

- [274] Xi Yang, Nima PourNejatian, Hoo Chang Shin, Kaleb E Smith, Christopher Parisien, Colin Compas, Cheryl Martin, Mona G Flores, Ying Zhang, Tanja Magoc, et al. 2022. GatorTron: A Large Language Model for Clinical Natural Language Processing. *medRxiv* (2022), 2022–02.
- [275] Zonghan Yang, Xiaoyuan Yi, Peng Li, Yang Liu, and Xing Xie. 2022. Unified Detoxifying and Debiasing in Language Generation via Inference-time Adaptive Optimization. *arXiv preprint arXiv:2210.04492* (2022).
- [276] Hongbin Ye, Ningyu Zhang, Shumin Deng, Xiang Chen, Hui Chen, Feiyu Xiong, Xi Chen, and Huajun Chen. 2022. Ontology-enhanced Prompt-tuning for Few-shot Learning. *Proceedings of the ACM Web Conference 2022* (2022).
- [277] Fangyi Yu, Lee Quartey, and Frank Schilder. 2022. Legal Prompting: Teaching a Language Model to Think Like a Lawyer. *arXiv preprint arXiv:2212.01326* (2022).
- [278] Elad Ben Zaken, Shauli Ravfogel, and Yoav Goldberg. 2021. Bitfit: Simple parameter-efficient fine-tuning for transformer-based masked language-models. *arXiv preprint arXiv:2106.10199* (2021).
- [279] Zhengran Zeng, Hanzhuo Tan, Haotian Zhang, Jing Li, Yuqun Zhang, and Lingming Zhang. 2022. An extensive study on pre-trained models for program understanding and generation. In *Proceedings of the 31st ACM SIGSOFT International Symposium on Software Testing and Analysis*. 39–51.
- [280] Xiaoming Zhai. 2022. ChatGPT user experience: Implications for education. *Available at SSRN 4312418* (2022).
- [281] Aston Zhang, Yi Tay, Shuai Zhang, Alvin Chan, Anh Tuan Luu, Siu Cheung Hui, and Jie Fu. 2021. Beyond fully-connected layers with quaternions: Parameterization of hypercomplex multiplications with $1/n$ parameters. *arXiv preprint arXiv:2102.08597* (2021).
- [282] Bo Zhang. [n. d.]. Preparing Educators and Students for ChatGPT and AI Technology in Higher Education. ([n. d.]).
- [283] Chaoning Zhang, Chenshuang Zhang, Sheng Zheng, Yu Qiao, Chenghao Li, Mengchun Zhang, Sumit Kumar Dam, Chu Myaet Thwal, Ye Lin Tun, Le Luang Huy, et al. 2023. A Complete Survey on Generative AI (AIGC): Is ChatGPT from GPT-4 to GPT-5 All You Need? *arXiv preprint arXiv:2303.11717* (2023).
- [284] Haojie Zhang, Ge Li, Jia Li, Zhongjin Zhang, Yuqi Zhu, and Zhi Jin. 2022. Fine-Tuning Pre-Trained Language Models Effectively by Optimizing Subnetworks Adaptively. *arXiv preprint arXiv:2211.01642* (2022).
- [285] Jieli Zhang, Zhongliang Zhou, Gengchen Mai, Lan Mu, Mengxuan Hu, and Sheng Li. 2023. Text2Seg: Remote Sensing Image Semantic Segmentation via Text-Guided Visual Foundation Models. *arXiv preprint arXiv:2304.10597* (2023).
- [286] Renrui Zhang, Jiaming Han, Aojun Zhou, Xiangfei Hu, Shilin Yan, Pan Lu, Hongsheng Li, Peng Gao, and Yu Qiao. 2023. Llama-adapter: Efficient fine-tuning of language models with zero-init attention. *arXiv preprint arXiv:2303.16199* (2023).
- [287] Rongsheng Zhang, Yinhe Zheng, Xiaoxi Mao, and Minlie Huang. 2021. Unsupervised domain adaptation with adapter. *arXiv preprint arXiv:2111.00667* (2021).
- [288] Zhuosheng Zhang, Aston Zhang, Mu Li, and Alex Smola. 2022. Automatic chain of thought prompting in large language models. *arXiv preprint arXiv:2210.03493* (2022).
- [289] Zhaowei Zhang, Hongyu Zhang, Beijun Shen, and Xiaodong Gu. 2022. Diet code is healthy: Simplifying programs for pre-trained models of code. In *Proceedings of the 30th ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering*. 1073–1084.
- [290] Hongyu Zhao, Hao Tan, and Hongyuan Mei. 2022. Tiny-Attention Adapter: Contexts Are More Important Than the Number of Parameters. *arXiv preprint arXiv:2211.01979* (2022).
- [291] Wenjie Zhao, Gaoyu Zhang, George Yuan, Jun Liu, Hongtao Shan, and Shuyi Zhang. 2020. The study on the text classification for financial news based on partial information. *IEEE Access* 8 (2020), 100426–100437.
- [292] Wayne Xin Zhao, Kun Zhou, Junyi Li, Tianyi Tang, Xiaolei Wang, Yupeng Hou, Yingqian Min, Beichen Zhang, Junjie Zhang, Zican Dong, et al. 2023. A Survey of Large Language Models. *arXiv preprint arXiv:2303.18223* (2023).
- [293] Denny Zhou, Nathanael Schärli, Le Hou, Jason Wei, Nathan Scales, Xuezhi Wang, Dale Schuurmans, Claire Cui, Olivier Bousquet, Quoc V Le, and Ed H. Chi. 2023. Least-to-Most Prompting Enables Complex Reasoning in Large Language Models. In *The Eleventh International Conference on Learning Representations*.
- [294] Jun-Jie Zhu, Jinyue Jiang, Meiqi Yang, and Zhiyong Jason Ren. 2023. ChatGPT and environmental research. *Environmental Science & Technology* (2023).
- [295] Daniel M Ziegler, Nisan Stiennon, Jeffrey Wu, Tom B Brown, Alec Radford, Dario Amodei, Paul Christiano, and Geoffrey Irving. 2019. Fine-tuning language models from human preferences. *arXiv preprint arXiv:1909.08593* (2019).
- [296] Mingyu Zong and Bhaskar Krishnamachari. 2022. Solving math word problems concerning systems of equations with gpt-3. In *Proceedings of the Thirteenth AAAI Symposium on Educational Advances in Artificial Intelligence*.
- [297] Guido Zuccon and Bevan Koopman. 2023. Dr ChatGPT, tell me what I want to hear: How prompt knowledge impacts health answer correctness. *arXiv preprint arXiv:2302.13793* (2023).