
Statistical Knowledge Assessment for Generative Language Models

Qingxiu Dong¹, Jingjing Xu², Lingpeng Kong³, Zhifang Sui¹ and Lei Li⁴

¹ National Key Laboratory for Multimedia Information Processing, Peking University

² Shanghai AI Lab ³ The University of Hong Kong ⁴ University of California, Santa Barbara
dqx@stu.pku.edu.cn, {jingjingxu, szf}@pku.edu.cn, lpk@cs.hku.hk, lilei@cs.ucsb.edu

Abstract

Given varying prompts regarding a factoid question, can a generative language model (GLM) reliably generate factually correct answers? Existing language models may generate distinct responses for different prompts. In this paper, we study the problem of quantifying knowledge contained in a GLM regarding a given set of facts. We propose KaRR, a statistical approach to assess factual knowledge for GLMs. The main idea is to estimate the ratio of GLM generating text corresponding to the answer entity given diverse prompts of subject and querying relation, versus it generating by random chances. Our assessment suite contains a comprehensive set of 994,123 entities and 600 relations, with 1,395,905 text aliases. We use our method to evaluate 14 GLMs of various sizes, including LLaMA, Alpaca, OPT, etc. Experiments show a strong correlation (0.43 Kendall’s τ) with human assessment of GLMs. Our results reveal that the knowledge in GLMs with the same backbone architecture adheres to the scaling law, while tuning on instruction-following data compromises the model’s capability to generate factually correct text consistently.

1 Introduction

Generative language models (GLMs) have achieved impressive performance on many tasks including text generation, question answering, and dialog generation [Radford et al., 2019, Brown et al., 2020, Thoppilan et al., 2022]. Recently, studies have shown that pretrained language models store factual knowledge in parameters [Dai et al., 2022a,b, Roberts et al., 2020], demonstrating considerable potential as assistants for responding to user questions based on their stored knowledge.

Despite the remarkable success of GLMs, critical concerns arise — GLMs often generate unreliable answers given varying prompts [Ji et al., 2023, Maharana et al., 2023, Chang and Bergen, 2023]. As shown in Fig. 1, reliability refers to the ability to consistently generate knowledge-correct text, posing a higher standard than accuracy. For example, although the Alpaca-7B [Taori et al., 2023] generates accurately (*playwright*) given a prompt of “William Shakespeare’s job”, it generates incorrect text (*boatman*) given a prompt of “the job of Swan of Avon is”¹. ChatGPT [OpenAI, 2022] correctly generates *playwright and teacher*. However, it answers “No” for the prompt “Is William Shakespeare a teacher?” In this paper, we assess the knowledge contained in a GLM and investigate whether a GLM is able to consistently generate factually correct answers given varying prompts. Knowledge assessment is a critical problem for GLMs. The assessment results directly affect the people’s trust in the GLM generated content. Once we identify inconsistency of GLM generation, we could potentially correct such knowledge in GLMs [Cao et al., 2021, Zhao et al., 2021, Dong et al., 2022].

¹Swan of Avon is a nickname of Shakespeare.

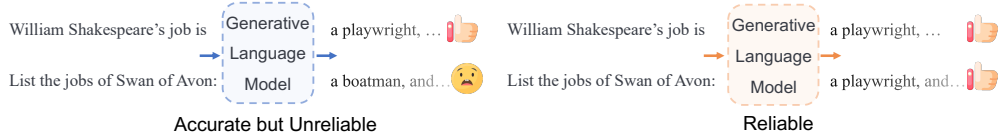


Figure 1: Accuracy pertains to generating the correct knowledge in response to a specific prompt, evaluated as a binary value (correct or incorrect). However, accuracy alone does not ensure reliability, as responses may vary with different prompt expressions. Reliability requires consistently generating correct knowledge for semantically similar prompts and is measured as a continuous value. Reliability sets a higher standard than accuracy by maintaining consistent accuracy across related prompts.

Important steps have been taken in the automatic knowledge evaluation for pre-trained language models [Petroni et al., 2019, Poerner et al., 2019, Jiang et al., 2020]. However, two issues still remain: 1) Prior methods focus on assessing the accuracy but not the consistency. Consequently, their probing results are prone to inflation or bias since models exhibit spurious correlations towards specific prompts [Poerner et al., 2019]. 2) Prior methods are designed for masked language models (MLMs) and are incapable of measuring GLMs. MLMs are evaluated by probing whether they predict the masked object in response to a given prompt in cloze form while the freely generated results of GLMs might be irrelevant to factual knowledge. Therefore, there is a need for a method that assesses both the accuracy and consistency of knowledge in GLMs.

In this paper, we introduce KaRR, a statistical approach to assess whether a generative language model contains reliable factual knowledge. We use triplets to represent probing facts (e.g. <William_Shakespeare, occupation, Playwright>). Our main idea is to estimate the ratio of the likelihood of GLM generating correct surface text for ground-truth object entities (Playwright) given varying prompts with subject entity (William_Shakespeare) and relation (occupation and the likelihood of generating entity text by pure chances. We distinguish the entities versus their surface text since each entity and relation may contain several text aliases. We introduce three latent variables to represent subject, object and relation, and a Bayesian network to represent the causal dependency among them and generated text. Our proposed KaRR quantifies the ratio of generating text with and without a specified relation/subject. This metric effectively captures the combined impact of both the subject and the relation in determining the model generation probability for plausible text of object (e.g., “playwright”, “dramatist”).

We develop a large-scale assessment suite with 994,123 entities, 600 relations, and their text aliases. Our suite is orders of magnitude larger than prior studies which only cover very few relations. We evaluate a comprehensive set of 14 GLMs using our method, including LLaMA [Touvron et al., 2023], OPT [Zhang et al., 2022], and others. Moreover, we also invite human experts to assess knowledge in these GLMs. Results reveal a strong correlation (0.43 Kendall’s τ) between KaRR and human evaluation. Our experimental results show that knowledge of GLMs with the same backbone architecture conforms to the scaling law.

In summary, our contributions are: 1) We propose KaRR, a statistical method to assess reliable knowledge contained in GLMs. We present a probabilistic graphical model to tackle varying text aliases regarding a fact. 2) We develop a large-scale assessment suite with 994,123 entities and 600 relations, which will be released to the community for further study. 3) We conduct a comprehensive evaluation of 14 recent GLMs. The experiments show that KaRR exhibits a strong correlation (0.43 Kendall’s τ) with human assessment, and achieves a lower variance to varying prompts, with a standard deviation of 0.82. Furthermore, our evaluation results are seldom affected by spurious correlations. Our results also reveal that instruction-tuning impairs the factual reliability of model generation.

2 Related Work

Knowledge Probing Previous works implement knowledge probing methods on MLMs. Petroni et al. [2019] introduces LAMA, studying whether MLMs correctly predict masked object entities in a cloze-style prompt. However, Elazar et al. [2021] shows that probing outcomes are inconsistent with various prompts, resulting in contradictory or unreliable results. Shin et al. [2020] show that well-engineered automatic prompts result in much higher results than handcrafted LAMA prompts. Therefore, ParaRel [Elazar et al., 2021] and LPAQA [Jiang et al., 2020] are proposed for MLMs. ParaRel measures the knowledge consistency of MLMs through the use of several paraphrased prompts. Subsequently, Hase et al. [2023] adopt their metric for paraphrase consistency as a measure of belief, and Raj et al. [2022] extend the metric to measure the general semantic consistency of

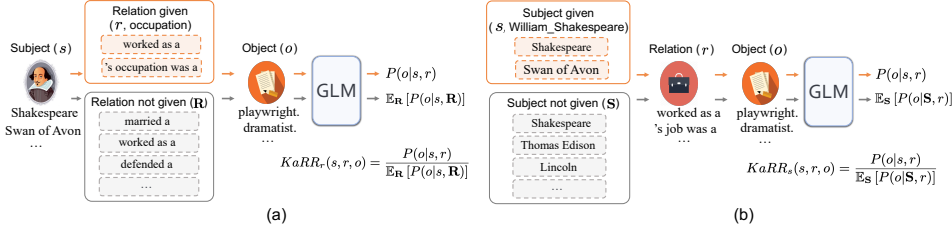


Figure 2: Illustration of KaRR using the fact <William_Shakespeare, occupation, playwright> as an example. S and R are latent variables representing the subject and relation, respectively; s , r and o refer to specific values of a subject, relation and object, respectively. (a) illustrates $KaRR_r$, which measures the impact of specifying r or not on GLM generating o given s . (b) illustrates $KaRR_s$, measuring the impact of specifying s or not on GLM generating o given r .

OPT [Zhang et al., 2022]. LPAQA focuses on prompts selection and ensembling to improve the prediction accuracy, by automatically discovering better prompts to query. However, it has not been comprehensively assessed to what extent GLMs contain knowledge about a fact.

Spurious Correlations in Probing Recent literature reveals that empirical probing results for masked entities may be affected by spurious correlations within the text [Poerner et al., 2019, Dong et al., 2022]. Spurious correlations occur when a model takes shortcuts to generate answers based on superficial information in surface forms, even without possessing relevant knowledge. For instance, a model with limited knowledge can correctly answer “French” when queried about the native language of actor “Jean Marais”, simply because “Jean Marais” is a typically French name. This phenomenon inflates probing results. On the one hand, Zhao et al. [2021] find that probing results are influenced by contextual biases, and Dong et al. [2022] also note frequency biases in rank-based probing. On the other hand, the subject’s class has been shown to introduce biases in predictions [Poerner et al., 2019]. To address these issues, Poerner et al. [2019] eliminate vulnerable relation classes, however, this approach restricts coverage and does not provide a comprehensive solution. In our statistical knowledge assessment, we account for both the subject and relation by decomposing their influence and using multiple aliases to minimize spurious correlations.

3 Statistical Knowledge Assessment

In this section, we present a statistical method for assessing knowledge in GLMs. We begin by introducing the graphical model for knowledge assessment, which serves as the foundation of our approach. Subsequently, we introduce the Knowledge Assessment Risk Ratio (KaRR), which we then convert into a computable form based on the graphical model.

3.1 Graphical Model for Knowledge Assessment




Figure 3: Graphical model for knowledge assessment.

Fig. 3 shows the graphical model that describes the knowledge symbols and their corresponding text forms. Here, we introduce latent variables S, R , and O to represent the symbolic subject, relation, and object, respectively. We use s, r , and o to refer to specific values of these variables, and the triplet (s, r, o) to denote a simple fact/knowledge, such as <William_Shakespeare, occupation, playwright>. We introduce α, β, γ to denote the random variables for the textual forms (textual aliases²) of (S) , (S, R) , (S, R, O) , respectively. Here, α, β , and γ refer to specific values of α, β , and γ , such as α being expressed as “Shakespeare”.

In general, fully observing the GLM probability on the latent variables of symbolic knowledge is infeasible. Although we can calculate the probability of the model generating “Shakespeare”, we cannot calculate the probability of it generating the symbol <William_Shakespeare>. This is because GLM pretraining focuses on textual input rather than symbolic representations, and each symbol may encompass a multitude of textual variations (paraphrased expressions) that signify the same entity or relation. Fortunately, we can establish a connection

²“Aliases” are alternative names for entities or relations, defined in Wikidata (<https://www.wikidata.org/wiki/Help:Aliases>).

between symbols and text forms, enabling the calculation of the model’s probability for specific textual forms. As demonstrated in the graphical model, a hollow circle indicates the variable is latent, and a shaded circle indicates the variable is observed. The parent node of α is \mathbf{S} , while \mathbf{S} and \mathbf{R} are parents to β . Additionally, γ has three parent nodes, namely \mathbf{S} , \mathbf{R} , and \mathbf{O} . This structure enables the calculation of model probability on symbols through their corresponding text forms. Therefore, the goal of consistent GLM knowledge assessment is to *estimate the model knowledge on symbols through the observable model probability across diverse corresponding textual forms*.

3.2 Knowledge Assessment Risk Ratio

For a given GLM \mathcal{M} , we define its factual knowledge correctness regarding a simple fact (s, r, o) as the extent to which it generates knowledge-correct answers to various textual prompts. Specifically, we assess the joint impact of subject and relation symbols on the GLM’s ability to generate the object symbol. This joint impact comprises two components: (1) the impact of specifying r or not on \mathcal{M} generating o given s and (2) the impact of specifying s or not on \mathcal{M} generating o given r .

Risk Ratio (RR) is a statistical measure employed to compare the likelihood of a specific event occurring in an exposed group versus an unexposed group [Sistrom and Garvan, 2004]. To quantify the impact of specifying symbols (or not) on the GLM generation accuracy, we adapt RR into a new metric for knowledge assessment called KaRR. For component (1), we consider the exposed group as the case where a relation symbol r is specified, and the unexposed group as cases where the relation is not specified. The outcome is the probability of model generating the object symbol o , representing any possible textual representation of o . We denote the RR of specifying relation r as KaRR_r :

$$\text{KaRR}_r(s, r, o) = \frac{P(o|s, r)}{\mathbb{E}_{\mathbf{R}}[P(o|s, \mathbf{R})]} = \frac{P(o|s, r)}{P(o|s)} \quad (1)$$

As the example illustrated in Fig. 2 (a), for the fact $\langle \text{William_Shakespeare}, \text{occupation}, \text{playwright} \rangle$, KaRR_r compare the probability of model generating any alias of the object entity playwright when the relation is specified as occupation (e.g., “*Shakespeare works as a*”, “*William Shakespeare’s occupation is a*”), with that when the relation is not specified, the expectation of model probability on generating an object alias under various possible relations such as “*Shakespeare married a*”. Similarly, the Risk Ratio (RR) for a specified subject s , denoted as KaRR_s , is defined as,

$$\text{KaRR}_s(s, r, o) = \frac{P(o|s, r)}{\mathbb{E}_{\mathbf{S}}[P(o|\mathbf{S}, r)]} = \frac{P(o|s, r)}{P(o|r)} \quad (2)$$

While the numerator for KaRR_s is the same as that for KaRR_r , it compares the numerator with the model’s probability when the subject is not specified. As shown in Fig. 2 (b), KaRR_s calculates the expected model probability of generating an object alias across multiple possible subjects (e.g., “*Dante worked as a*” and “*Thomas Edison worked as a*”). To effectively represent the combined impact of KaRR_s and KaRR_r , we calculate the joint influence as the geometric mean of them, and name it Knowledge Assessment Risk Ratio (KaRR). Formally, we define KaRR on (s, r, o) as,

$$\text{KaRR}(s, r, o) = \sqrt[2]{\text{KaRR}_r(s, r, o) \cdot \text{KaRR}_s(s, r, o)} \quad (3)$$

Intuitively, there are two components in KaRR: (1) the ratio of the chance of GLM generating o given s and r , and that of GLM generating o only given s ; (2) the ratio of the chance of M generating o given s and r , and that of GLM generating o only given r . In Appendix B, we also provide a comprehensive causal explanation for our knowledge assessment approach.

3.3 Computing KaRR using Graphical Model

The Knowledge Assessment Risk Ratio (KaRR), as indicated in Eq. (3), is formulated based on knowledge symbols, while directly computing the KaRR score using the definition (on symbols) is unfeasible. In light of this, our graphical model for knowledge assessment facilitates the implementation of KaRR by employing model probabilities on the text.

Based on Fig. 3, we can use the graphical model to evaluate the numerator of KaRR_s and KaRR_r in Eq. (1) and Eq. (2). To do so, we can use the following formula:

$$P(o|s, r) = \sum_{k=1}^{|\beta|} P(o, \beta_k|s, r) = \sum_{k=1}^{|\beta|} P(\beta_k|s, r) \cdot P(o|s, r, \beta_k) \quad (4)$$

The probability is computed by summing over all possible values of β , namely the set of aliases for (s, r) . For each specific value of β , β_k , we apply the product rule of probability and decompose

$P(o, \beta_k | s, r)$ into the probability of observing the relation β_k given s and r , denoted by $P(\beta_k | s, r)$, and the probability of observing the object o given the subject s , relation r , and the specific relation type β_k , denoted by $P(o | s, r, \beta_k)$. We use $P_{\mathcal{M}}$ to denote the generation probability of model \mathcal{M} . According to the law of total probability for marginal distribution, $P(o | s, r, \beta_k)$ can be expanded as,

$$P(o | s, r, \beta_k) = \sum_{j=1}^{|\gamma|} P(o, \gamma_j | s, r, \beta_k) = \sum_{j=1}^{|\gamma|} P_{\mathcal{M}}(\gamma_j | s, r, \beta_k) P(o | \gamma_j) \quad (5)$$

In the context of calculating the denominator of KaRR_r , it is important to note that each subject symbol may possess multiple aliases. For instance, the subject symbol s , `<William_Shakespeare>`, may be expressed with varying probabilities as α_1 ("*Shakespeare*"), α_2 ("*Swan of Avon*"), α_3 ("*William Shakespeare*"), and so forth. To account for this variability, we expand the denominator as demonstrated in Eq. (6) by leveraging principles of marginal distribution and Bayes Rule.

$$P(o | s) = \sum_{i=1}^{|\alpha|} P(o, \alpha_i | s) = \sum_{i=1}^{|\alpha|} P(\alpha_i | s) P(o | s, \alpha_i) \quad (6)$$

$P(o | s, \alpha_i)$ is obtained by summing the joint probabilities of o and γ_j given s, α_i under all possible γ_j . Mathematically, this can be expressed as:

$$P(o | s, \alpha_i) = \sum_{j=1}^{|\gamma|} P(o, \gamma_j | s, \alpha_i) = \sum_{j=1}^{|\gamma|} P_{\mathcal{M}}(\gamma_j | s, \alpha_i) P(o | s, \alpha_i, \gamma_j) \quad (7)$$

Subsequently, we obtain $\text{RR}_r(s, r, o)$ as follows,

$$\text{KaRR}_r(s, r, o) = \frac{P(o | s, r)}{P(o | s)} = \frac{\sum_{k=1}^{|\beta|} P(\beta_k | s, r) \sum_{j=1}^{|\gamma|} P_{\mathcal{M}}(\gamma_j | s, r, \beta_k) P(o | \gamma_j)}{\sum_{i=1}^{|\alpha|} P(\alpha_i | s) \sum_{j=1}^{|\gamma|} P_{\mathcal{M}}(\gamma_j | s, \alpha_i) P(o | s, \alpha_i, \gamma_j)} \quad (8)$$

For the denominator of KaRR_s , in the above example, $P(o | r)$ represents the expected probability of generating an alias for playwright when presented with a sentence that semantically conveys the occupation relation, over all the subject symbols. Consequently, we can expand $P(o | r)$ as $\sum_{u=1}^{|\mathbf{s}_u|} P(s_u | r) P(o | s_u, r)$. The computation of $P(o | s_u, r)$ is identical to that in Eq. (4). Combining these results, we can determine the value of $\text{KaRR}_s(s, r, o)$ as Eq. (9).

$$\text{KaRR}_s(s, r, o) = \frac{\sum_{k=1}^{|\beta|} P(\beta_k | s, r) \sum_{j=1}^{|\gamma|} P_{\mathcal{M}}(\gamma_j | s, r, \beta_k) P(o | \gamma_j)}{\sum_{u=1}^{|\mathbf{s}_u|} P(s_u | r) \cdot \sum_{k=1}^{|\beta|} P(\beta_k | s_u, r) \sum_{j=1}^{|\gamma|} P_{\mathcal{M}}(\gamma_j | s_u, r, \beta_k) P(o | \gamma_j)} \quad (9)$$

The KaRR score is calculated as the geometric mean of the results obtained from Eq. (8) and (9), and its detailed implementation is shown in Appendix A.

4 Experiments

We implement KaRR to examine knowledge in 14 GLMs and analyze its correlation with human assessment on model knowledge. Compared to large language models, knowledge graphs store a vast amount of structured information explicitly, which is ensured to be correct through manual construction and verification. Therefore, we implement KaRR based on existing large knowledge graphs, including T-REx [EISahar et al., 2018] and Wikidata [Vrandečić and Krötzsch, 2014].

4.1 Data and Settings

In the following paragraphs, we provide details on the models and settings used in our experiments, with additional information provided in Appendix C.

Symbolic Facts We utilize T-REx knowledge graph [EISahar et al., 2018] as our primary source of symbolic knowledge. T-REx comprises 11 million triples that are aligned with 3.09 million Wikipedia abstracts; its quality is validated by extensive crowdsourcing evaluation. In our main experiments, we consider all 600 English relations available in T-REx and sample a maximum of 20 facts per relation, resulting in a total of 10,691 facts for knowledge assessment.

Multiple Entity Aliases For the text forms of subjects and objects involved in calculating KaRR_s and KaRR_r , we search the entity aliases from Wikidata with Wikidata Integrator³. Ultimately, we obtain 1,349,474 aliases for 968,426 subjects and 368,511 aliases for 207,985 objects.

³<https://github.com/SuLab/WikidataIntegrator>

Method	Subj. Alias	Obj. Alias	Rel. Alias	Rel. Cvg.	Method	Recall	Kendall's τ	p-value
LAMA@1	✗	✗	✗	6.83%	LAMA@1	83.25%	0.17	0.10
LAMA@10	✗	✗	✗	6.83%	LAMA@10	65.81%	0.08	0.23
ParaRel	✗	✗	✓	6.33%	ParaRel	69.15%	0.22	0.02
KaRR	✓	✓	✓	100%	K-Prompts	78.00 %	0.32	0.03
					KaRR	95.18%	0.43	0.03

(a) Basic information.

(b) Human evaluation results.

Table 1: Basic information and human evaluation results. The abbreviations Subj., Rel., Obj., and Rel. Cvg. represent Subject, Relation, Object, and the coverage of English relations in T-REx, respectively. The methods evaluated through human evaluation are assessed using the same set of 410 randomly sampled facts. K-Prompts is implemented using shared basic information with KaRR.

Multiple Relation Templates Previously, Elazar et al. [2021] manually wrote paraphrased relation templates for 41 relations, covering 6% of the relation classes in T-REx. To expand the coverage and diversity of templates for **R**, we incorporated placeholders "[X]" and "[Y]" into the relation aliases in Wikidata. Subsequently, we utilized a fine-tuned Flan-T5 model⁴ for grammar correction to generate corresponding template candidates. We then filtered the templates with missing subjects or objects, and the remaining template candidates were all validated and corrected manually by two human annotators. Finally, we obtained 3,488 templates for 600 relations, with an average of 5.82 paraphrased templates per relation, achieving full coverage of the English relation classes in T-REx.

Overall KaRR Score The KaRR score assigned to each fact is a continuous value. If the KaRR score exceeds a predefined threshold, the fact is considered consistently known by the GLM. When applying our method to a set of facts for the general knowledge reliability of GLMs, the overall KaRR score is calculated as the proportion of facts with a KaRR score exceeding the threshold.

As the statistical information shown in Tab. 1 (a), our method exhibits a broader coverage of relations and entity aliases compared to previous studies. To implement KaRR, we initially sample the knowledge triplet from T-REx, which comprises the symbolic subject s , relation r , and object o . By leveraging entity aliases and relation templates that we have devised, we can generate multiple expressions for each element in the KaRR metric. For instance, α is implemented as the set of all searched subject aliases from our aliases datastore, and β represents the combination of all possible relation templates and subject aliases (by replacing the placeholder with the subject aliases).

4.2 Models

We explore the following GLMs for knowledge assessment: (1) traditional pretrained GLMs ($< 5B$), including GPT [Radford and Narasimhan, 2018], XLNet [Yang et al., 2019], GPT2-XL [Radford et al., 2019], GPT-NEO [Black et al., 2021], T5-large and T5-3B [Raffel et al., 2020]; (2) medium size GLMs ($\geq 5B$, $< 50B$), including GLM [Du et al., 2022], Dolly⁵, Moss⁶, LLaMA [Touvron et al., 2023] and its variants Alpaca [Taori et al., 2023] and Vicuna [Chiang et al., 2023]; (3) large GLMs ($\geq 50B$), including the 65B LLaMA [Touvron et al., 2023] and the 175B OPT [Zhang et al., 2022].

4.3 Results of Knowledge Assessment

We present the results of our statistical knowledge assessment on various GLMs in Tab. 2 (a). Among the small and the medium size GLMs, most models' KaRR scores are between 3 to 13 on the KaRR score, indicating that they often struggle to generate factually accurate sentences consistently. However, it is surprising to note that the 2.65B GPT-NEO outperforms the T5-3B with a 3.92 KaRR score. We hypothesize that this difference in performance may be attributed to the pretraining corpora. Specifically, GPT-NEO (trained on the Pile dataset [Gao et al., 2020]) primarily utilizes high-quality data, while T5 (pretrained on a combination of corpora) uses a more varied set of data sources. For larger GLMs, a comparison between the original LLaMA and the Alpaca model reveals that fine-tuning on instruction-following data could lead to improved alignment at the cost of sacrificing the ability to generate consistent and correct knowledge. Compared to the KaRR score of the original LLaMA, the KaRR score of Vicuna indicates that fine-tuning GLM on data collected from a more knowledgeable model could augment its knowledge. This finding is not commonly observed in previous studies but is consistent with our case study presented in Appendix D.

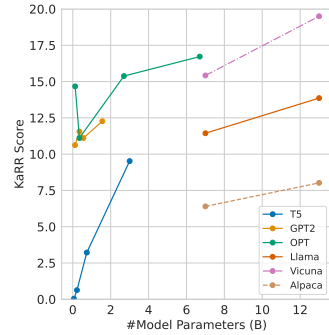
⁴<https://huggingface.co/pszemraj/flan-t5-large-grammar-synthesis>

⁵<https://github.com/databricks/dolly>

⁶<https://github.com/OpenMLab/MOSS>

Model	Model Size	KaRR Score	Model	Model Size	KaRR Score
< 5B			≥ 5B, < 50B		
GPT	0.12B	9.57	GLM	10B	5.59
XLNet	0.12B	5.86	Dolly	12B	15.60
T5-large	0.74B	3.22	LLaMA	13B	13.86
GPT2-XL	1.56B	12.27	Alpaca	13B	8.24
GPT-NEO	2.65B	13.44	Vicuna	13B	19.50
T5-3B	3B	9.52	Moss	16B	11.20
≥ 50B, < 100B			≥ 100B		
LLaMA	65B	14.56	OPT	175B	23.06

(a) Statistical knowledge assessment results on 14 GLMs.



(b) Scaling law of model knowledge.

Table 2: Statistical knowledge assessment on 14 GLMs. KaRR scores indicate that fine-tuning on instruction-following data (Alpaca) jeopardizes the reliability of knowledge-correct generation of the original LLaMA, whereas utilizing data from a more knowledgeable model enhances model knowledge (Vicuna). Additionally, the same backbone architecture adheres to the scaling law.

Fig. 2 demonstrates that our statistical evaluation approach supports the law of model scaling, where larger models tend to possess more factual knowledge. This observation is consistent with our intuition about the relationship between model size and knowledge representation. However, the benefits of scaling vary across models. For instance, T5 shows a significant improvement in KaRR score when scaled from T5-small to T5-3B, while GPT2 and OPT also benefit from scaling, albeit to a lesser extent. Notably, we observe that when the model size is smaller than 1B, GPT2 and OPT significantly outperform T5 in generating correct knowledge. This finding highlights the importance of selecting an appropriate model size for a given task and dataset, as larger models may not always provide the best performance on generation factuality. Overall, the results of our statistical evaluation method are consistent with our intuition, and highlight the importance of large-scale pretraining on knowledge-rich corpora for building high-performing language models.

4.4 Human Evaluation for Knowledge Assessment

To further validate our approach, we conduct the human evaluation for model knowledge assessment. Utilizing the mean score from human evaluation as the benchmark for knowledge assessment on individual facts, we examine the validity of our knowledge assessment outcomes from two perspectives: 1) the τ [Kendall, 1938] correlation between the KaRR score on each fact and the human-annotated score on each fact and 2) the Recall for identifying knowledge that is not consistently known (facts with human average scores below 0.5) when employing the GPT2-XL model.

Human Evaluation Our human evaluation for the reliable knowledge generation ability of GPT2-XL contains two procedures: 1) Annotating: We have three annotation volunteers with NLP backgrounds. Each annotator is asked to write three prompts to probe the model knowledge from multiple aspects for each fact. They are instructed to refine the prompts based on model generations until the generations are of the same type as the target answer (but not necessarily the target entity). To ensure a fair comparison with LAMA [Petroni et al., 2019], which is only applicable to 41 high-frequency relations, we restricted the relation classes to those included in LAMA. In total, we obtained 4,140 valid manually-annotated prompts for 410 facts in T-REx. 2) Rating: We then asked three other annotators to rate the knowledge (0 or 1) in GPT2-XL according to the model generations. We used the average score of their ratings on all the prompts as the gold standard for knowledge assessment on each fact. Detailed annotation procedures and scoring criteria are shown in Appendix E.

Baselines As a comparison, we also calculated the correlation between three baseline methods and human evaluation results. The baseline methods, LAMA [Petroni et al., 2019] and ParaRel [Elazar et al., 2021], were designed for MLMs, we modified them for GLMs by deleting words following the objects in the templates. We implement the *Consistent-Acc* score [Elazar et al., 2021] for ParaRel, which combines knowledge and consistency during evaluation. LAMA@1 and LAMA@10 refer to the results obtained by checking whether the top-1 or top-10 generations contain the object, respectively. Furthermore, we propose a K-Prompts baseline that computes the average probability of k randomly sampled prompts. It can be considered an adapted version of LPAQA [Shin et al., 2020] for GLMs, utilizing a collection of ensembled high-quality prompts identical to those used in KaRR. Here, we adopt K-Prompts as a simple baseline to evaluate the knowledge assessment in the absence of latent variables. All baseline methods are executed on the same set of facts as KaRR. LAMA and ParaRel are implemented as per the original papers [Petroni et al., 2019, Elazar et al., 2021], and

Method	Var (\downarrow)	Std (\downarrow)	Method	SP (\downarrow)	ΔP (\downarrow)
LAMA@1	1.90	1.37	LAMA@1	3.81	0.00
LAMA@10	5.14	2.27	LAMA@10	64.29	47.31
ParaRel	0.77	0.94	ParaRel	2.66	-0.51
K-Prompts	2.34	5.47	K-Prompts	0.00	-7.54
KaRR	0.67	0.82	KaRR	1.94	-14.94

(a) Evaluation variance towards varied prompts.

(b) Spurious correlation of knowledge assessment.

Table 3: Evaluation variance and spurious correlation analysis. Compared to LAMA and ParaRel, our statistical knowledge assessment is more robust and less influenced by the spurious correlation.

K-Prompts is implemented using shared aliases and templates with KaRR. We set 22 as the threshold of KaRR, which is chosen by aligning the proportion of GPT2-XL known facts distinguished by humans on a sampled set of 200 facts. Similarly, the threshold for K-Prompts has been set to 0.13, as a result of aligning with the proportion recognized by humans.

Results Tab. 1 (b) demonstrate that KaRR performs exceptionally well in knowledge detection, achieving a Recall of 95.18%. This highlights the potential of our statistical knowledge assessment to identify unreliable knowledge in GLMs, which is essential for downstream knowledge updating [Cao et al., 2021]. Furthermore, the evaluation results of KaRR show a strong correlation (Kendall’s 0.43) with human evaluation results. However, averaging the results of paraphrased prompts (ParaRel and K-Prompts) is not effective due to varying prompt frequencies and diverse impacts of different prompts on the model’s probability of generating an "object". Since human annotation for model knowledge assessment is demanding and time-consuming, our statistical evaluation method provides a fast and accurate solution for the automatic knowledge assessment of GLMs.

5 Variance toward Prompts

Since the aim of the evaluation is to reflect the correctness of a GLM’s knowledge towards various prompts, the evaluation result should be stable and not easily affected by the text forms used. In this section, we investigate the evaluation variance of our method. In contrast to ParaRel [Elazar et al., 2021], which studies the consistency of models under paraphrased inputs, we investigate the variance of the knowledge assessment method. Specifically, we evaluate the evaluation variance, which measures how much the same evaluation results vary when different prompts are used.

Following the same procedures as described in Sec. 4.4, we sampled 4,100 facts and employed the paraphrased relation templates from ParaRel [Elazar et al., 2021] to replace the relation templates in the baseline methods and KaRR, resulting in three distinct relation templates for each fact. We evaluated the effectiveness of the different evaluation approaches on each of the three paraphrased relation templates separately, the model is GPT2-XL. Additionally, we measured the evaluation variance by calculating the variance and standard deviation of the evaluation scores.

Tab. 3 (a) illustrates that the results of LAMA and K-Prompts exhibit significant variations when the probing prompt is changed to a paraphrased one, despite the overall prompt similarity in ParaRel being 0.96. In contrast, our statistical knowledge assessment approach demonstrates robustness to prompt variation, with a 0.67 score variance and a 0.82 standard deviation.

6 Spurious Correlation in Knowledge Assessment

Both LAMA and KaRR use text as a proxy to estimate unobservable knowledge. However, spurious correlations can lead to inflated probing results where GLMs "guess" the probing answer based on these correlations. For example, when probed with the relation template "[X]’s birthplace is [Y]", GLMs assign "America" the highest probability, regardless of the subject entity [Zhao et al., 2021]. To address this issue, we investigate spurious correlations in knowledge assessment, particularly focusing on the correlation between relation templates and high-frequency objects.

We synthesize false facts by replacing objects with high-frequency but factually incorrect ones. They are non-existent knowledge in GLMs as they are not part of the pretraining corpus. Ideally, the scores on synthetic facts should be 0 if the assessment is not influenced by spurious correlations. We obtain high-frequency objects by identifying the top-5 GPT2-XL predictions on 3 ParaRel relation templates without subjects, yielding 14 relations with valid high-frequency objects out of 41 relations. For each relation, we randomly sample 100 facts and replace the original object with the high-frequency object (if the original object is not the high-frequency object). We employ two metrics: Spurious Positive




Figure 4: Ablation study on the sampling parameter K. With the increase in K, the variability in knowledge assessment decreases, while its correlation with human evaluation improves.




Figure 5: KaRR scores on different relations when scaling up OPT from 350M to 2.7B and 175B.

(SP) and Positive Gap (ΔP). SP quantifies the proportion of false synthetic facts that are incorrectly recognized as known by the GLMs, and ΔP is calculated by SP minus model’s true positive rate on the actual facts. Ideally, SP and (ΔP) should be close to 0 and (ΔP) should be negative; otherwise, it indicates a significant bias in the evaluation result due to spurious correlations.

Tab. 3 (b) presents results on GPT2-XL. LAMA scores are highly correlated to the prompt, leading to the high ΔP and SP. For example, 3.81% and 64.29% false facts are estimated as learned facts by LAMA@1 and LAMA@10, respectively. In contrast, K-Prompts and KaRR are less influenced by spurious correlations, with an SP score lower than 2% and a negative ΔP . By assessing the model knowledge from various perspectives, the ultimate scores of KaRR indicate the comprehensive knowledge mastery of GLMs. Therefore, our results are less influenced by co-occurrence shortcuts.

7 Discussion

In this section, we explore the effects of the sampling parameter and the influence of model size.

Influence of Sampling Parameter When calculating KaRR, the sampling number K makes a trade-off between the evaluation accuracy and efficiency. We perform an ablation study on the results of KaRR implemented with different K. Following the same settings of Sec. 5 and Sec. 4.4, we analyze the evaluation variance and evaluation accuracy of the overall KaRR scores of GPT2-XL. Fig. 4 demonstrates that increasing K decreases variance in knowledge assessment, but the impact is minor since evaluation variance is already low. The sampling parameter has a more significant effect on evaluation accuracy, with higher accuracy correlating with higher K values. When K is less than 4, the evaluation results are unreliable due to weak correlation with human evaluation results.

Impact of Scaling on GLM Knowledge To further study the impact of scaling up on GLM knowledge, we compare the KaRR scores of OPT-350M, OPT-2.7B and OPT-175B on 100 best-known and 100 worst-known relations of OPT-350M. Fig. 5 shows that larger models exhibit better and more consistent knowledge-correct generation ability on the relations where small models also show some knowledge-correct generation ability. Moreover, larger models surpass small models in terms of knowledge on a wider range of relations. Therefore, the scaling law of GLM knowledge is attributed to the broader and more extensive knowledge of larger GLMs.

8 Conclusion

We investigate the consistent knowledge-correct generation ability of GLMs and propose a statistical knowledge assessment approach. By implementing our approach on 14 GLMs, we conduct an impartial evaluation of the factual reliability of current models. Our experiments reveal that our method yields a high correlation with human evaluation, and addresses the issues of variance and spurious correlation during knowledge assessment. Additionally, our assessment suite including millions of entity and relation aliases, enabling large-scale knowledge evaluation. Overall, our method evaluates the reliable knowledge generation ability of GLMs, providing a robust foundation for future developments in knowledge refinement, augmentation, and beyond.

Limitations

Despite the effectiveness of KaRR for GLM knowledge assessment, there are several limitations. First, KaRR scores are calculated for each fact using GLM inference on various prompts, which can be time-consuming and may limit the scalability of the evaluation. Second, the assessment is currently limited to scenarios where the logits of GLMs are available. This excludes many real-world applications where the models’ internal representations may not be accessible or may be difficult to obtain. Overall, these limitations highlight the need for further research to address the knowledge assessment approaches for GLMs.

References

- Sid Black, Gao Leo, Phil Wang, Connor Leahy, and Stella Biderman. GPT-Neo: Large Scale Autoregressive Language Modeling with Mesh-Tensorflow, March 2021. URL <https://doi.org/10.5281/zenodo.5297715>. If you use this software, please cite it using these metadata.
- Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. Language models are few-shot learners. In Hugo Larochelle, Marc’Aurelio Ranzato, Raia Hadsell, Maria-Florina Balcan, and Hsuan-Tien Lin, editors, *Advances in Neural Information Processing Systems 33: Annual Conference on Neural Information Processing Systems 2020, NeurIPS 2020, December 6-12, 2020, virtual*, 2020. URL <https://proceedings.neurips.cc/paper/2020/hash/1457c0d6bfc4967418bfb8ac142f64a-Abstract.html>.
- Nicola De Cao, Wilker Aziz, and Ivan Titov. Editing factual knowledge in language models. In *Proceedings of EMNLP*, 2021.
- Tyler A Chang and Benjamin K Bergen. Language model behavior: A comprehensive survey. *arXiv preprint arXiv:2303.11504*, 2023.
- Wei-Lin Chiang, Zhuohan Li, Zi Lin, Ying Sheng, Zhanghao Wu, Hao Zhang, Lianmin Zheng, Siyuan Zhuang, Yonghao Zhuang, Joseph E. Gonzalez, Ion Stoica, and Eric P. Xing. Vicuna: An open-source chatbot impressing gpt-4 with 90%* chatgpt quality, March 2023. URL <https://lmsys.org/blog/2023-03-30-vicuna/>.
- Damai Dai, Li Dong, Yaru Hao, Zhifang Sui, Baobao Chang, and Furu Wei. Knowledge neurons in pretrained transformers. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 8493–8502, Dublin, Ireland, May 2022a. Association for Computational Linguistics. doi: 10.18653/v1/2022.acl-long.581. URL <https://aclanthology.org/2022.acl-long.581>.
- Damai Dai, Wenbin Jiang, Qingxiu Dong, Yajuan Lyu, Qiaoqiao She, and Zhifang Sui. Neural knowledge bank for pretrained transformers. *arXiv preprint arXiv:2208.00399*, 2022b.
- Qingxiu Dong, Damai Dai, Yifan Song, Jingjing Xu, Zhifang Sui, and Lei Li. Calibrating factual knowledge in pretrained language models. In *the Conference on Empirical Methods in Natural Language Processing (EMNLP) - Findings*, December 2022.
- Zhengxiao Du, Yujie Qian, Xiao Liu, Ming Ding, Jiezhong Qiu, Zhilin Yang, and Jie Tang. Glm: General language model pretraining with autoregressive blank infilling. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 320–335, 2022.
- Yanai Elazar, Nora Kassner, Shauli Ravfogel, Abhilasha Ravichander, Eduard Hovy, Hinrich Schütze, and Yoav Goldberg. Measuring and improving consistency in pretrained language models. *TACL*, 9, 2021.

Hady ElSahar, Pavlos Vougiouklis, Arslan Remaci, Christophe Gravier, Jonathon S. Hare, Frédéric Laforest, and Elena Simperl. T-rex: A large scale alignment of natural language with knowledge base triples. In Nicoletta Calzolari, Khalid Choukri, Christopher Cieri, Thierry Declerck, Sara Goggi, Kôiti Hasida, Hitoshi Isahara, Bente Maegaard, Joseph Mariani, Hélène Mazo, Asunción Moreno, Jan Odijk, Stelios Piperidis, and Takenobu Tokunaga, editors, *Proceedings of LREC*, 2018.

Leo Gao, Stella Biderman, Sid Black, Laurence Golding, Travis Hoppe, Charles Foster, Jason Phang, Horace He, Anish Thite, Noa Nabeshima, et al. The pile: An 800gb dataset of diverse text for language modeling. *arXiv preprint arXiv:2101.00027*, 2020.

Peter Hase, Mona Diab, Asli Celikyilmaz, Xian Li, Zornitsa Kozareva, Veselin Stoyanov, Mohit Bansal, and Srinivasan Iyer. Methods for measuring, updating, and visualizing factual beliefs in language models. In *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics*, pages 2706–2723, 2023.

Paul W Holland. Statistics and causal inference. *Journal of the American statistical Association*, 81(396):945–960, 1986.

Ziwei Ji, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Ye Jin Bang, Andrea Madotto, and Pascale Fung. Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12):1–38, 2023.

Zhengbao Jiang, Frank F. Xu, Jun Araki, and Graham Neubig. How can we know what language models know? *TACL*, 2020.

Maurice G Kendall. A new measure of rank correlation. *Biometrika*, 30(1/2):81–93, 1938.

Adyasha Maharana, Amita Kamath, Christopher Clark, Mohit Bansal, and Aniruddha Kembhavi. Exposing and addressing cross-task inconsistency in unified vision-language models. *arXiv preprint arXiv:2303.16133*, 2023.

OpenAI. <https://openai.com/blog/chatgpt>, 2022.

Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. Bleu: a method for automatic evaluation of machine translation. In *Proc. of ACL*, pages 311–318, 2002.

Fabio Petroni, Tim Rocktäschel, Sebastian Riedel, Patrick Lewis, Anton Bakhtin, Yuxiang Wu, and Alexander Miller. Language models as knowledge bases? In *Proceedings of EMNLP-IJCNLP*, 2019.

Nina Poerner, Ulli Waltinger, and Hinrich Schütze. E-bert: Efficient-yet-effective entity embeddings for bert. *arXiv preprint arXiv:1911.03681*, 2019.

Alec Radford and Karthik Narasimhan. Improving language understanding by generative pre-training. 2018.

Alec Radford, Jeff Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners. 2019.

Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J Liu. Exploring the limits of transfer learning with a unified text-to-text transformer. *Journal of Machine Learning Research*, 21:1–67, 2020.

Harsh Raj, Domenic Rosati, and Subhabrata Majumdar. Measuring reliability of large language models through semantic consistency. *arXiv preprint arXiv:2211.05853*, 2022.

Adam Roberts, Colin Raffel, and Noam Shazeer. How much knowledge can you pack into the parameters of a language model? In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 5418–5426, Online, November 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.emnlp-main.437. URL <https://aclanthology.org/2020.emnlp-main.437>.

- Taylor Shin, Yasaman Razeghi, Robert L Logan IV, Eric Wallace, and Sameer Singh. Autoprompt: Eliciting knowledge from language models with automatically generated prompts. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 4222–4235, 2020.
- Christopher L Sistrom and Cynthia W Garvan. Proportions, odds, and risk. *Radiology*, 230(1):12–19, 2004.
- Rohan Taori, Ishaan Gulrajani, Tianyi Zhang, Yann Dubois, Xuechen Li, Carlos Guestrin, Percy Liang, and Tatsunori B. Hashimoto. Stanford alpaca: An instruction-following llama model. https://github.com/tatsu-lab/stanford_alpaca, 2023.
- Romal Thoppilan, Daniel De Freitas, Jamie Hall, Noam Shazeer, Apoorv Kulshreshtha, Heng-Tze Cheng, Alicia Jin, Taylor Bos, Leslie Baker, Yu Du, YaGuang Li, Hongrae Lee, Huaixiu Steven Zheng, Amin Ghafouri, Marcelo Menegali, Yanping Huang, Maxim Krikun, Dmitry Lepikhin, James Qin, Dehao Chen, Yuanzhong Xu, Zhifeng Chen, Adam Roberts, Maarten Bosma, Yanqi Zhou, Chung-Ching Chang, Igor Krivokon, Will Rusch, Marc Pickett, Kathleen S. Meier-Hellstern, Meredith Ringel Morris, Tulsee Doshi, Renelito Delos Santos, Toju Duke, Johnny Soraker, Ben Zevenbergen, Vinodkumar Prabhakaran, Mark Diaz, Ben Hutchinson, Kristen Olson, Alejandra Molina, Erin Hoffman-John, Josh Lee, Lora Aroyo, Ravi Rajakumar, Alena Butryna, Matthew Lamm, Viktoriya Kuzmina, Joe Fenton, Aaron Cohen, Rachel Bernstein, Ray Kurzweil, Blaise Aguera-Arcas, Claire Cui, Marian Croak, Ed H. Chi, and Quoc Le. Lamda: Language models for dialog applications. *ArXiv preprint*, abs/2201.08239, 2022. URL <https://arxiv.org/abs/2201.08239>.
- Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, et al. Llama: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971*, 2023.
- Denny Vrandečić and Markus Krötzsch. Wikidata: a free collaborative knowledgebase. *Communications of the ACM*, 57(10):78–85, 2014.
- Zhilin Yang, Zihang Dai, Yiming Yang, Jaime G. Carbonell, Ruslan Salakhutdinov, and Quoc V. Le. Xlnet: Generalized autoregressive pretraining for language understanding. In Hanna M. Wallach, Hugo Larochelle, Alina Beygelzimer, Florence d’Alché-Buc, Emily B. Fox, and Roman Garnett, editors, *Advances in Neural Information Processing Systems 32: Annual Conference on Neural Information Processing Systems 2019, NeurIPS 2019, December 8-14, 2019, Vancouver, BC, Canada*, pages 5754–5764, 2019.
- Susan Zhang, Stephen Roller, Naman Goyal, Mikel Artetxe, Moya Chen, Shuohui Chen, Christopher Dewan, Mona Diab, Xian Li, Xi Victoria Lin, et al. Opt: Open pre-trained transformer language models. *arXiv preprint arXiv:2205.01068*, 2022.
- Zihao Zhao, Eric Wallace, Shi Feng, Dan Klein, and Sameer Singh. Calibrate before use: Improving few-shot performance of language models. In *International Conference on Machine Learning*, pages 12697–12706. PMLR, 2021.

Appendix

A KaRR Implementation

To implement KaRR, we define a binary function $\delta_{text}(symbols)$ that takes the symbols as input and returns a value of either 0 or 1 depending on whether any alias of all the symbols appears in the text. By applying this function, we can transform $P(\beta_k|s, r)$ and $P(\alpha_i|s)$ into a format that is compatible with the model probability.

$$P(\beta_k|s, r) = \frac{P(\beta_k, s, r)}{P(s, r)} = \frac{P_{\mathcal{M}}(\beta_k) \cdot \delta_{\beta_k}(s, r)}{P(s, r)} \quad (10)$$

$$P(\alpha_i|s) = \frac{P(\alpha_i, s)}{P(s)} = \frac{P_{\mathcal{M}}(\alpha_i) \cdot \delta_{\alpha_i}(s)}{P(s)} \quad (11)$$

If the aliases β_k contain any alias of both s and r , then $P_{\mathcal{M}}(\gamma_j|s, r, \beta_k)$ is equal to $P_{\mathcal{M}}(\gamma_j|\beta_k)$. Otherwise, $P_{\mathcal{M}}(\gamma_j|s, r, \beta_k)$ is zero. Similarly, if the aliases α_i contain any alias of s , then $P_{\mathcal{M}}(\gamma_j|s, \alpha_i)$ is equal to $P_{\mathcal{M}}(\gamma_j|\alpha_i)$; otherwise, $P_{\mathcal{M}}(\gamma_j|s, \alpha_i)$ is zero.

Furthermore, the probabilities $P(o|\gamma_j)$ and $P(o|s, \alpha_i, \gamma_j)$ are both equal to $\delta_{\gamma_j}(o)$, where δ_{γ_j} is the binary function that checks whether any alias of the object o appears in the text γ_j . By substituting the transformed formulas into the equation Eq. (9), we obtain an implementation of KaRR:

$$\text{KaRR}_r(s, r, o) = \frac{P(s) \sum_{k=1}^{|\beta|} P_{\mathcal{M}}(\beta_k) \cdot \delta_{\beta_k}(s, r) \sum_{j=1}^{|\gamma|} P_{\mathcal{M}}(\gamma_j|\beta_k) \delta_{\gamma_j}(o)}{P(s, r) \sum_{i=1}^{|\alpha|} P_{\mathcal{M}}(\alpha_i) \cdot \delta_{\alpha_i}(s) \sum_{j=1}^{|\gamma|} P_{\mathcal{M}}(\gamma_j|\alpha_i) \delta_{\gamma_j}(o)} \quad (12)$$

Based on the graphical model depicted above, **S** and **R** are conditionally independent when **O** is unknown. Consequently, we can simplify the expression $\frac{P(s)}{P(s, r)}$ to $\frac{1}{P(r)}$. Therefore,

$$\text{KaRR}_r(s, r, o) = \frac{1}{P(r)} \cdot \frac{\sum_{k=1}^{|\beta|} P_{\mathcal{M}}(\beta_k) \cdot \delta_{\beta_k}(s, r) \sum_{j=1}^{|\gamma|} P_{\mathcal{M}}(\gamma_j|\beta_k) \delta_{\gamma_j}(o)}{\sum_{i=1}^{|\alpha|} P_{\mathcal{M}}(\alpha_i) \cdot \delta_{\alpha_i}(s) \sum_{j=1}^{|\gamma|} P_{\mathcal{M}}(\gamma_j|\alpha_i) \delta_{\gamma_j}(o)} \quad (13)$$

The calculation of subject replacement is analogous. When **O** is unknown, $\frac{1}{P(s_u|r) \cdot P(s, r) \cdot \frac{1}{P(s_u, r)}}$ equals to $\frac{1}{P(s)}$, then we have,

$$\begin{aligned} \text{KaRR}_s(s, r, o) &= \frac{\sum_{k=1}^{|\beta|} P_{\mathcal{M}}(\beta_k) \cdot \delta_{\beta_k}(s, r) \sum_{j=1}^{|\gamma|} P_{\mathcal{M}}(\gamma_j|\beta_k) \delta_{\gamma_j}(o)}{\sum_{u=1}^{|\mathbf{s}_u|} P(s_u|r) \cdot P(s, r) \sum_{k=1}^{|\beta|} P_{\mathcal{M}}(\beta_k) \frac{1}{P(s_u, r)} \cdot \delta_{\beta_k}(s_u, r) \sum_{j=1}^{|\gamma|} P_{\mathcal{M}}(\gamma_j|\beta_k) \delta_{\gamma_j}(o)} \\ &= \frac{1}{P(s)} \cdot \frac{\sum_{k=1}^{|\beta|} P_{\mathcal{M}}(\beta_k) \cdot \delta_{\beta_k}(s, r) \sum_{j=1}^{|\gamma|} P_{\mathcal{M}}(\gamma_j|\beta_k) \delta_{\gamma_j}(o)}{\sum_{u=1}^{|\mathbf{s}_u|} \sum_{k=1}^{|\beta|} P_{\mathcal{M}}(\beta_k) \cdot \delta_{\beta_k}(s_u, r) \sum_{j=1}^{|\gamma|} P_{\mathcal{M}}(\gamma_j|\beta_k) \delta_{\gamma_j}(o)} \end{aligned} \quad (14)$$

In practical scenarios, computing all possible values of α or β for the KaRR metric can be prohibitively expensive. To address this challenge, we can make the assumption that the symbols **S** and **R** follow a unified distribution, and then use Markov Chain Monte Carlo (MCMC) sampling to estimate the denominators by sampling K subjects or objects. This results in $\frac{1}{P(s_u)}$ or $\frac{1}{P(r_u)}$ being added to the denominators, offset by the $\frac{1}{P(r)}$ and $\frac{1}{P(s)}$ terms in Eq. (13) and (14). By employing this approach, we are able to improve the accuracy of the KaRR metric within a limited sampling budget and derive a new score based on the sampled subset. Our default sampling parameter, K , is set to 4.

The KaRR score is calculated as the geometric mean of the results obtained from Eq. (13) and (14).

B Causal Explanation for KaRR

In this section, we discuss the relationship between KaRR and causal estimation. Taking KaRR_r as an instance, we provide the causal explanation of our approach.

From a causal perspective, our problem can be defined as: “given a subject s , the causal effect of given the relation r or not for the GLM generating the object o ”. Assuming there are no confounders,

given a fixed s , the treatment (T) is whether a specific r is provided or not, the outcome (Y) is the probability of the model generating o , then the Average Treatment Effect (ATE) [Holland, 1986] is defined as:

$$\mathbb{E}[Y(1)] - \mathbb{E}[Y(0)] = \mathbb{E}[Y \mid T = 1] - \mathbb{E}[Y \mid T = 0] = P(o|s, r) - \mathbb{E}_{\mathbf{R}}[P(o|s, \mathbf{R})] \quad (15)$$

Note that in ATE, the minuend and subtrahend correspond to the denominators and numerator of KaRR_r in Eq. (1), respectively. Assuming no confounders are present, the only distinction lies in the mathematical operation utilized (subtraction or division). Both subtraction and division can be used to compare the difference between the exposed group and the unexposed group, with the only distinction being the difference in the threshold for the resulting value. As a result, ATE helps clarify and validate KaRR from a causal perspective.

C Experimental Details

C.1 Settings for Different GLMs

In this paper, we load the models (except the OPT-175B) from the Huggingface Models⁷ and implement the OPT-175B according to Metaseq⁸. The specific version of Vicuna employed in this paper is lmsys/vicuna-7b-delta-v1.1.

C.2 OOV Objects

We collect a vast range of aliases for objects, certain aliases may contain out-of-vocabulary (OOV) words. Consequently, the model generation probability for these aliases cannot be computed. To address this issue, we follow the approach suggested in LAMA [Petroni et al., 2019] and exclude OOV object aliases when implementing KaRR on each GLM.

C.3 Scoring Details

According to the KaRR implementation above, we can obtain the KaRR score on each fact. If the KaRR is higher than a threshold, the fact is regarded as a consistently known one by the GLM; otherwise not. When implementing the statistical framework on a large set of facts for the general knowledge correctness of GLMs, we get the proportion of known facts over the 10,691 symbolic facts as the (KaRR higher than the threshold) as the overall KaRR score. We set 22 as the threshold for KaRR, and the threshold is chosen by aligning the proportion of GPT2-XL known facts distinguished by humans on a sampled set of facts. Also as a result of alignment with human-recognized proportion, the threshold for the K-Prompts baseline implemented in this paper is set to 0.13.

C.4 Computing Resources

We load the GLMs from the Huggingface Models, the detailed configurations of evaluated LMs are listed on the corresponding Webpage for each model (<https://huggingface.co/models>). We conduct all the experiments on 5 A40 GPUs and 8 A100 GPUs.

D Case Study on GLM Generation

Please refer to Tab. 4.

E Human Annotation Specification

E.1 Main Procedures

In our Human Annotation Assessment (HKA), there are two primary procedures: generation-guided prompting and knowledge mastery scoring. Each procedure is carried out by a group of three volunteer annotators who have backgrounds in NLP.

⁷<https://huggingface.co/models>

⁸<https://github.com/facebookresearch/metaseq/tree/main>

Model	Question	Generation
LLaMA	The occupation of Swan of Avon was:	The Swan of Avon was a British cargo ship...
	William Shakespeare worked as a	playwright and actor in London
	Joseph Armstrong worked for:	The Daily Telegraph (London)
	The companies where Joseph Armstrong worked in are:	1. The Armstrong Company, 1900-1901 2. The Armstrong Company, 1901-1902...
Alpaca	The occupation of Swan of Avon was:	that of a boatman...
	William Shakespeare worked as a	a playwright, poet, and actor...
	Joseph Armstrong worked for:	Joseph Armstrong is a former employee of the following companies: Joseph Armstrong currently works for:...
	The companies where Joseph Armstrong worked in are:	Joseph Armstrong’s average salary was \$100,000...
Vicuna	The occupation of Swan of Avon was:	The Swan of Avon was a paddle steamer...
	William Shakespeare worked as a	a playwright and actor...
	Joseph Armstrong worked for:	Joseph Armstrong was a member of the following organisations:...
	The companies where Joseph Armstrong worked in are:	* Armstrong, Whitworth & Co....

Table 4: The generation cases of the original 7B LLaMA model and its instruction-tuned variants Alpaca and Vicuna. For each model, we provide the same Knowledge-related prompts.

E.1.1 Generation-Guided Prompting

In our HKA, a critical procedure is generation-guided prompting. That is, annotators continually modify their prompts according to both manual and automatic feedback from the model generations.

This approach differs from previous scenarios where annotators merely wrote relation templates for each fact to create prompts [Petroni et al., 2019]. Instead, we ask each annotator to write prompts from various perspectives, including prompting for entities, relations, or the entire fact. Moreover, we perform manual answer-type checking throughout the annotation process. Adjusting prompts according to the model’s generation is essential for accurately assessing the model’s knowledge. For example, because of the prompt template “[X] is a [Y]” in LAMA, 100% of the top-1 prediction for facts of relation ‘P31’ is a definite article or a predicate. This induces the underestimation of model knowledge while implementing LAMA@1. And for ‘P19’, although the prompt in LAMA, “[X] was born in [Y].” is a plausible one, most of the top-10 generations are dates or months. This indicates the significance of model-generation-guided prompt refining.

Manual Answer-type Checking To ensure that prompt accurately prompts the model to generate the relevant text of the current fact in the next word, instead of freely describing other content, we conduct manual answer-type checking. Only when the model generations are mostly of the same type of target answers, the prompts are deemed valid ones; otherwise, they will be further refined to a valid one.

Automatic Diversity Checking To explore the model generations towards different prompts, we implement an automatic diversity check for the prompts. Regarding the other valid prompts as references and the new prompt as the candidate, we calculate BLEU score [Papineni et al., 2002] and set a similarity threshold to 0.8 to filter similar prompts. Besides, annotates are asked to not only write prompts to probe the object entity of each fact but also write prompts to probe the subject, the relation, or the whole fact.

E.1.2 Knowledge Mastery Scoring

In the knowledge mastery scoring procedure, we enlist an additional three annotators to evaluate the model’s knowledge. Provided with a prompt and the corresponding GPT2-XL model generation (top-1 generation using beam search), annotators are asked to score the model’s knowledge mastery

on a scale between 0 (unknown) and 1 (known). Ultimately, the scores for prompts related to the same fact are averaged to determine the model’s knowledge mastery score for that particular piece of factual knowledge.

E.2 Annotation Specification

As it introduced above, our annotation mainly contains two parts: A. prompting, B. scoring. All the annotators are volunteer students majoring in Computational Linguistics in China. We told them the annotation intention and explained how the data would be used.

E.2.1 Prompting

Given a triplet fact, annotators are instructed to do their best to determine whether the knowledge is stored in the GLM and compose three prompts (from various perspectives) for each fact to probe it.

During the prompting process, annotators should create prompts that not only probe the object entity but also examine the subject, the relation, or the entire fact. For instance, for the fact <Obama, birthplace, Hawaii>, possible prompts could include:

- *One U.S. president was born in Hawaii, and his name was ____ (This prompt probes for the subject)*
- *Obama was born in ____ (this prompt probes for the object).*
- *Hawaii is Obama’s ____ (this prompt probes for the relation).*
- *The state where Barack Obama born is ____ (this prompt probes for the object).*
- *President Obama was born in Honolulu, ____ (this prompt probes for the object).*

To ensure the diversity of prompts, we employ the BLEU metric [Papineni et al., 2002] for automatic diversity checks. Specifically, we consider all valid prompts as references and the new prompt as the candidate, setting a threshold of 0.5 to filter out prompts that are similar to existing ones in terms of surface forms.

Annotators assess whether a given prompt effectively elicits knowledge based on the model-generated results. If the generated content is primarily related to the target aspect, the prompt is considered valid. For instance, when prompting for object entities, if the model’s output consistently includes a set of entities (both correct and incorrect ones are acceptable, as long as the type is broadly relevant), the prompt is deemed valid. However, if the generated text or tokens consist of stopwords or unrelated entity types, the prompt is classified as invalid. For example, the fact <China, capital, Beijing>, *The capital city of China is ____* is not a valid prompt for GPT2 as the generations are: “inching”, “home to”, “is now a”, “is described as” or “is the birthplace”. While *The capital city of China is called ____* is a valid prompt as the model generations are city names like: “Shanghai”, “Beijing” or “Hangzhou”.

E.2.2 Scoring

The scoring criteria are as follows:

- "Invalid": Factually unrelated generations or improper prompts.
- 1: The output produced during the greedy decoding process contains at least one alias of the correct answer, without any factually incorrect answers.
- 0: The generated content does not include any related entities or relations, nor are there any correct answers expressed in paraphrased words or phrases. In other words, the correct answer is not present in any of the generated responses.

E.3 Annotation Results

We ultimately obtained 4,140 manually annotated prompts for 410 valid facts. The overall sentence similarity among the prompts for the same fact is 0.78, which demonstrates good diversity compared to the 0.95 overall similarity of prompts in ParaRel [Elazar et al., 2021]. The Kappa score between human raters is 0.4.

While HKA provides relatively trustworthy and reliable results for assessing model knowledge correctness, it places high demands on human annotators and is time-consuming to implement for each fact in the model-specific annotation process. In our experience, annotating 100 facts on GPT2-XL takes an average of 4 hours per annotator. Consequently, HKA is more suitable as a solution for obtaining gold-standard results for knowledge assessment on a smaller scale. The Kendall's τ correlation between the results of a knowledge-assessing approach and HKA reflects the evaluation's accuracy.